

“A nickel ain’t worth a dime anymore.”

— Yogi Berra

“You know that I write slowly. This is chiefly because I am never satisfied until I have said as much as possible in a few words, and writing briefly takes far more time than writing at length.”

— Karl Friedrich Gauss
(1777-1855)

Lecture V THE GREEDY APPROACH

An algorithmic approach is called “greedy” when it makes decisions for each step based on what seems best at the current step. Moreover, once a decision is made, it is never revoked. It may seem that this approach is rather limited. Nevertheless, many important problems have special features that allow optimal solutions using this approach. Since we do not revoke our greedy decisions, such algorithms tend to be simple and efficient.

*greedy with no
regrets*

To make this concept of “greedy decisions” concrete, suppose we have some “gain” function $G(x)$ which quantifies the gain we expect with each possible decision x . View the algorithm as making a sequence x_1, x_2, \dots, x_n of decisions, where each $x_i \in X_i$ for some set X_i of feasible choices at the i th step. Greediness amounts to choosing the $x_i \in X_i$ which maximizes the value $G(x)$.


The greedy method is supposed to exemplify the idea of “local search”. But closer examination of many greedy algorithms will reveal some global information being used. Such global information is usually minimal or easily obtained, such as doing some global sorting step. Indeed, the preferred data structure for delivering this sorting information is the priority queue.

We begin with three simple classes of greedy problems: a toy version of bin packing, the coin changing problem and problems involving intervals. Next we discuss the Huffman coding problem and minimum spanning trees. An abstract setting for the minimum spanning tree problem is based on **matroid theory** and the associated **maximum independent set problem**. This abstract framework captures the essence of a large class of problems with greedy solutions.

§1. Joy Rides and Bin Packing

We start with an example of a greedy algorithm to solve a simple problem which we call **linear bin packing**. This problem belongs to a major topic in combinatorial algorithms called bin packing.

¶1. Amusement Park Problem. Suppose we have a joy ride in an amusement park where riders arrive in a queue. We want to assign riders into cars, where the cars are empty as they



arrive and we can only load one car at a time. There is a weight limit $M > 0$ for all cars. The number of riders in a car is immaterial, as long as their total weight is $\leq M$ pounds. We may assume that no rider has weight $> M$. The way that riders are put into cars is controlled by two policies:

- (1) **Online policy**: this means we must make a decision for each rider at the moment that he or she arrives at the head of the queue. This decision may depend on earlier riders in the queue, but not depend on subsequent riders. The decision is binary: “take the current car” or “take the next car”.
- (2) **First come, first ride policy** (FCFR). In other words, there is no “jumping the queue” in getting into a car.

Imagine what can happen if we only have the online policy, without the FCRC policy. After we ask a rider to “take the next car”, we go to the next rider and say “take the current car”. This would violate the FCRC policy. Therefore, whenever we make the decision “take the next car”, we must immediately dispatch the current car, and call for the next empty car.

Example. Let the weight limit be $M = 400$ (in pounds) and the weights of the riders in the queue be

$$w = (30, 190, 80, 210, 100, 80, 50) \quad (1)$$

where 30 represents the front of the queue. We can load the riders into 3 cars as follows:

$$Solution_1 : (30, 190, 80; 210, 100, 80; 50)$$

where the semi-colons separate successive car loads. For easy read, we henceforth drop the last digit of 0 from these weights, and simply write

$$Solution_1 : (3, 19, 8; 21, 10, 8; 5). \quad (2)$$

What is our goal in this problem? *It is to minimize the number of cars used.* $Solution_1$ uses three cars.

We may verify that $Solution_1$ conforms to the Online and FCRC policies. We call $Solution_1$ the **greedy solution** for the instance (1). In the greedy solution, we always make the decision “take the current car” if it does not violate the weight limitation. It is easy to see that the greedy solution is always exist and is unique.

But $Solution_1$ is not the only one to satisfy our policies. Here are two others:

$$Solution_2 : (3, 19; 8, 21; 10, 8, 5).$$

$$Solution_3 : (3, 19; 8, 21; 10, 8; 5).$$

They do not improve upon the greedy solution in terms of the number of cars. In fact, $Solution_4$ is worse because it uses 4 cars. Nevertheless we are prompted to ask if there exists instances where a non-greedy solution is better than the greedy one?

Leaving this question aside for now, let us illustrate our introductory remark about global information in greedy algorithms. Suppose we place riders into the cars in two phases. In phase one, we sort the weights w in (1), giving us a new queue:

$$Sort(w) = (1, 3, 5, 8, 8, 19, 21). \quad (3)$$

In phase two, we apply the greedy algorithm to $\text{Sort}(w)$, giving us the solution

$$\text{Solution}_4 : (1, 3, 5, 8, 8; 19, 21).$$

This uses only two cars, improving our greedy algorithm. Decisions exploiting sorting is using global information about the queue w . Thus Solution_4 has violated both our policies.

Suppose we now mildly relax our two policies:

- (1') **Online Peek policy**: The decision for the current rider may depend on all the riders up to the present, plus the next rider. Thus, we can “peek” to see who comes next.
- (2') **First Come Second Ride policy** (FCSR): In the queue $w = (w_1, \dots, w_n)$, rider w_i ($i = 1, \dots, n-1$) is guaranteed to be *at most* one car behind rider w_{i+1}

The following solution to (1)

$$\text{Solution}_5 : (3, 19, 8, 10; 21, 8, 5).$$

This is an improvement over the greedy solution, but it violates the FCSR policy since $w_5 = 10$ rides ahead of $w_4 = 21$. But it is obtained by an algorithm that follows the modified policies (1') and (2'): when rider $w_4 = 21$ came to the front of the queue, the current car has the riders $(3, 19, 8)$. So 21 does not fit into the current car. But peeking at the next rider $w_5 = 10$, we saw that it can fit into the current car. So we tell w_4 to “take the next car” and then tell w_5 to “take the current car”.

¶2. **Linear Bin Packing**. Formally, the joy ride problem is given an integer M and a sequence $w = (w_1, w_2, \dots, w_n)$ of weights satisfying $0 < w_i \leq M$ ($i = 1, \dots, n$). A **linear solution** simply breaks w into some $k \geq 1$ groups,

$$(w_1, w_2, \dots, w_{t(1)}; w_{t(1)+1}, \dots, w_{t(2)}; \dots; w_{t(k-1)+1}, \dots, w_n)$$

where the groups are determined by a sequence of $k+1$ **breakpoints** given by

$$0 = t(0) < t(1) < t(2) < \dots < t(k) = n. \quad (4)$$

Moreover, the i -th group is the subsequence $w(t(i-1), t(i)]$ where

$$w(i, j] = (w_{i+1}, w_{i+2}, \dots, w_j), \quad (i \leq j).$$

E.g., the greedy solution (2) has three breakpoints: $t(1) = 3, t(2) = 6, t(3) = 7$. The solution is **feasible** if each group has total weight at most M ,

$$M \geq w_{t(i-1)+1} + w_{t(i-1)+2} + \dots + w_{t(i)}$$

(for $i = 1, \dots, k$). A feasible solution is **optimal** if k is minimum over all feasible solutions. The **linear bin packing problem** is to compute an optimal linear solution for any input M, w .

¶3. **Greedy Algorithm**. It is instructive to write out the Greedy Algorithm for linear bin packing problem (a.k.a. joy ride problem) in **pseudo-code**. An important criterion for pseudo-code that it exposes the control-loop structures of the algorithm. Critical variables used in these control-loops should also be exposed. We are happy with English descriptions of variables, etc.

Let $w = (w_1, w_2, \dots, w_n)$ be the input sequence of weights and C denote a container (or bin or car) that is being filled with elements of w . Let the variable W keep track of the sum of the weights in C .

GREEDY ALGORITHM FOR LINEAR BIN PACKING:
Input: (M, w) where $w = (w_1, \dots, w_n)$, each $w_i \leq M$.
Output: Container sequence $(C_1; C_2; \dots; C_m)$ representing a linear solution.
 ▷ *Initialization*
 $C \leftarrow \emptyset, W \leftarrow 0$
 ▷ *Loop*
 for $i = 1$ to n
 ▷ *INVARIANT:* $M \geq W = \sum_{w \in C} x$
 if $(i = n \text{ or } M < W + w_i)$
 Append C to the output sequence
 $C \leftarrow \emptyset, W \leftarrow 0$
 $W \leftarrow W + w_i; C \leftarrow C \cup \{w_i\}$.

Pseudo-code is aimed at human understanding. It is deliberately short of any actual programming language. Why? Because real programming languages are really meant for computers (in the form of a compiler) to understand, but it is suboptimal for human understanding. Our pseudo-code exploits the power of mathematical notations *and* the linguistic structures of English which humans understand best. Of course, English can be replaced by any other natural language. Although there are many possible levels of detail, let us informally say that a pseudo-code is “detailed enough” when it achieves two purposes:

(P1) It can be “literally” transcribed into common programming languages, by exploiting common data types like arrays or linked lists.

(P2) We can analyze its complexity up to Θ -order. Operations on the common data types might actually contain iterative loops – but these are well-understood and can be given complexity bounds without a need to exposed them.

We claim that our Greedy Algorithm above is detailed enough. We ask the student to try their hand at showing (P1). Let us argue (P2): we claim that our greedy algorithm takes $O(n)$ time. There is a loop with n iterations. Inside the loop, each step is $O(1)$. For instance, we may assume that the container C is represented by a linked list, and so the step $C \leftarrow C \cup \{w_i\}$ amounts to appending to a linked list.

We need an important assumption which is often taken for granted: we assume the real RAM computational model of Chapter 1. In this model, we can compare and perform arithmetic operations on any two real numbers in constant time.

¶4. **Optimality of Greedy Algorithm.** It may not be obvious why the greedy algorithm produces an optimal linear solution. In any case, it is instructive to prove this.

Theorem 1 *The greedy solution is optimal for linear bin packing.*

Proof. Suppose the greedy algorithm outputs k bins as defined by the sequence of breakpoints

$$0 = t(0) < t(1) < t(2) < \dots < t(k) = n,$$

Let us compare the greedy solution with some optimal solution with breakpoints

$$0 = s(0) < s(1) < s(2) < \dots < s(\ell) = n.$$

By way of contradiction, assume

$$\ell < k. \tag{5}$$

Now we compare $s(i)$ with $t(i)$ for $i = 1, \dots, \ell$. Note that

(a) We have $s(1) \leq t(1)$ because $t(1)$ is obtained by the greedy method.

(b) We have $s(\ell) > t(\ell)$ because

$$t(\ell) < t(k) = n = s(\ell).$$

Therefore there is a smallest index i^* such that $s(i^*) > t(i^*)$. Clearly, $1 < i^* \leq \ell$. Then

$$s(i^* - 1) \leq t(i^* - 1) < t(i^*) < s(i^*). \tag{6}$$

The weights in the i^* -th bin of the optimal solution is given by the subsequence $w(s(i^* - 1) .. s(i^*))$. But according to (6), this subsequence contains

$$w(t(i^* - 1) .. t(i^*) + 1]. \tag{7}$$

This expression is meaningful since $t(i^*) + 1 \leq s(i^*) \leq n$. By definition of the greedy algorithm, the total weight in (7) exceeds M (otherwise the greedy algorithm would have added $w_{t(i^*)+1}$ to its i^* th car). This is our desired contradiction. **Q.E.D.**

§5. General Bin Packing. The joy ride problem is a restricted version of the following **bin packing problem**: *given a collection of items, to place them into as few bins as possible*. Each item has a weight (a positive real number) and the bins are identical, with a limited capacity. More precisely, given a multiset set $S = \{w_1, \dots, w_n\}$ of positive weights, and a bin capacity $M > 0$, we want to partition S into a minimum number of subsets such that the total weight in each subset is at most M . The subsets constitute the contents of a bin. Assume that each $w_i \leq M$. Unlike the joy ride problem, the weights are not ordered. A solution to this general bin packing problem is called a **globally optimal solution**. To avoid confusion with “linear bin packing”, or for emphasis, “bin packing” may also be called “**global** bin packing”.

Here is a simple **brute force algorithm** to solve global bin packing: given an instance (M, w) , we generate all $n!$ permutations of $w = (w_1, \dots, w_n)$, and for each permutation $\sigma(w)$, we compute the greedy solution of $\sigma(w)$. We return any greedy solution with the minimum number of bins.

Preview: Since $n!$ is exponential in n , the brute force algorithm takes exponential time. Since $n! = \Omega(n^k)$ for any constant k , it is not a polynomial-time algorithm. This problem, suitably formulated, is an example of an **NP-hard problem**. It is widely believed NP-hard problems do not have polynomial-time algorithms. We will return to this important topic at the end of the book.

Unlike linear bin packing, there are no ordering constraint on the weights. Although optimality in this general setting is very hard to achieve, the simple greedy solution for linear bin packing illustrates this idea: *we can convert a hard problem into a simple problem by introducing constraints on the possible solutions*. Moreover, we shall see that this simple problem might be turn out to be a good approximation for the optimal solution.

E.g., if $S = \{1, 1, 1, 3, 2, 2, 1, 3, 1\}$ and $M = 5$ then one solution is $\{3, 2\}, \{2, 3\}, \{1, 1, 1, 1, 1\}$, illustrated in Figure ?? . This solution uses 3 bins, and it is clearly optimal since each bin is filled to capacity.

¶6. Scale Invariance and Unit Weight Sequence Suppose A is an algorithm for bin packing. Given an input instance (M, w) , A will output a solution $Solution_A(M, w)$ to (M, w) . But we will be mainly interested in the number of bins (or cars) used by $Solution_A(M, w)$. So we let $A(M, w)$ denote the **size** of the $Solution_A(M, w)$, i.e., the number of bins. For instance, if A is the greedy algorithm G_1 , then $G_1(M, w)$ is the size of the greedy solution. Likewise $Opt(M, w)$ is the size given by an optimal algorithm for bin packing (see ¶5).

For any $c > 0$, we can transform an input instance (M, w) into another instance $(M/c, w/c)$ where $w/c = (w_1/c, \dots, w_n/c)$ if $w = (w_1, \dots, w_n)$. This is called a **scale transformation**. An algorithm A is **scale invariant** if it has the property that $A(M, w) = A(M/c, w/c)$ for all M, w, c . The Greedy algorithm G_1 and any optimal algorithm Opt are scale invariant (Exercise). Indeed, it seems unnatural to even conceive of algorithms that are not scale invariant! Henceforth, unless otherwise noted, we shall assume that our algorithms for any variation of the bin packing problem to be scale invariant. An instance (M, w) is called a **unit weight instance** if $M = 1$, and the corresponding w (with each $w_i \leq 1$) is called a **unit weight sequence**. It is convenient to go back and forth between some general capacity M and unit capacity 1. Moreover, if $M = 1$, we may omit it: thus we speak of an input instance w instead of $(1, w)$, and write $A(w)$ or $G_1(w)$ instead of $A(1, w)$ or $G_1(1, w)$.

¶7. How good is the Greedy Solution? How good is the greedy algorithm G_1 when viewed as an algorithm for general bin packing? We are not interested in goodness in terms of computational complexity. Instead we are interested in the quality of the output of G_1 , namely the size of its solution, $G_1(w)$. We shall compare $G_1(w)$ to $Opt(w)$, the optimal solution. Since $G_1(x)/Opt(w) \geq 1$, we want to upper bound the ratio $G_1(w)/Opt(w)$.

Theorem 2 For any unit weight sequence w ,

$$Opt(w) \geq 1 + \lfloor G_1(w)/2 \rfloor \quad (8)$$

Moreover, for each n , there are weight sequences with $Opt(w) = n$ for which (8) is an equality.

Proof. Suppose $G_1(w) = k$. Let the weight of i th output bin be W_i for $i = 1, \dots, k$. The following inequality holds:

$$W_i + W_{i+1} > 1. \quad (9)$$

To see this, note that the first weight v to be put into the $i+1$ st bin by the greedy algorithm must satisfy $W_i + v > 1$. This implies (9) since $W_{i+1} \geq v$. Summing up (9) for $i = 1, 3, 5, \dots, 2 \lfloor k/2 \rfloor - 1$, we obtain $\sum_{i=1}^k W_i > \lfloor k/2 \rfloor$. The last inequality implies that the optimal solution needs more than $\lfloor k/2 \rfloor$ bins, i.e., $Opt(w) \geq 1 + \lfloor k/2 \rfloor$. This proves (8). To see that the inequality (8) is sharp, consider the following¹ unit weight sequence of length n :

$$w = \begin{cases} (1, \frac{1}{n}, 1, \frac{1}{n}, \dots, 1, \frac{1}{n}) & \text{if } n = \text{even}, \\ (\frac{1}{n}, 1, \frac{1}{n}, \dots, 1, \frac{1}{n}) & \text{if } n = \text{odd}. \end{cases}$$

¹Thanks to Jason Y. Lee (2008) for this simplification.

The greedy solution uses n bins, but clearly $\text{Opt}(w) = 1 + \lfloor n/2 \rfloor$.

Q.E.D.

¶8. Approximate Bin Packing. We said the general bin packing problem is very hard (no polynomial-time algorithms are known). So it is essential to seek good approximations to general bin packing. Let A be any bin packing algorithm. To evaluate the performance of A , we look at the worst case ratio $A(w)/\text{Opt}(w)$, as w range over all unit input sequences. Let us consider the first definition that comes to mind: the **absolute approximation ratio** of A is defined as $\alpha_{abs}(A) := \sup_w A(w)/\text{Opt}(w)$. The problem is that $\alpha_{abs}(A)$ can be determined by those value of w for which $\text{Opt}(w)$ is bounded (cf. Theorem 2), not by those w where $\text{Opt}(w) \rightarrow \infty$.

For instance, suppose that we have algorithm A such that $A(w) \leq 2\text{Opt}(w) + 1$ for all n . Moreover, for each $n \in \mathbb{N}$, there is some unit weight sequences w such $\text{Opt}(w^n) = n$. We would conclude that $\alpha_{abs}(A) = 3$, and this is achieved by the weight sequence w where $\text{Opt}(w) = 1$. Yet, we would intuitively, like to say that “ $\alpha(A) = 2$ ”. That leads to our main definition: the (asymptotic) **approximation ratio** of A , defined as²

$$\alpha(A) := \limsup_{n \rightarrow \infty} a_n$$

where $a_n = \sup \left\{ \frac{A(w)}{\text{Opt}(w)} : \text{Opt}(w) \geq n \right\}$.

Lemma 3 *The greedy algorithm G_1 for Linear Bin Packing has approximation ratio*

$$\alpha(G_1) = 2. \quad (10)$$

Proof. Note that Theorem 2 implies that $\text{Opt}(w) \geq 1 + \lfloor G_1(w)/2 \rfloor \geq \frac{G_1(w)}{2} + \frac{1}{2}$. This implies

$$\frac{G_1(w)}{\text{Opt}(w)} \leq 2 - \frac{1}{\text{Opt}(w)}. \quad (11)$$

This implies $\alpha(G_1) \leq 2$ (Careful: we are not entitled to say “ $\alpha(G_1) < 2$ ”). Moreover, Theorem 2 also shows that for each odd integer $n \geq 2$, there are inputs with $\text{Opt}(w) = n$ for which (11) is an equality. This implies $\alpha(G_1) \geq 2 - \varepsilon$ for any $\varepsilon > 0$. By definition of \limsup , $\alpha(G_1) \geq 2$. This proves (10). **Q.E.D.**

¶9. First Fit Algorithm. In order to beat the approximation factor of 2 in (10), we must give up the “linearity constraint” of linear bin packing. We consider the following “non-linear” bin packing algorithm:

²The **limit superior** of an infinite sequence of numbers (x_1, x_2, \dots) is given by $\limsup_{n \rightarrow \infty} := \lim_{n \rightarrow \infty} \{\sup \{x_i : i \geq n\}\}$.

FIRST FIT ALGORITHM

Given $w = (w_1, \dots, w_n)$:

Initialize n empty bins: B_1, B_2, \dots, B_n .

For $i = 1, \dots, n$,

place w_i into the first bin B_j that it fits into.

Return the sequence of non-empty bins (B_1, \dots, B_m) .

We leave it as an Exercise to give a more precise pseudo-code for this algorithm – recall the goals of pseudo-code in ¶3.

This is called the **First Fit Algorithm** because each item is placed into the first (smallest j) bin B_j in which it fits. Note that each B_j , as long as it is not full, remains open for packing additional weights. So the worst case space used by this algorithm is linear in n . Let $FF(w)$ denote the number of bins used by the First Fit Algorithm.

In this context of approximation bin packing algorithms, what we called the Greedy algorithm $G(w)$ is also known as Next Fit (NF) Algorithm. Another such heuristic is the Best Fit (BF) in which we pack w_i into the bin whose residual capacity exceeds w_j by the smallest positive amount.

Example: Let $M = 7$ and $w = (3, 5, 4, 1, 3, 2, 3)$. It is easy to check that $\text{Opt}(M, w) = 3$. and greedy achieves $G_1(M, w) = 5$. The following simulation shows the First Fit Algorithm achieves $FF(M, w) = 4$:

i	B_1	B_2	B_3	B_4
1	③			
2	3	⑤		
3	3 ④	5		
4	3 4	5 ①		
5	3 4	5 1	③	
6	3 4	5 1	3 ②	
7	3 4	5 1	3 2	③

The behavior of FF has been well-studied, by Johnson, Ullman and others in the 1970s:

Theorem 4 $\alpha(FF) = 17/10$.

Let us first prove a lower bound on $\alpha(FF)$, somewhat less than the optimal bound of 1.7:

$$\alpha(FF) \geq 5/3 > 1.6\bar{6}. \quad (12)$$

Let w be a sequence of $18n$ weights given in 3 groups:

$$w = (\underbrace{0.15, \dots, 0.15}_{6n}, \underbrace{0.34, \dots, 0.34}_{6n}, \underbrace{0.51, \dots, 0.51}_{6n})$$

Clearly, $\text{Opt}(w) = 6n$ since the total weight is $0.15 + 0.34 + 0.51 = 1.0$ and so $6n$ bins is necessary and sufficient. But $FF(w) = 10n$. To see this, FF puts the first group into n bins (each bin $0.15 \times 3 = 0.9$ full), second group in $3n$ bins (each bin $0.34 \times 2 = 0.68$ full), and the last group in $6n$ bins (each bin 0.51 full). So $FF(w)/\text{Opt}(w) = 10n/6n = 5/3$, as claimed in (12).

The upper bound in Theorem 4 is from Ullman (1971). We give a compact proof from György Dósa and Jiří Sgall (2006). The **value** of each item $w_i \in w$ is $v(w_i) := \frac{6}{5}w_i + b(w_i)$ where $b(x)$ is the “bonus” defined by

$$b(x) = \begin{cases} 0 & \text{if } x \leq \frac{1}{6}, \\ \frac{3}{5}(x - \frac{1}{6}) & \text{if } \frac{1}{6} < x < \frac{1}{3}, \\ 0.1 & \text{if } \frac{1}{3} \leq x \leq \frac{1}{2}, \\ 0.4 & \text{if } x > \frac{1}{2}. \end{cases}$$

The bonus function $b(x)$ is continuous except at $x = 1/2$. For any set of weights B , let $b(B) = \sum_{x \in B} b(x)$ and $s(B) = \sum_{x \in B} x$.

GRAPH of $b(x)$:

For any set B of items, if $s(B) \leq 1$ then $v(B) \leq 1.7$: because $\frac{6}{5}s(B) \leq 1.2$ and it is sufficient to show that $v(B) \leq 0.5$. (see paper)

The next idea is to introduce sorting. Heuristically, it seems a good idea to pack the larger weights before the smaller one. Applying this heuristic to G_1 and to FF , we obtain what are known as **Next Fit Decreasing** and **First Fit Decreasing** Algorithms, respectively. Thus,

$$NFD(w) = G_1(\text{Sort}(w)), \quad FFD(w) := FF(\text{Sort}(w))$$

where $\text{Sort}(w)$ return the sequence w in non-increasing sorted order. Our previous bound that $\alpha(G_1) = 2$ is no longer valid, but we have this³ lower bound:

Lemma 5 $\alpha(NFD) \geq 5/3 = 1.6\bar{6}$.

Proof. Consider this sorted sequence of $3n$ weights in three equal size groups:

$$w = \left(\frac{1}{2} + e, \dots, \frac{1}{2} + e; \frac{1}{3} + e, \dots, \frac{1}{3} + e; \frac{1}{7} + e, \dots, \frac{1}{7} + e \right).$$

where $\frac{1}{2} + \frac{1}{3} + \frac{1}{7} + 3e = 1$. So $e = 1/126$. Thus $\text{Opt}(w) = n$. The greedy algorithm uses n bins for the first group. The first weight of $\frac{1}{3} + e$ fits into the previous bin. The next $n - 1$ weights fit into $(n - 1)/2$ bins (assume n is odd). The first weight of $\frac{1}{7} + e$ fits into the previous bin. The remaining $n - 1$ weights fit into $(n - 1)/6$ bins (assume $n - 1$ is divisible by 6). Thus $G_1^*(w) = n + (n - 1)/2 + (n - 1)/6 = (5n - 2)/3$. Thus $G_1^*(w)/\text{Opt}(w) = 5/3 - (2/3n)$.

Q.E.D.

Theorem 6

$$\alpha(FFD) = 11/9 = 1.2\bar{2}.$$

³From notes of Peter Sanders.

The lower bound comes from the following weight sequence $w = (w_1, \dots, w_{30n})$ where

$$w_i = \begin{cases} \frac{1}{2} + \varepsilon & (i = 1, \dots, 6n), \\ \frac{1}{4} + 2\varepsilon & (i = 6n + 1, \dots, 12n), \\ \frac{1}{4} + \varepsilon & (i = 12n + 1, \dots, 18n), \\ \frac{1}{4} - 2\varepsilon & (i = 18n + 1, \dots, 30n). \end{cases}$$

The optimal packing uses $9n$ bins: we can pack $6n$ bins with weights $(\frac{1}{2} + \varepsilon, \frac{1}{4} + \varepsilon, \frac{1}{4} - 2\varepsilon)$, and $3n$ bins with weights $(\frac{1}{4} + 2\varepsilon, \frac{1}{4} + 2\varepsilon, \frac{1}{4} - 2\varepsilon, \frac{1}{4} - 2\varepsilon)$. However, we see that FFD uses $11n$ bins: it fills the first $6n$ bins with $(\frac{1}{2} + \varepsilon, \frac{1}{4} + 2\varepsilon)$, the next $2n$ bins with $(\frac{1}{4} + \varepsilon, \frac{1}{4} + \varepsilon, \frac{1}{4} + \varepsilon)$, and the last $3n$ bins with $(\frac{1}{4} - 2\varepsilon, \frac{1}{4} - 2\varepsilon, \frac{1}{4} - 2\varepsilon, \frac{1}{4} - 2\varepsilon)$.

But for upper bound, we shall be contented with a simple argument⁴ for this weaker upper bound:

$$\alpha(FFD) \leq 1.5. \quad (13)$$

It suffices to prove that for any unit weight sequence w ,

$$FF^*(w) \leq 1.5 \cdot \text{Opt}(w) + 1. \quad (14)$$

Suppose $FF^*(w) = k$. Consider the first $m = \lfloor 2k/3 \rfloor$ bins, say B_1, \dots, B_m . We claim that $\text{Opt}(w) \geq m$. There two cases. (1) Suppose bin B_m contains a weight $> 1/2$ (for $i = 1, \dots, m$). Thus $\text{Opt}(w) \geq m$ as claimed. (2) Suppose bin B_m contains only weights that are $\leq 1/2$. Therefore, the remaining $k - m$ bins contains only weights that are $\leq 1/2$. This implies each of these $k - m$ bins contains at least two weights. So there are $2(k - m) = 2 \lceil k/3 \rceil \geq 2k/3 \geq m$ weights in these bins. Take any m of these weights v_1, v_2, \dots, v_m . None of these weights fits into the first m bins. Therefore $v_i + |B_i| > 1$ for $i = 1, \dots, m$ (where $|B_i|$ is the total weight contained in B_i), i.e., $\text{Opt}(w) \geq m$. So we have proved that $\text{Opt}(w) \geq m$. It follows that $\text{Opt}(w) \geq m = \lfloor 2k/3 \rfloor \geq (2k - 2)/3$, or $k = FF^*(w) \leq 1.5 \cdot \text{Opt}(w) + 1$. This conclude our demonstration of (13).

The efficient implementation of Step 4 in the First Fit algorithm is interesting: we could use a while loop, checking B_1, B_2, B_3, \dots in this order until we find a bin that can accept w_i . Each iteration of the while-loop is $O(n)$ time, with overall complexity $O(n^2)$. But let us improve this to $O(\log n)$, giving an overall complexity of $O(n \log n)$. We represent the bins as leaves of a (static) binary tree T of height $O(\lg n)$. The i th leaf stores the set of weights in bin B_i . Define the **residual capacity** of B_i to be $R_i := M - W_i$ where W_i is the sum of the weights in B_i . Each node u of T stores a key $u.\text{key}$ which is the maximum residual capacity of the bins in the subtree at u . Initially, each W_i is zero, and so $u.\text{key} = M$ for each node u . For each w_i , we can find the first bin B_j whose residual R_j is at least as large as w_i as follows:

```

INSERT( $w_i$ ):
1. Initialize  $u$  to the root of  $T$ .
2. While ( $u$  is not a leaf)
3.   If ( $u.\text{left}.\text{key} \geq w_i$ )
       $u \leftarrow u.\text{left}$ 
4.   else
       $u \leftarrow u.\text{right}$ 
5.  $u.\text{key} \leftarrow u.\text{key} - w_i$    $\triangleleft$   $u$  is a leaf
6. While ( $u$  is not the root)
7.    $u \leftarrow u.p$    $\triangleleft$  move to the parent
8.    $u.\text{key} \leftarrow \max\{u.\text{left}.\text{key}, u.\text{right}.\text{key}\}$ 

```

⁴Thanks to Chien-Chin Huang (fall 2013).

The first while-loop (Line 2) finds the leftmost bin B_j where w_i can be inserted; the second while-loop (Line 6) updates the keys along the path from B_j to the root.

¶10. **Exact General Bin Packing.** What if we just want to determine the value of $Opt(w)$? As far as we know, this is no easier than actually computing the feasible solution with $Opt(w)$ bins. This can be done by using the linear bin packing solution as a subroutine: we just cycle through each of the $n!$ permutations of $w = (w_1, \dots, w_n)$, and for each compute the greedy solution in $O(n)$ time. The optimal solution is among them. This yields an $\Theta(n \cdot n!) = \Theta((n/e)^{n+(3/2)})$ time algorithm. Here, we assume that we can generate all n -permutations in $O(n!)$ time. This is a nontrivial assumption; see §8 for how to do this.

The Θ -form of Stirling's approximation

We can improve the preceding algorithm by a factor of n , since without loss of generality, we may restrict to permutations that begins with an arbitrary w_1 (why?). Since there are $(n-1)!$ such permutations, we obtain:

Lemma 7 *The general bin packing problem can be optimally solved in $O(n!) = O((n/e)^{n+(1/2)})$ time in the real RAM model.*

We can extend this idea to improve the complexity by any polynomial factors in n (Exercise). Observe that by imposing restrictions on the space of possible solutions, we have turned a difficult problem like general bin packing into a feasible one like linear bin packing. The latter problem may be interesting on its own merit, but we see that it can also be used as a subroutine for solving the original problem.

¶11. **Two-Car Loading.** Consider the extension of linear bin packing where we simultaneously load two cars. Call these two cars the **front** and **rear cars**. This is a realistic joy ride scenario. It will mildly violate the first-come first-serve policy: a rider may be assigned to the rear car, while the next rider in the queue may be assigned to the front car. But this is the worst that can happen (people behind you in the queue can never be ahead by more than one car). If neither car can accommodate the new rider, we must **dispatch** the front car, so that the rear car comes to the front position and a new car empty becomes the rear car. We continue to have the “online restriction”, i.e., we must make the decision for each rider in the queue without using knowledge of who comes afterward. As usual, we assume that decisions are **irrevocable**. Once a rider has been assigned a car, it cannot be changed.

What is a good design G_2 for 2-car loading? The goal, as usual, is to minimize the number of cars used for any given input sequence w . we want to ensure G_2 is at least as good as G_1 , the loading policy of the original greedy algorithm (¶3). More precisely, for all $w = (w_1, \dots, w_n)$, we require

$$G_2(w) \leq G_1(w). \quad (15)$$

A trivial way to ensure (15) is to just imitate G_1 ! This means (15) is actually an equality, but it is not very interesting. We want a policy G_2 where, in addition to (15), there are many inputs w where $G_2(w) < G_1(w)$, with quantifiable advantages.

Consider the following definition of G_2 : *Load each rider into the front car if it fits, otherwise load into the rear car if it fits. If neither fits, dispatch the front car.* We leave as an exercise to show (15). But G_2 can be strictly better than G_1 : For instance, if $w = (30, 190, 80, 210, 90, 80, 50)$ and $M = 400$ in our example in (1), then the first fit policy gives an improvement: $G_2(w) = 2 < 3 = G_1(w)$.

We can extend the 2-car loading framework: these extensions):

- Allow decisions to be revoked. This means that, upon seeing the next rider in the queue, we are allowed to move one or more riders between the front and rear cars. An even stronger notion of revoking is to exchange a rider in one of the two loading cars with the rider at the head of the queue. Note that this stronger notion can be applied even in 1-car loading.
- Assume that two loading cars are in “parallel tracks” (left or right tracks). That means we can dispatch either car first. Note that this extension permits loading policies which are arbitrarily unfair in the sense that a rider may be placed into a car that is arbitrarily far ahead of someone who arrived earlier in the queue. So we might want to restrict the admissible loading policies.

¶12. **Extensions.** There are many ways to extend the bin packing problem. Let us consider the higher dimensional version of this problem. In particular, the packing of rectangles into rectangular bins. One interesting simplified 2-dimension problem is to consider one bin of unit width but infinite height: we want to pack the rectangles so as to minimize the height of the packing. A popular computer game is based on this problem.

What kind of bin packing is Tetris? Is it an online problem?

Let us consider another simplified 2-dimensional problem: the bins are unit squares, and the input is a sequence of squares. Let $w = (w_1, w_2, \dots, w_n)$ where $0 < w_i \leq 1$ represents an input square of width w_i . Things are complicated enough in 2D that we will begin with assuming that the w_i are already sorted by non-increasing weights.

We start with the simplest greedy packing heuristic that we can analyze. Consider the⁵ following **cell packing** idea: Suppose we subdivide each bin into **cells** using the usual quadtree method: a bin is a cell, and given any cell, we can recursively split it into 4 congruent subcells. We want to put the w_i 's into cells, at most one per cell! Moreover, each cell is guaranteed to be at least $1/4$ full. The latter condition is easy to fulfill: if a cell is not $1/4$ full, we can split it first, and use one of its child. It is not hard to prove that this gives an approximation ratio $\alpha(A) = 1/4$.

EXERCISES

Exercise 1.1: Show that the following algorithms are scale invariant:

- Greedy algorithm $G_1(M, w)$,
- Greedy sorted $G_1(M, \text{Sort}(w))$
- Optimal algorithm $\text{Opt}(M, \text{Sort}(w))$

◇

Exercise 1.2: Construct an algorithm $A(M, w)$ for linear bin packing that is NOT scale invariant.

◇

Exercise 1.3: We consider linear bin packing problem in which the weights w_i 's can be negative.

- Show that $G_1(w)$ is no longer optimal for linear bin packing.
- How bad can $G_1(w)$ compared to the optimal linear bin packing solution? Please quantify “badness” in some reasonable way.

◇

⁵Thanks to Shravas Rao (Fall'2013) for this suggestion.

Exercise 1.4: There are two places where our optimality proof for the greedy algorithm breaks down when there are negative weights. What are they? \diamond

Exercise 1.5: Consider the following “generalized greedy algorithm” in case w_i ’s can be negative. A solution to linear bin packing be characterized by the sequence of breakpoints, $0 = n_0 < n_1 < n_2 < \dots < n_k = n$ where the i th car holds the weights

$$[w_{n_{i-1}+1}, w_{n_i+2}, \dots, w_{n_i}].$$

Here is a greedy way to define these indices: let n_1 to be the largest index such that $\sum_{j=1}^{n_1} w_j \leq M$. For $i > 1$, define n_i to be the largest index such that $\sum_{j=n_{i-1}+1}^{n_i} w_j \leq M$. Note that this algorithm is no longer “online”. Either prove that this solution is optimal, or give a counter example. \diamond

Exercise 1.6: Give an $O(n^2)$ algorithm for linear bin packing when there are negative weights. HINT: Assume that when you solve the problem for (M, w) , you also solve it for each (M, w') where w' is a suffix of w . This is really the idea of dynamic programming (Chapter 7). \diamond

Exercise 1.7: Improve the bin packing upper bound in Lemma 7 to $O((n/e)^{n-(1/2)})$. HINT: Repeat the trick which saved us a factor of n in the first place. Fix two weights w_1, w_2 . Consider two cases: either w_1, w_2 belong to the same bin or they do not. \diamond

Exercise 1.8: Prove that $(n-1)! = \Theta((n/e)^{n-\frac{1}{2}})$. NOTE: Stirling’s approximation implies $(n-1)! = O((n/e)^{n-\frac{1}{2}})$. So you only need to show that $(n-1)! = \Omega((n/e)^{n-\frac{1}{2}})$. \diamond

Exercise 1.9: The previous exercise fixes two weights w_1, w_2 .
 (a) Generalize the idea by fixing any fixed number $k \geq 3$ of weights.
 (b) Further generalize the idea by fixing a non-constant number $k = k(n)$ of weights. \diamond

Exercise 1.10: We have the 2-car loading problem, but now imagine the 2 cars move along two independent tracks, say the left track and right track. Either car could be sent off before the other. We still make decision for each rider in an online manner, but our i th decision x_i now comes from the set $\{L, R, L^+, R^+\}$. The choice $x_i = L$ or $x_i = R$ means we load the i th rider into the left or right car (resp.), but $x_i = L^+$ means that we send off the left car, and put the i -th rider into a new car in its place. Similarly for $x_i = R^+$. Consider the following heuristic: let $C_0 > 0$ and $C_1 > 0$ be the residual capacities of the two open cars. Try to put w_i into the car with the smaller residual capacity. If neither fits, we send off the car with the smaller residual capacity (and put w_i into its replacement car). Call this the **best fit dual-track 2-car strategy**, and let $G_2''(w)$ be the number of cars used by this strategy on input w (say capacity is 1). This is open ended: compare $G_2''(w)$ to G_1 and G_2 . \diamond

Exercise 1.11: Let $G_2(w)$ and $G_2'(w)$ denote (respectively) the number of cars used when loading according to the First Fit and Best Fit Policies.
 (a) Show an example where $G_2(w) > G_2'(w)$.
 (b) Show an example where $G_2(w) < G_2'(w)$. \diamond

Exercise 1.12: In the text, we compare our 2-car loading policy against an optimal bin packing solution. Now we want to compare our 2-car loading policy with the performance of a **clairvoyant 2-car algorithm**. Clairvoyant means that the algorithm can see into the future (and thus knows the entire queue). However, it must still respect the online requirement – each rider must be assigned a car and this cannot be revoked later. \diamond

Exercise 1.13: (Open ended) Explore the revoking of decisions in 1- or 2-car loading. \diamond

Exercise 1.14: (Open ended) Quantify the improvements possible when loading 2 cars in parallel tracks instead of loading 2 cars in a single track. \diamond

Exercise 1.15: Weights with structure: suppose that the input weights are of the form $w_{i,j} = u_i + v_j$ and (u_1, \dots, u_m) and (v_1, \dots, v_n) are two given sequences. So w has mn numbers. Moreover, each group must have the form $w(i, i', j, j')$ comprising all $w_{k,\ell}$ such that $i \leq k \leq i'$ and $j \leq \ell \leq j'$. Call this a “rectangular group”. We want the sum of the weights in each group to be at most M , the bin capacity. Give a greedy algorithm to form the smallest possible number of rectangular groups. Prove its correctness. \diamond

Exercise 1.16: For $0 < r < 1$, let $Opt(r)$ denote the maximal number of identical disks of radius r that can be packed into a unit disk.

- (a) Device a greedy online algorithm for this problem, and analyze its performance relative to $Opt(r)$,
 (b) How good is your algorithm when $r = 1/2$? \diamond

Exercise 1.17: A **vertex cover** for a bigraph $G = (V, E)$ is a subset $C \subseteq V$ such that for each edge $e \in E$, at least one of its two vertices is contained in C . A **minimum vertex cover** is one of minimum size. Here is a greedy algorithm to find a vertex cover C :

GREEDY VERTEX COVER($G(V, E)$):

1. Initialize C to the empty set.
2. Choose a vertex $v \in V$ with the largest degree.
 Add vertex v to the set C , and remove vertex v and all edges that are incident on it from graph G .
3. Repeat step 2 until the edge set is empty.
4. The final set C is a vertex cover of the original G .

- (a) Show a graph G for which this greedy algorithm fails to give a minimum vertex cover.
 (b) Let $\mathbf{x} = (x_1, \dots, x_n)$ where each x_i is associated with vertex $i \in V = \{1, \dots, n\}$. Consider the following set of inequalities:

- For each $i \in V$, introduce the inequality

$$0 \leq x_i \leq 1.$$

- For each edge $i-j \in E$, introduce the inequality

$$x_i + x_j \geq 1.$$

If $\mathbf{a} = (a_1, \dots, a_n) \in \mathbb{R}^n$ satisfies these inequalities, we call \mathbf{a} a **feasible solution**. If each a_i is either 0 or 1, we call \mathbf{a} a 0 – 1 feasible solution. Clearly, there is a bijective

correspondence between the set of vertex covers and the set of 0 – 1 feasible solutions. denote the corresponding 0 – 1 feasible solution. We also write $|\mathbf{a}|$ for the sum $a_1 + \dots + a_n$. Suppose $\mathbf{x}^* = (x_1^*, \dots, x_n^*) \in \mathbb{R}^n$ is a feasible solution that minimizes the $|\mathbf{x}|$, i.e., for all feasible \mathbf{x} ,

$$|\mathbf{x}^*| \leq |\mathbf{x}|. \quad (16)$$

Call \mathbf{x}^* an **optimum vector**. Construct a graph $G = (V, E)$ where the optimum vector \mathbf{x}^* is not a 0 – 1 feasible solution.

(c) Show that \mathbf{x}^* is half-integer valued.

(d) Given a vector \mathbf{x} , define the **rounded vector** $\lfloor \mathbf{x} \rfloor$ where each component x_i is rounded to $\lfloor x_i \rfloor \in \{0, 1\}$. Note that rounding is usually defined up to some tie-breaking rule, viz., how to round 0.5. Show that, with a suitable tie-breaking rule, if \mathbf{x} is feasible solution then $\lfloor \mathbf{x} \rfloor$ is feasible 0 – 1 solution.

(e) Suppose C^* is a minimum vertex cover. Show that $|\lfloor \mathbf{x}^* \rfloor| \leq 2|C^*|$.

(f) Consider the following trivial vertex cover algorithm:

```

TRIVIAL VERTEX COVER( $G(V, E)$ ):
1. Initialize  $C$  to the empty set
2. While  $E \neq \emptyset$ 
3.   Remove  $u-v$  from  $E$ 
4.    $C \leftarrow C \cup \{u\}$    ◀ arbitrarily pick  $u$  or  $v$ 
5.   Remove from  $E$  any edge incident on  $u$  or  $v$ 
6. Return  $C$ 

```

Show that the output C is at most twice the size of the minimum vertex cover.

(g) Given that the trivial algorithm in (f) is also a 2-approximation algorithm for Vertex Cover, is there a reason to use the rounding algorithm $\lfloor \mathbf{x}^* \rfloor$ of (e)? \diamond

Exercise 1.18: Let \mathbf{x}^* (16) be the optimum solution in the previous question. numbers. Show that coordinates of \mathbf{x}^* must be half-integers $(0, \frac{1}{2}, 1)$. \diamond

Exercise 1.19: For $k \geq 1$, a **k -coloring** of a bigraph $G = (V, E)$ is a function $C : V \rightarrow \{1, \dots, k\}$. The coloring is **proper** if $u-v \in E$ implies $C(u) \neq C(v)$. The **chromatic number** of G is the smallest k such that there exists a proper k -coloring of G ; this number is denoted $\chi(G)$. Computing the chromatic number of bigraphs is one of the pure graph problems for which we do not know any polynomial time solution. But like the bin packing problem we can “linearize it”: Given an enumeration $\mathbf{v} = (v_1, \dots, v_n)$ of the vertices of G , we define the following **greedy coloring** of G relative to \mathbf{v} . For $i = 1, \dots, n$, let $N_i = \{v_j : j < i, v_j - v_i \in E\}$ denote the set of vertices with index less than i that are adjacent to v_i . For each $i = 1, \dots, n$, in order, we color v_i using the following rule: $C(v_1) = 1$. For $i \geq 2$, $C(v_i)$ is the smallest integer $k \geq 1$ such that $k \notin C(N_i) := \{C(v) : v \in N_i\}$.

(a) Show that if the enumeration is arbitrary, then the greedy coloring may be arbitrarily bad compared to $\chi(G)$.

(b) Suppose we first sort the vertices in order of increasing degrees. How bad can the greedy coloring be in this case?

(c) Show that there exists an enumeration whose greedy coloring is optimal.

(d) Using (d), establish an upper bound on the complexity of computing chromatic numbers. \diamond

END EXERCISES

§2. Coin Changing Problem

In the US, if you paid your grocer bill of \$5.75 with a ten dollar bill, you are very likely to get back four singles (\$1.00 bills) and a quarter (25 cents) as change. You would be surprised if you got three singles and 5 quarters, for example. This expectation suggests that a deterministic algorithm that all cashiers use. In fact, it is a greedy algorithm! What problem does this greedy algorithm solve? It is popularly known as the **coin changing problem** although it is clear that bills are also counted as coins here.

*greed and money,
yes, they go
together*

Let us explore some questions associated with coin changing. But first, we need some terminology. A **currency system** is a vector of increasing positive integers, $D = [d_1, d_2, \dots, d_m]$ where each $d_i \in D$ is called a **denomination**. The number m of denominations is the **order** of the system. Notice that we use square brackets $[\dots]$ for currency systems. For instance, the US currency system in daily use may be given by: $D = [1, 5, 10, 25, 100, 500, 1000, 2000, 5000, 10000]$ or more compactly,

$$D = [1, 5, 10, 25, \$1, \$5, \$10, \$20, \$50, \$100]. \quad (17)$$

In (17), we have omitted the rare two-dollar bill which has been a US denomination since 1976. For $x \in \mathbb{N}$, a **D -solution** for x is any vector $s \in \mathbb{N}^m$ such that the dot product $\langle s, D \rangle := \sum_{i=1}^m s_i d_i$ equals x . Call x the **value** of s . We say D is **complete** if every positive integer has a D -solution. It is easy to see that D is complete iff $d_1 = 1$. In other words, the humble penny is what makes our currency system complete.

Two D -solutions s, s' are **equivalent** denoted $s \equiv s'$, if they have the same value. For example, if $s = (0, 0, 0, 1, 4, 0, 0, 0, 0, 0)$ is a solution in the US currency system (17), then its value is $\langle s, D \rangle = \$4.25$ (the change we expected in the initial example). An equivalent solution would be $s' = (0, 0, 0, 5, 3, 0, 0, 0, 0, 0)$, i.e.,

$$(0, 0, 0, 1, 4, 0, 0, 0, 0, 0) \equiv (0, 0, 0, 5, 3, 0, 0, 0, 0, 0),$$

Henceforth, we omit reference to D if it is understood (so we speak of “solution” instead of “ D -solution”, etc).

Call s a **greedy solution** for x if s is lexicographically the largest among solutions for x : that means that if s' is another solution, then the *last* non-zero entry in the vector difference $s - s'$ is positive. For example, $s' = (0, 0, 0, 5, 3, 0, 0, 0, 0, 0)$ then s is lexicographically larger than s' . Note that we look at last (not first) non-zero entry because we order the vector D from smallest to largest denomination.

¶13. The Canonicity Problem. The **size** of a solution $s = (s_1, \dots, s_m)$ is given by $|s| := \sum_{i=1}^m s_i$. Thus “size” is the “number of coins” we get in change during a transaction. An **optimal solution** for x is any solution s for x of minimum size. Let $\text{Opt}_D(x)$ (resp., $\text{Grd}_D(x)$) denote any optimal (resp., the greedy) solution, assuming x is a multiple of d_1 . Clearly, we have $|\text{Opt}_D(x)| \leq |\text{Grd}_D(x)|$. We say D is **canonical** if $|\text{Grd}_D(x)| = |\text{Opt}_D(x)|$ for all $x \in \mathbb{N}$ that is a multiple of d_1 . For instance, $D = [1, 2, 4, 5]$ is not canonical. To see this, notice that greedy solution for $x = 8$ is $(1, 1, 0, 1)$ of size 3. The optimal solution is $(0, 0, 2, 0)$ with size 2. On the other hand $D = [1, 2, 4, 8]$ is canonical because for any x , the greedy solution (s_1, s_2, s_3, s_4) has $s_4 = \lfloor x/8 \rfloor$ which is necessary for optimality, and the remainder $x \bmod 8$ has a unique optimal solution in the system $[1, 2, 4]$. This argument shows that any binary currency system $[1, 2, 4, \dots, 2^{m-1}, 2^m]$ is canonical. The key open problem in coin changing is *to characterize all canonical systems*. Surprisingly, this problem remain unsolved.



*The 2 dollar bill
was introduced in
1862, discontinued
in 1966, and
reintroduced in
1976 for the US
Bicentennial.*

the main problem!

Besides canonicity, another desirable property is uniqueness: D has **uniqueness** if $\text{Opt}_D(x)$ is unique for every x that is a multiple of d_1 . For example, the system $[1, 2, 3]$ is a canonical system, but it is non-unique since $x = 4$ has two solutions, $(1, 0, 1) \equiv (0, 2, 0)$.

If $s = (s_1, \dots, s_m)$ and $s' = (s'_1, \dots, s'_m)$ then we write $s \leq s'$ to mean that $s_i \leq s'_i$ for all i . Call s a **subsolution** of s' . The following is easy to see: *Optimal solutions are closed under the subsolutions*. In other words, subsolutions of an optimal solution are optimal. This is a form of the Dynamic Programming Principle which we will address in Chapter 7.

To show non-canonicity, we need counter examples. If $|\text{Grd}_D(x)| > |\text{Opt}_D(x)|$, we call x a **counter example** for (the canonicity of) D . Suppose x is a minimum counter example for $D = [d_1, d_2, \dots, d_m]$, i.e., any value less than x is no counter example. Tien and Hu noted that if $\text{Opt}(x) = (s_1^*, \dots, s_m^*)$ is an optimum solution for such an x then

$$s_i \cdot s_i^* = 0 \quad (\text{for all } i = 1, \dots, m) \quad (18)$$

where $\text{Grd}_D(x) = (s_1, \dots, s_m)$. Suppose not. Say $s_i \cdot s_i^* > 0$. Then we can replace x by $x' = x - d_i$ to get a strictly smaller counter example. Why? Clearly, $\text{Grd}(x')$ is obtained from $\text{Grd}(x)$ by subtracting 1 from the i -th component. Also, $|\text{Opt}(x')| \leq |\text{Opt}(x)| - 1$. We know that $|\text{Opt}(x)| < |\text{Grd}(x)|$. Thus $|\text{Opt}(x')| \leq |\text{Opt}(x)| - 1 < |\text{Grd}(x)| - 1 = |\text{Grd}(x')|$. Thus x' is also a counter example, contradicting the minimality of x .

Let $q(D) = (q_1, q_2, \dots, q_m)$ where $q_i = \lceil d_{i+1}/d_i \rceil$ for each $i = 1, \dots, m-1$. Also let $q_m = \infty$. Thus $q(D)$ is roughly the multiple by which successive denominations increase. Then for any x , we claim

$$\text{Grd}(x) \leq q(D) \quad (19)$$

in a componentwise manner: if $\text{Grd}(x) = (s_1, \dots, s_m)$ then $s_i \leq q_i$ for all i . If (19) is violated at some i , say $s_i > q_i$, then we can construct a solution s' that is lexicographically larger than s (in particular, we can make $s'_{i+1} > s_i$). This contradicts the definition of s as the greedy solution.

¶14. Is the US Currency System canonical? Not all currency system is canonical – the pre-1971 British currency system is non-canonical (Exercise). Experience suggests that the US Currency system is canonical. But how do we prove this?

have you ever met a counter example?

Towards this end, let us say one currency systems D' is an **extension** of another D when D is a prefix of D' . Consider the following sequence of extensions:

$$D_1 = [1, 2, 4], \quad D_2 = [1, 2, 4, 5], \quad D_3 = [1, 2, 4, 5, 8].$$

We already noted that D_1 is canonical and D_2 is not. Thus, there are non-canonical extensions of canonical ones. Moreover, it can be shown that D_3 is canonical (Cai-Zheng). Thus, there are also canonical extensions of non-canonical ones. These examples hint at the difficulty of canonicity.

Consider two kinds of extensions of $D = [d_1, \dots, d_m]$:

- **Type A extension** is $D' = [d_1, \dots, d_m, d_{m+1}]$ where $d_{m+1} = d_m \cdot q$ for some $q \geq 2$. E.g., $[1, 5, 10]$ is a type A extension of $[1, 5]$.
- **Type B extension** is $D' = [d_1, \dots, d_m, d_{m+1}, d_{m+2}]$ where

$$\left. \begin{aligned} d_{m+1} &= d_m \cdot q && (\text{for some } q \geq 2), \\ d_{m+2} &= d_{m+1} \cdot q' + d_m \cdot r && (\text{for some } q' \geq 1, r \geq 1. \end{aligned} \right\} \quad (20)$$

Call (q, q', r) a set of **parameters** of the Type B extension. E.g., $[1, 5, 10, 20, 50]$ is a type B extension of $[1, 5, 10]$ with parameters $(q, q', r) = (2, 2, 1)$.

If $r > q$, then $(q, q' + 1, r - q)$ is also a set of parameters for the extension. By repeating this process, we conclude that there is a set of parameters in which $r < q$. Call it a **reduced set of parameters** for the Type B extension.

Examples: Trivially, $D = [d_1]$ and $D = [1, d_2]$ is canonical for any $d_1 \geq 1$ and $d_2 > 1$. The system $D = [1, 2, 3]$ is a Type B extension with parameters $(q, q', r) = (2, 1, 1)$. and $D'' = [1, 5, 10, 25]$ is a Type B extension of $[1, 5]$ with parameters $(q, q', r) = (2, 2, 1)$. In fact, the US currency system (17) is a sequence of Type A and Type B extensions of $[1]$. Indeed, this remains true even if we add the two dollar bill to the system. We now characterize how canonicity is preserved under these extensions:

Theorem 8 Let D be a canonical system, D' be a Type A extension of D , and D'' be a Type B extension of D with reduced parameters (q, q', r) .

- (i) If D is canonical, then D' is canonical.
- (ii) If D is canonical, then D'' is canonical iff $r < q \leq r + q'$.
- (iii) If D is uniquely canonical, then D'' is uniquely canonical iff $r < q < r + q'$.

Proof. Let us initially assume $D = [q_1]$ (we will see that this is without loss of generality). Then $D' = [d_1, qd_1]$ ($q \geq 2$) and $D'' = [d_1, qd_1, (qq' + r)d_1]$ ($r < q$).

(i) Suppose x is the minimum counter example for D' where $q \geq 2$. If the greedy D' -solution for x is $\text{Grd}_{D'}(x) = (s'_1, s'_2)$ then $s'_2 \geq 1$. If $\text{Opt}_{D'}(x) = (s_1^*, s_2^*)$ then by property (18), $s_2^* = 0$. Thus, we can view $\text{Opt}_{D'}(x)$ as a D -solution for x . Since D is canonical, this proves that $|\text{Opt}_{D'}(x)| \geq |\text{Grd}_D(x)|$. Thus,

$$\begin{aligned}
 |\text{Grd}_{D'}(x)| &> |\text{Opt}_{D'}(x)| && x \text{ is a counter example} \\
 &\geq |\text{Grd}_D(x)| && \text{just noted} \\
 &= s'_1 + qs'_2 && \text{since } \text{Grd}_D(x) = (s'_1 + qs'_2) \\
 &= |\text{Grd}_{D'}(x)| + (q-1)s'_2 \\
 &> |\text{Grd}_{D'}(x)| && q \geq 2 \text{ and } s'_2 \geq 1
 \end{aligned}$$

which is a contradiction.

(ii) Note that D' is canonical (by part(i)), and $r < q$ holds by definition of reduced parameters. We must show that $q \leq r + q'$ holds iff D'' is canonical. There are two cases: **CASE 1**, $q > r + q'$: In this case, we must provide a counter example. This is readily provided by the x such that

$$\text{Grd}_{D''}(x) = (q - r, 0, 1).$$

Then $(q - r, 0, 1) \equiv (0, q' + 1, 0)$ and we see that $|(q - r, 0, 1)| = q - r + 1 > q' + 1 = |(0, q' + 1, 0)|$.

CASE 2, $q \leq r + q'$: by way of contradiction, assume there is a counter example. Kozen-Zaks (Exercise) shows that the minimum counter example x has the form $\text{Grd}_{D''}(x) = (a, 0, 1)$ for some $a \geq 0$. But we know $(a, 0, 1) \equiv (a + r, q', 0)$. The optimal solution $s^* = (s_1^*, s_2^*, s_3^*)$ satisfies $s_3^* = 0$ (by (18)). Since (s_1^*, s_2^*) is equal to $\text{Grd}_{D'}(x)$ (by canonicity of D'), we deduce

$$s^* = \begin{cases} (a + r, q', 0) & \text{if } a + r < q \\ (a + r - q, 1 + q', 0) & \text{else.} \end{cases}$$

If $a + r < q$, $|s^*| = a + r + q' \geq a + 1 = |\text{Grd}_{D''}(x)|$. If $a + r \leq q$, $|s^*| = a + (r - q) + 1 + q' \geq a + 1 = |\text{Grd}_{D''}(x)|$. In either case, we have shown that $|s^*| \geq |\text{Grd}_{D''}(x)|$, i.e., the D'' -greedy solution for x is optimal. This contradicts the assumption that x is a counter example.

(iii) This is similar to part(ii). We must show that $q < r + q'$ iff D'' is uniquely canonical. Again, consider two cases. **CASE 1**, $q \geq r + q'$. We see that $(q - r, 0, 1) \equiv (0, q' + 1, 0)$ implies that the greedy solution $(q - r, 0, 1)$ has size that is either greater than or equal to $(0, q' + 1, 0)$. This shows it is either suboptimal or non-unique. This proves that D'' is not uniquely canonical. **CASE 2**, $q < r + q'$: We want to show that D'' is uniquely canonical. By way of contradiction, assume an x where $|\text{Opt}_{D''}(x)| \leq |\text{Grd}_{D''}(x)|$. The same argument as before will lead to the contradiction that $|\text{Opt}_{D''}(x)| > |\text{Grd}_{D''}(x)|$.

The arguments for parts(i)-(iii) assume $D = [d_1]$. But it is easy to see that if D is an arbitrary canonical system, there is no change in any argument, except for slightly more tedious notations such as $s = (0, \dots, 0, q - r, 0, q)$ instead of $s = (q - r, 0, q)$. **Q.E.D.**

As a result of this theorem, we conclude that the US currency system (17) is uniquely canonical.

¶15. Historical Notes. Chang and Gill [3] seems to be the first to study the coin change problem algorithmically. Kozen and Zaks characterized canonical system of up to order $m = 3$ [5], and our Type A and Type B extensions have roots there. Cai and Zheng [2] characterized canonical systems of order up to $m = 5$.

EXERCISES

Exercise 2.1: In a certain country, its currency system was originally $D = [1, 5, 10, 25, 50, 100, 200]$. A new king came along. As a mathematician, he wanted to honor $\pi = 3.1415\dots$ and so he decreed a new denomination 314. Is the new currency system canonical? \diamond

Exercise 2.2: A certain sovereign state had a complete and canonical currency system $D = (d_1, \dots, d_m)$. After a period of hyper-inflation, the state decreed that its pennies ($d_1 = 1$) are no longer legal currency. Is the new currency still canonical? What if the next denomination d_2 is also no longer legal? \diamond

Exercise 2.3: In 1971, the British denomination converted to a decimal system. The old system⁶ has these denominations:

$$\frac{1}{2}, 1, 3, 6, 12(\text{=shilling}), 24(\text{=florin}), 30(\text{=half-crown}), 60(\text{=crown}), 240(\text{=pound}).$$

- (a) Show that the old system is non-canonical.
 (b) Determine the largest possible value of $\text{Grd}(x) - \text{Opt}(x)$ in the old system. \diamond

Exercise 2.4: Show by a direct argument that the binary system $D = (1, 2, 4, \dots, 2^m)$ is a canonical system that is also unique. Direct argument means to avoid quoting our theorem on extensions. \diamond

Yogi Berra was referring to this in the introductory quote!

⁶There was an obsolete coin, the Guinea which is equal to 21 shillings.

Exercise 2.5: (Kozen-Zaks)

(a) Let $D = [1, d_2, d_3]$ where $d_3 = qd_2 + r$ and $0 \leq r < d_2$. Show that $(q+1)d_2$ is the minimum counter example.

(b) If $D = [1, d_2, \dots, d_m]$ is non-canonical, the minimum counter example $s = (s_1, \dots, s_m)$ satisfies $s_3 < x < s_{m-1} + s_m$. \diamond

Exercise 2.6: (Panagiotis Karras) The following problem arises in “compressing databases”. We are given a sequence $w = (w_1, \dots, w_n)$ of numbers and some $\epsilon > 0$. We say a sequence $x = (x_1, \dots, x_m)$ is an ϵ -**approximation** of w of **order** m if there is a sequence of m breakpoints (as in (4))

$$0 = t(0) < t(1) < t(2) < \dots < t(m) = n$$

such that for each original number w_i , the unique x_j ($j = 1, \dots, m$) such that $t(j-1) < i \leq t(j)$ provides an ϵ -approximation to w_i in the sense that

$$|w_i - x_j| \leq \epsilon.$$

Intuitively, this says that we can approximate the sequence w by a histogram with m steps. Let $\text{Min}(w, \epsilon)$ denote the minimum order of an ϵ -approximation of w . Design and prove an $O(n)$ greedy algorithm to compute the $\text{Min}(w, \epsilon)$. \diamond

Exercise 2.7: Suppose that our supermarket checkout suddenly ran out of quarters. All the other denominations remain available.

(a) Is the greedy algorithm still optimal?

(b) More generally, we say a canonical currency system $D = [d_1, \dots, d_m]$ is **robust** if it remains canonical under the loss of any single denomination d_i . Is the US currency system robust? \diamond


END EXERCISES

§3. Interval Problems

An important class of greedy algorithms involves intervals. Typically, we think of an interval $I \subseteq \mathbb{R}$ as a time interval, representing the duration of some activity.

¶16. Intervals as Activities. We will use half-open intervals of the form $I = [s, f)$ where $s < f$ to represent an activity that starts at time s and finishes before time f . Recall that $[s, f)$ is the set $\{t \in \mathbb{R} : s \leq t < f\}$, (we might say our activity intervals are “open ended”). Two activities **conflict** if their time intervals are not disjoint. We use half-open intervals instead of closed intervals so that the finish time of an activity can coincide with the start time of another activity without causing conflict. A set $S = \{I_1, \dots, I_n\}$ of intervals is said to be **compatible** if the intervals in S are pairwise disjoint (i.e., the activities in S are mutually conflict-free).

We begin with the **activities selection problem**, originally studied by Gavril. Imagine you have the choice to do any number of the following fun activities in one afternoon:



beach 12 : 00 – 4 : 00,
 swim 1 : 15 – 2 : 45,
 tennis 1 : 30 – 3 : 20,
 movie 3 : 00 – 4 : 30,
 movie 4 : 30 – 6 : 00.

You are not allowed to do two activities simultaneously. Assuming that your goal is to maximize your number of fun activities, which activities should you choose? Formally, the activities selection problem is this: *given a set*

$$A = \{I_1, I_2, \dots, I_n\}$$

of intervals, to compute a compatible subset of S that is optimal. Here optimal means “of maximum cardinality”. E.g., in the above fun activities example, an optimal solution would be to swim and to see two movies. It would be suboptimal to go to the beach. What would a greedy algorithm for this problem look like? Here is a generic version:

GENERIC GREEDY ACTIVITIES SELECTION:

Input: a set A of intervals

Output: $S \subseteq A$, a set of compatible intervals

▷ *Initialization*

Sort A according to some numerical criterion.

Let (I_1, \dots, I_n) be the sorted sequence.

Let $S = \emptyset$.

▷ *Main Loop*

For $i = 1$ to n

If $S \cup \{I_i\}$ is compatible, add I_i to S

Return(S)

Thus, S is a partial solution that we are building up. At stage i , we consider interval I_i , to either **accept** or **reject** it. Accepting means to make it part of current solution S . Notice the difference and similarities between this greedy solution and the one for joy rides.

But what greedy criteria should we use for sorting? Here are four possibilities. Note that sometimes, a criterion leads to a tie. In this case, you break ties arbitrarily. Notice that we say “sort in increasing order”, instead of the more accurate “sort in non-decreasing” order⁷. I prefer⁷ this direct formulation, *always assuming the rule to break ties arbitrarily*.

- (a) Sort I_i ’s in order of increasing finish times:
- swim, tennis, beach, movie 1, movie 2*

⁷The more accurate way of speaking slows down our processing of such statements.

(b) Sort I_i 's in order of increasing start times:

beach, swim, tennis, movie 1, movie 2

(c) Sort I_i 's in order of duration where the duration of activity I_i is $f_i - s_i$. Note that *movie 1, movie 2* and *swim* are tied, but breaking ties arbitrarily:

movie 1, movie 2, swim, beach, tennis

(d) Sort I_i 's in order of increasing conflict degree.

The conflict degree of I_i is the number of I_j 's which conflict with I_i . If the conflict degree of I_i is zero, clearly we can always add I_i into our set. Thus the ordering is:

movie 2, movie 1 or swim, beach or tennis

My class (Fall 2011) voted on which of these criteria yields the optimal solution. Tally: (a) 11, (b) 0, (c) 16, (d) 20.

We can combine these criteria: e.g., if one criterion leads to a tie, we can use another criterion to break ties. We now show that the first criterion (sorting by increasing finish times) leads to an optimal solution. In the Exercises, you will provide counter examples to the optimality of the other three criteria. Because of the possibility of ties, we distinguish two kinds of counter examples: “strong” counter examples do not depend on how ties are broken; “weak” counter example depends on how the ties are broken.

Let us prove the optimality of sorting by increasing finish times. We use induction, reminiscent of the joy ride proof. Let $S = (I_1, I_2, \dots, I_k)$ be the solution given by our greedy algorithm. If $I_i = [s_i, f_i)$, we may assume

$$f_1 < f_2 < \dots < f_k.$$

Suppose $S' = (I'_1, I'_2, \dots, I'_\ell)$ is an optimal solution where $I'_i = [s'_i, f'_i)$ and again $f'_1 < f'_2 < \dots < f'_\ell$. By optimality of S' , we have $k \leq \ell$. CLAIM: We have the inequality $f_i \leq f'_i$ for all $i = 1, \dots, k$. We leave the proof of this CLAIM to the reader. Let us now derive a contradiction if the greedy solution is not optimal: assume $k < \ell$ so that I'_{k+1} is defined. Then

$$\begin{aligned} f_k &\leq f'_k && \text{(by CLAIM)} \\ &\leq s'_{k+1} && \text{(since } I'_k, I'_{k+1} \text{ have no conflict)} \end{aligned}$$

and so I'_{k+1} is compatible with $\{I_1, \dots, I_k\}$. This is a contradiction since the greedy algorithm halts after choosing I_k because there are no other compatible intervals.

What is the running time of this algorithm? In deciding if interval I_i is compatible with the current set S , it is enough to only look at the finish time f of the last accepted interval. This can be done in $O(1)$ time since this comparison takes $O(1)$ and f can be maintained in $O(1)$ time. Hence the algorithm takes linear time after the initial sorting.

¶17. **Extensions, variations.** There are many possible variations and generalizations of the activities selection problem. Some of these problems are explored in the Exercises.

- Suppose your objective is not to maximize the number of activities, but to maximize the total amount of time spent in doing activities. In that case, for our fun afternoon example, you should go to the beach and see the second movie.
- Suppose we generalize the objective function by adding a weight (“pleasure index”) to each activity. Your goal now is to maximize the total weight of the activities in the compatible set.
- We can think of the activities to be selected as a uni-processor scheduling problem. (You happen to be the processor.) We can ask: what if you want to process as many activities

as possible using two processors? Does our original greedy approach extend in the obvious way? (Find the greedy solution for processor 1, then find greedy solution for processor 2).

- Alternatively, suppose we ask: what is the minimum number of processors that suffices to do all the activities in the input set?
- Suppose that, in addition to the set A of activities, we have a set C of classrooms. We are given a bipartite graph with vertices $A \cup C$ and edges is $E \subseteq A \times C$. Intuitively, $(I, c) \in E$ means that activity I can be held in classroom c . We want to know whether there is an assignment $f : A \rightarrow C$ such that (1) $f(I) = c$ implies $(I, c) \in E$ and (2) $f^{-1}(c)$ is compatible. REMARK: scheduling of classrooms in a school is more complicated in many more ways. One additional twist is to do weekly scheduling, not daily scheduling.

EXERCISES

Exercise 3.1: We gave four different greedy criteria for the activities selection problem.

- (a) We know one of them is optimal, but do not if sorting in increasing order degree-of-conflict is optimal. Show that the two of the other three criteria are suboptimal. EXTRA CREDIT: provide a *single* set of activities that serves as *strong* counter example to each of three criteria. It is possible to find a single counter example with 9 activities.
- (b) Naturally, each of the above four criteria has an inverted version in which we sort in decreasing order. Again, one of these inverted criteria is actually optimal, and the other three suboptimal. Prove the optimality of one, and provide counter examples for the other three. \diamond

Please draw figures as counter examples!

Exercise 3.2: Suppose the input $A = (I_1, \dots, I_n)$ for the activities selection problem is already sorted, by increasing order of their start times, i.e., $s_1 \leq s_2 \leq \dots \leq s_n$. Give an algorithm to compute a optimal solution in $O(n)$ time. Show that your algorithm is correct. \diamond

Exercise 3.3: Consider the activities selection problem with the following optimality criterion: to maximize the length $|A|$ of a set $A \subseteq S$ of activities. Define the **length** $|A|$ of a compatible set S to be the length of all the activities in S , where the length of an activity $I = [s, f]$ is just $|I| = f - s$. In case S is not compatible, its length is 0. Write $L(S)$ for the maximum length of any $A \subseteq S$. Let $A_{i,j} = \{I_i, I_{i+1}, \dots, I_j\}$ for $i \leq j$ and $L_{i,j} = L(A_{i,j})$.

(a) Show by a counter example that the following “dynamic programming principle” fails:

$$L_{i,j} = \max \{L_{i,k} + L_{k+1,j} : i \leq k \leq j-1\} \quad (21)$$

Would assuming that S is sorted by its start or finish times help?

- (b) Give an $O(n \log n)$ algorithm for computing $L_{1,n}$. HINT: order the activities in the set S according to their finish times, say,

$$f_1 \leq f_2 \leq \dots \leq f_n.$$

For $i = 1, \dots, n$, let L_i be the maximum length of a subset of $\{I_1, \dots, I_i\}$. Use an incremental algorithm to compute L_1, L_2, \dots, L_n in this order. \diamond

Exercise 3.4: Give a divide-and-conquer algorithm for the problem in previous exercise, to find the maximum length feasible solution for a set S of activities. (This approach is harder and less efficient!) \diamond

Exercise 3.5: Interval problems often arises from scheduling.

- (a) There is a 5 player game that lasts 48 minutes. In this game, any number of players can be swapped at any time. Suppose there are 8 friends what wants to play this game. Give a schedule for swapping players so that each of the 8 friends has the same amount of play time.
- (b) Suppose there is a n player game that lasts t minutes. Again, any number of players can be swapped at any time. There are m friends who wants to play this game. Prove that there is always a schedule to let each friend have the same amount of playtime.
- (c) Design an algorithm for (b) to schedule the swaps so that every one has the same amount of play time. \diamond

END EXERCISES

§4. Huffman Code

The problem of compressing information is central to computing and information processing. We shall study one problem whose solution is based on the greedy paradigm.

¶18. An Encoding Problem. It is best to begin with an informally stated problem:

- (P) Given a string s of characters (or letters or symbols) taken from an alphabet Σ , choose a *variable length code* C for Σ so as to minimize the space to encode the string s .

Before making this problem precise, it is helpful to know a principal application for such a problem. A computer file may be regarded as a string s , so problem (P) can be called the **file compression problem**. Often, the characters in computer files are extended ASCII characters. In modern computers, these files are called **text files**. Each extended ASCII character is represented by an 8-bit binary string, and so the alphabet Σ *may be* identified with the set $\{0,1\}^8$ of size $2^8 = 256$. But we prefer to view Σ as some abstract set⁸ of symbols, and view the extended ASCII encode as a function

$$C_{asc} : \Sigma \rightarrow \{0,1\}^8.$$

The ASCII code is a *fixed-length binary code*, i.e., $|C_{asc}(x)| = 8$ for all $x \in \Sigma$. So the ASCII encode of a file of m characters is a binary string of length $8m$. Can we do better?

Problem (P) suggests the use of **variable length code** to take exploit the relative frequency of characters in Σ . For instance, in typical English texts, the letters ‘e’ and ‘t’ are most frequent and it is a good idea to use shorter codes for them. On the other hand, infrequent letters like ‘q’ or ‘z’ could have longer codes. According to one statistics (see §5, Exercises), the frequencies of ‘e’, ‘t’, ‘q’, ‘z’ are 12.02, 9.10, 0.11, 0.07, respectively. An example of a variable length code is Morse code (see Notes at the end of this section). To see what additional properties are needed in variable-length codes, let us give some definitions:

⁸This abstraction is at the human-level of understanding, a pre-requisite to any formalization. Our familiar Latin alphabet a to z and Arabic numerals 0 to 9 are “directly given” in Σ . Formally, we could identify Σ with the set $\{0,1\}^8$ – this might be convenient in the current context. But it will be disconnected with the broader human context in which this enterprise will eventually be embedded in.

Was Gauss referring to the difficulty of compression in the opening quotation?

A (binary) **code** for Σ is an injective function

$$C : \Sigma \rightarrow \{0, 1\}^*.$$

A string of the form $C(x)$ ($x \in \Sigma$) is called a **code word**. The string $s = x_1x_2 \cdots x_m \in \Sigma^*$ is then encoded as

$$C(s) := C(x_1)C(x_2) \cdots C(x_m) \in \{0, 1\}^*. \quad (22)$$

This raises the problem of **decoding** $C(s)$, *i.e.*, recovering s from $C(s)$. For a general code C and s , we cannot expect unique decoding. A basic requirement for decodability is the ability to detect the boundary between code words in $C(s)$. One solution is to introduce a new symbol ‘\$’ and insert it between successive $C(x_i)$ ’s. If we insist on the binary alphabet for the code, this may force us to convert, say, ‘0’ to 00, ‘1’ to 01 and ‘\$’ to 11. This doubles the number of bits, probably wasteful.

¶19. Prefix-free codes. Our preferred solution for unique decoding is that C be **prefix-free**. This means that if $a, b \in \Sigma$ are distinct letters then $C(a)$ is not a prefix of $C(b)$. Clearly, this ensures unique decoding. With suitable preprocessing (basically to construct the “code tree” for C , defined next) decoding can be done quite simply, in an on-line fashion.

We represent a prefix-free code C by a binary tree T_C with $|\Sigma|$ leaves. Each leaf in T_C is labeled by a character $b \in \Sigma$ such that the path from the root to b is encoded by $C(b)$ in the natural way: the path length has length $|C(b)|$, with the 0-bit (resp., 1-bit) representing a left (resp., right) branch. We call T_C a **code tree** for C .

Figure 2 shows two such trees representing prefix codes for the alphabet $\Sigma = \{a, b, c, d\}$. The first code, for instance, corresponds to $C(a) = 00$, $C(b) = 010$, $C(c) = 011$ and $C(d) = 1$.




Figure 2: Code trees for two prefix-free codes: assume $f(a) = 6$, $f(b) = 2$, $f(c) = 1$ and $f(d) = 2$.

For $|\Sigma| \geq 2$, if a code tree T_C has any internal node of degree 1, we can clearly “prune” that node to obtain a more efficient code. When T_C has no nodes to prune, we obtain a full binary tree. *Henceforth, we assume that $|\Sigma| \geq 2$ and all code trees are full binary trees.* The case $|\Sigma| = 1$ is clearly very special, and we may ignore it. Returning to the informal problem (P), we can now interpret this problem as the construction of the best prefix-free code C for s , *i.e.*, the code that minimizes the length $|C(s)|$ of $C(s)$. Clearly, the only statistics important about s is the frequency $f_s(x)$ of each letter x in s , *i.e.*, the number of occurrences of x in s . In general, call a function⁹ of the form

$$f : \Sigma \rightarrow \mathbb{N} \quad (23)$$

⁹Sometimes, frequency is regarded as a fraction between zero and one. But we view it as an counting value: perhaps “census function” is a better term here.

a **frequency function**. So we now regard the input data to our problem to be a frequency function $f = f_s$ rather than the string s . Relative to f , the **cost** of C is defined to be

$$COST(f, C) := \sum_{a \in \Sigma} |C(a)| \cdot f(a). \quad (24)$$

Clearly $COST(f_s, C)$ is the length of $C(s)$. Finally, the **cost** of f is defined by minimization over all choices of C :

$$COST(f) := \min_C COST(f, C)$$

over all prefix-free codes C on the alphabet Σ . A code C is **optimal** for f if $COST(f, C)$ attains this minimum. It is easy to see that an optimal code tree must be a *full* binary tree (i.e., non-leaves must have two children).

Consider a simple example¹⁰ where s is “abadacadaba”. So $\Sigma = \{a, b, c, d\}$ with frequencies of the characters a, b, c, d equal to 6, 2, 1, 2 (respectively). For the codes in Figure 2, the cost of the first code is $6 \cdot 2 + 2 \cdot 3 + 1 \cdot 3 + 2 \cdot 1 = 23$. The second code is better, with cost 19.

¶20. **Huffman Code Algorithm.** We can now restate the informal problem (P) as the precise **Huffman coding problem**:

(H) Given a frequency function $f : \Sigma \rightarrow \mathbb{N}$, find an optimal prefix-free code C for f .

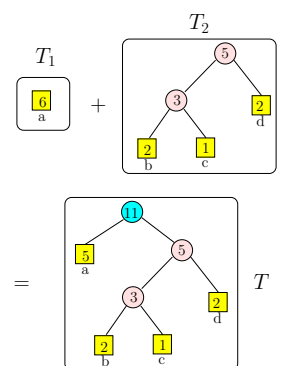
Relative to a frequency function f on Σ , we associate a **weight** $W(u)$ with each node u of the code tree T_C : the weight of a leaf is just the frequency $f(x)$ of the character x at that leaf, and the weight of an internal node is the sum of the weights of its children. Let $T_{f,C}$ denote such a **weighted code tree**. In general, a weighted code tree is just a code tree together with weights on each node satisfying the property that the weight of an internal node is the sum of the weights of its children. For example, see Figure 2 where the weight of each node is written next to it. The **weight** of $T_{f,C}$ is the weight of its root, and its **cost** $COST(T_{f,C})$ defined as the sum of the weights of all its *internal* nodes. In Figure 2(a), the internal nodes have weights 3, 9, 11 and so the $COST(T_{f,C}) = 3 + 9 + 11 = 23$. In general, the reader may verify that

$$COST(f, C) = COST(T_{f,C}). \quad (25)$$

We need the **merge** operation on code trees: if T_i is a code tree on the alphabet Σ_i ($i = 1, 2$) and $\Sigma_1 \cap \Sigma_2$ is empty, then we can merge them into a code tree T on the alphabet $\Sigma_1 \cup \Sigma_2$ by introducing a new node as the root of T and with T_1, T_2 as the two children of the root. We may write $T = T_1 + T_2$. Note that we do not care whether T_1 or T_2 is the left of T . If T_1, T_2 are weighted code trees, the result T is also a weighted code tree.

We now present a greedy algorithm for the Huffman coding problem:

¹⁰Regard it as toddler’s version of the magical incantation “abracadabra” or its more palindromic form “abradacadabra”. The palindromic nature on this incantation is supposed to be part of the magic.



tree merging

HUFFMAN CODE ALGORITHM:

Input: Frequency function $f : \Sigma \rightarrow \mathbb{N}$.Output: Optimal code tree T^* for f .

1. Let Q be a set of weighted code trees. Initially, Q is the set of $n = |\Sigma|$ trivial trees, each tree with only one node representing a single character in Σ .
2. While Q has more than one tree,
 - 2.1. Choose $T, T' \in Q$ with the minimum and next-to-minimum weights, respectively.
 - 2.2. Merge T, T' and insert the result $T + T'$ into Q .
 - 2.3. Delete T, T' from Q .
3. Now Q has only one tree T^* . Output T^* .

A **Huffman tree** is defined as a weighted code tree that *could* be output by this algorithm. We say “could” because the Huffman code algorithm is nondeterministic – when two trees have the same weight, the algorithm may pick either one. For instance, Figure 2(a) is non-Huffman, but Figure 2(b) is Huffman.

Let us illustrate the algorithm with a familiar 12-letter string, **hello world!**. The alphabet Σ for this string and its frequency function may be represented by the following two arrays:

arguably the most famous string in computing

letter	h	e	l	o	␣	w	r	d	!
frequency	1	1	3	2	1	1	1	1	1

Note that the exclamation mark (!) and blank space (␣) are counted as letters in the alphabet Σ . The final Huffman tree is shown in Figure 3. The number shown inside a node u of the tree is the **weight** of the node. This is just sum of the frequencies of the leaves in the subtree at u . Each leaf of the Huffman tree is labeled with a letter from Σ .




Figure 3: Huffman Tree for “hello world!”: weights are inside each node, but ranks (0, 1, ..., 16) are beside the nodes.

Figure 3 shows the Huffman tree produced by our algorithm on our famous string. In addition, we display, next to each node, its “rank” (0, 1, 2, ..., 16). The rank of a node specifies the order in which nodes were extracted from the priority queue. For instance, the leaves **h** (rank 0) and **e** (rank 0) were the first two to be extracted in the queue. Their merge produced a node of rank 8. Note that the root is the last (rank 16) to be extracted from the queue.

With this rank information, we can re-trace the step-by-step execution of the Huffman code algorithm. In the next section, we will exploit rank information in a more significant way. For the time being, we draw Huffman trees without paying much attention to its “orientation” i.e., the ordering of the 2 children of an internal node. When we treat dynamic Huffman tree in the next section, we will start paying attention to this issue. E.g., in Figure 3, the left child of the root has rank 15. Thus the Huffman code for the letter ℓ is 00. But we could also let the node of rank 14 be the left child of the root (this will be required in dynamic Huffman trees). The Huffman code of ℓ is then 10.

¶21. Implementation and complexity. The input for the Huffman algorithm may be implemented as an array $f[1..n]$ where $f[i]$ is the frequency of the i th letter and $|\Sigma| = n$. We construct the Huffman code tree T whose leaves are weighted by frequencies $f[a]$ ($a \in \Sigma$), and internal nodes have weights that are the sum of the weights of the children. This algorithm can be implemented using a priority queue Q containing a set of binary tree nodes. Recall (§III.2) that a priority queue supports two operations, (a) inserting a keyed item and (b) deleting the item with smallest key. The frequency of the code tree serves as its key. Any balanced binary tree scheme (such as AVL trees in Lecture IV) will give an implementation in which each queue operation takes $O(\log n)$ time. Hence the overall algorithm takes $O(n \log n)$.

Now that we have constructed the code tree, let us see how to use it for coding and decoding. Decoding is easy: just start from the root of T , we follow a path to a leaf by turning left or right according to the scanned bits. Once we read a leaf, we can output the character stored at the leaf, and start again at the root. What about using T for encoding a string $s \in \Sigma^*$? For each $a \in \Sigma$, we need to determine the code word $C(a) \in \{0, 1\}^*$. How could we do it? We must assume a separate “character map” \mathbf{Cm} that takes each $a \in \Sigma$ to the corresponding leaf in T . Then by following parent pointers from $\mathbf{Cm}(a)$ to the root, we obtain the code word (in reverse).

¶22. Correctness. We show that the produced code C has minimum cost. This depends on the following simple lemma. Let us say that a pair of nodes in T_C is a **deepest pair** if they are siblings and their depth is equal to the depth of T_C . In a full binary tree, there is always a deepest pair.

Lemma 9 (Deepest Pair Property) *For any frequency function f , there exists a code tree T that is optimal for f , with the further property that T has a deepest pair b, c where b is some least frequent character, and c is some next-to-least frequent character.*

Proof. The proof follows from a simple observation: suppose a and b are two leaves in T such that the depth of a is at most the depth of b : $D(a) \leq D(b)$. If their frequencies satisfy $f(a) \leq f(b)$, then we can exchange $a \leftrightarrow b$ in the tree without increasing its cost, $COST(f, T)$. If T is optimal for f , the exchange will preserve optimality. By making two such exchanges in an initially optimal T , we produce an optimal tree with the desired “deepest pair” property. **Q.E.D.**

Note that this lemma does not claim that every optimal code tree has the deepest pair property. See Exercise for a counter example.

We are ready to prove the correctness of Huffman’s algorithm. Suppose by induction hypothesis that our algorithm produces an optimal code whenever the alphabet size $|\Sigma|$ is less

than n . The basis case, $n = 1$, is trivial. Now suppose $|\Sigma| = n > 1$. After the first step of the algorithm in which we merge the two least frequent characters b, b' , we can regard the algorithm as constructing a code for a modified alphabet Σ' in which b, b' are replaced by a new character $[bb']$ with modified frequency f' such that $f'([bb']) = f(b) + f(b')$, and $f'(x) = f(x)$ otherwise. By induction hypothesis, the algorithm produces the optimal code C' for f' :

$$\text{COST}(f') = \text{COST}(f', C'). \quad (26)$$

From the code C' , we can obtain a code C for Σ as follows: the code tree T_C is obtained from $T_{C'}$ by replacing the leaf $[bb']$ by an internal node with two children representing b and b' , respectively. Thus we have

$$\text{COST}(f, C) = \text{COST}(f', C') + f(b) + f(b'). \quad (27)$$

By our deepest pair lemma, and using the fact that the COST is a sum over the weights of internal nodes, we conclude that

$$\text{COST}(f) = \text{COST}(f') + f(b) + f(b'). \quad (28)$$

[More explicitly, this equation says that if T is the optimal weighted code tree for f and T has the deepest pair property, then by removing the deepest pair with weights $f(b)$ and $f(b')$, we get an optimal weighted code tree for f' .] From equations (26)–(27) and (28), we conclude $\text{COST}(f) = \text{COST}(f, C)$, i.e., C is optimal. ■

¶23. Representation of the Code Tree for Transmission. Suppose I want to send you the string s . It is not enough to send you its code $C(s)$. I must also send you some representation of the code tree T_C . Let this representation be a binary string α_C . We will now provide one description of α_C that is essentially optimal in its length. We split the binary α_C into two parts:

$$\alpha_C = \beta_C \gamma_C \quad (29)$$

where β_C encodes the "shape" of the code tree T_C , and γ_C is just a list of the characters in Σ , in the order they appear as leaves of T_C .

We first explain γ_C . Let the code be $C : \Sigma \rightarrow \{0, 1\}^*$. The alphabet Σ occurs in some larger context. This context is provided by a **standard character set** which we denote by Σ_0 . For instance, Σ_0 may be the ASCII set or Unicode (see Figure 8 and notes below). We assume that Σ is a subset of Σ_0 . For the present purposes, assume $\Sigma_0 \subseteq \{0, 1\}^N$ for some fixed N . This N is the common knowledge used by both the transmitter and receiver. As a practical matter, it is important that we allow Σ to be a proper subset of Σ_0 . Typically, Σ is just the set of characters that actually occur in the string we are transmitting.

E.g., Let C be Huffman code in Figure 2(b), Σ_0 be the extended ASCII set (so $N = 8$). If $\Sigma = \{a, b, c, d\}$, then the codes for these letters in hexadecimal are $a=0x61, b=0x62, c=0x63, d=0x64$ (see Figure 8 in the next section). Therefore $\gamma_C = 0x61626364$. In full glory, $\gamma_C = 0110'0001'0110'0010'0110'0011'0110'0100$.

It remains to explain the string β_C in (29). We give a progression of ideas that lead to the final form. The initial idea is simple: let us prescribe a systematic way to traverse T . Starting from the root, we use a depth-first traversal, always go down the left child first. Each edge is traversed twice, initially downward and later upward. Then if we "spit" out a 0 for going down an edge and "spit" out a 1 for going up an edge, we would have faithfully output a description of the shape of T by the time we return to the root for the second time. Figure 4 illustrates this traversal of the Huffman tree of Figure 2(a), and shows resulting binary sequence

$$0010'0101'1101. \quad (30)$$

$4n - 4 = 12 \text{ bits!}$

Where have we
exploited this fact?




Figure 4: Compressed bit representation for the Huffman tree Figure 2a

Another remark is this: this encoding is **self-limiting** in the sense that, if we know the beginning of the encoding, then we would know when we have reached the end of the encoding. This property is useful for many applications; in particular, we need this property in (29): we must know when β_C ends, and γ_C begins.

To improve this representation, observe that a contiguous sequence of ones can be replaced by a single 1 since we know where to stop when going upward from a leaf (we stop at the first node whose right child has not been visited). This also takes advantage of the fact that we have a full binary tree. Previously we used $2n - 2$ ones. With this improvement, we only use n ones (corresponding to the leaves). The representation now has only $3n - 2$ bits. Then (30) is now represented by

$$0010'0101'01. \quad (31)$$

$3n - 2 = 10 \text{ bits!}$

Finally, we note that each 1 is immediately followed by a 0 (since the 1 always leads us to a node whose right child has not been visited, and we must immediately go down to that child). The only exception to this rule is the final 1 when we return to the root; this final 1 is not followed by a 0. We propose to replace all such 10 sequences by a plain 1. Since there are n ones (corresponding to the n leaves), we would have eliminated $n - 1$ zeros in this way. This gives us the final representation with $2n - 1$ bits. The scheme (31) is now shortened to:

$$0010'111. \quad (32)$$

$2n - 1 = 7 \text{ bits!}$

See the final illustration in Figure 4. The final scheme (32) will be known as the **compressed bit representation** β_T of a full binary tree T . In case T is the code tree for a Huffman code C , β_T is the desired β_C of (29).

Recall that $|\Sigma| \geq 2$ and hence T has more than one node: in this case, β_T always begins with a 0 and ends with a 1. So the shortest such string is 011, represent a full binary tree with two leaves. If T has only one node, it is natural to represent it as $\beta_T = 1$. Some additional simple properties are summarized as follows:

Lemma 10 Let T be a full binary tree T with $n \geq 1$ leaves, and β_T is its compressed bit representation.

(i) The length of β_T is $2n - 1$, with $n - 1$ zeros and n ones.

- (ii) The number of zeros is at least the number of ones in any proper prefix of β_T .
 (iii) The set $S \subseteq \{0, 1\}^*$ of all such compressed bit representations β_T forms a prefix-free set.
 (iv) There is a linear time algorithm that checks whether a given binary string s belongs to the set S , i.e., if s has the form β_T for some T .

We leave the proof as an Exercise. This has the following application:

Theorem 11 *There is a protocol to transmit a binary string α_C representing any Huffman code $C : \Sigma \rightarrow \{0, 1\}^*$ on $|\Sigma| = n$ letters such that*

- (i) *The length of α_C is $(2n - 1) + Nn = n(N + 2) - 1$.*
 (ii) *A receiver can recover the code C from α_C in linear time, without prior knowledge of Σ except that $\Sigma \subseteq \Sigma_0 \subseteq \{0, 1\}^N$.*

Proof. Using the notation of (29), $\alpha_C = \beta_C \gamma_C$ where $|\beta_C| = 2n - 1$ when the code tree T_C has n leaves, and $|\gamma_C| = nN$ under our assumption that $\Sigma \subseteq \Sigma_0 \subseteq \{0, 1\}^N$. This proves (i). For part (ii), the receiver can use the prefix-free property of α_C to detect the end of α_C while processing β_C . In linear time, it reconstruct the shape of T_C and thus knows n . Since the receiver knows N , it can also parse each symbol of Σ in the rest of β_C . **Q.E.D.**

Canonical Huffman Trees. A full binary tree is said to be **canonical** if in each level, all the leaves appear to the left of all the internal nodes. If a full binary tree is non-canonical, we can try to make it canonical by **swapping**: pick any two nodes u, v at the same level, and swap them. The subtrees T_u, T_v rooted at u and v are implicitly swapped by this operation (so T_u, T_v is unchanged). By repeated swaps, we can obtain a canonical tree. For example, the tree in Figure 2(a) is non-canonical: at level 1, the leaf **d** appears after the internal node **9**. After swapping them, the tree become canonical. Likewise, the tree in Figure 2(b) can be made canonical with one swap. The tree in Figure 3 is a canonical (despite the way we draw all leaves at the “same” level).

Canonical trees has a self-limiting encoding that uses only $n+2$ bits where n is the number of leaves (Exercise). We want to exploit this in the preceding lemma on α_C . Suppose T is Huffman tree. We note that swaps does not change the cost of the code tree (it is less clear whether swaps preserves Huffman-ness of trees). By repeated swaps, we finally reach a “canonical Huffman tree”.

Remarks: The publication of the Huffman algorithm in 1952 by D. A. Huffman was considered a major achievement. This algorithm is clearly useful for compressing binary files. See “Conditions for optimality of the Huffman Algorithm”, D.S. Parker (*SIAM J. Comp.*, 9:3(1980)470–489, *Erratum* 27:1(1998)317), for a variant notion of cost of a Huffman tree and characterizations of the cost functions for which the Huffman algorithm remains valid.

¶24. Notes on Morse Code. In the Morse¹¹ code, letters are represented by a sequence of dots and dashes: $a = \cdot -$, $b = - \cdot \cdot \cdot$ and $z = - - \cdot \cdot$. The code is also meant to be sounded: dot is pronounced ‘dit’ (or ‘di-’ when non-terminal), dash is pronounced ‘dah’ (or ‘da-’ when non-terminal). So the famous distress signal “S.O.S” is di-di-di-da-da-da-di-di-dit. Thus ‘a’ is di – dah, ‘z’ is da – da – di – dit. The code does not use capital or small letters. Here is the full alphabet:

¹¹Samuel Finley Breese Morse (1791-1872) was Professor of the Literature of the Arts of Design in the University of the City of New York (now New York University) 1832-72. It was in the university building on Washington Square where he completed his experiments on the telegraph.

Letter	Code	Letter	Code
A	. -	B	- . . .
C	- . . .	D	- . .
E	.	F	. . - .
G	- - .	H
I	. .	J	. - - -
K	- . -	L	. - . .
M	- -	N	- .
O	- - -	P	. - - .
Q	- - . -	R	. - .
S	. . .	T	-
U	. . -	V	. . . -
W	. - -	X	- . . -
Y	- . - -	Z	- - . .
0	- - - - -	1	. - - - -
2	. . - - -	3	. . . - -
4	. . . -	5
6	-	7	- - . . .
8	- -	9	- - - . . .
Fullstop (.)	. - . - . -	Comma (,)	- - . . - -
Query (?)	. . - - . .	Slash (/)	-
BT (pause)	- . . . -	AR (end message)	. - . . .
SK (end contact)	. . . - . -		

Table 1: Morse Code

Note that Morse code assigns a dot to **e** and a dash to **t**, the two most frequent English letters. These two assignments dash any hope for a prefix-free code. So how can do you send or decode messages in Morse code? Spaces! Since spaces are not part of the Morse alphabet, they have an informal status as an explicit character (so Morse code is not strictly a binary code). There are 3 kinds of spaces: space between *dit*'s and *dah*'s within a letter, space between letters, and space between words. Let us assume some **unit space**. Then the above three types of spaces are worth 1, 3 and 7 units, respectively. These units can also be interpreted as “unit time” when the code is sounded. Hence we simply say **unit** without prejudice. Next, the system of dots and dashes can also be brought into this system. We say that spaces are just “empty units”, while *dit*'s and *dah*'s are “filled units”. *dit* is one filled unit, and *dah* is 3 filled units. Of course, this brings in the question: why 3 and 7 instead of 2 and 4 in the above? Today, Morse code is still required of HAM radio operators and is useful in emergencies.

EXERCISES

Exercise 4.1: Give a Huffman code for the string “hello! this is my little world!”.

◇

Exercise 4.2: What is the length of the Huffman code for the string $s =$ “please compress me”. Show your hand computation. Do not forget the empty space character.

◇

Exercise 4.3: Consider the following letter frequencies:

$$a = 5, b = 1, c = 3, d = 3, e = 7, f = 0, g = 2, h = 1, i = 5, j = 0, k = 1, l = 2, m = 0, \\ n = 5, o = 3, p = 0, q = 0, r = 6, s = 3, t = 4, u = 1, v = 0, w = 0, x = 0, y = 1, z = 1.$$

Please determine the cost of the optimal tree. NOTE: you may ignore letters with the zero frequency.

◇

Exercise 4.4: Give an example of a prefix-free code $C : \Sigma \rightarrow \{0, 1\}^*$ and a frequency function $f : \Sigma \rightarrow \mathbb{N}$ with the property that (i) $COST(C, f)$ is optimal, but (ii) C could not have arisen from the Huffman algorithm. Try to minimize $|\Sigma|$.

◇

Exercise 4.5: True or False? If T and T' are two optimal prefix-free code for the frequency function $f : \Sigma \rightarrow \mathbb{N}$, then T and T' are isomorphic as unordered trees. Prove or show counter example. NOTE: a binary tree is an ordered tree because the two children of a node are ordered. \diamond

Exercise 4.6: In the text, we prove that for any frequency function f , there is an optimal code tree in which there is a deepest pair of leaves whose frequencies are the least frequent and the next-to-least frequent. Consider this stronger statement: *if T is any optimal code tree for f , there must be a deepest pair whose frequencies are least frequent and next-to-least frequent.* Prove it or show a counter example. \diamond

Exercise 4.7: Let $C : \Sigma \rightarrow \{0, 1\}^*$ be any prefix-free code whose code tree T_C is a full-binary tree. Prove that there exists a frequency function $f : \Sigma \rightarrow \mathbb{N}$ such that C is optimal. \diamond

Exercise 4.8:

(a) Draw the full binary tree corresponding to its compressed bit representations:

$$\alpha_1 = 0010'1100'1011'1 \quad \alpha_2 = 0100'1001'0011'111$$

- (b) What is β_T where T is the full binary tree with 6 leaves and every right child is a leaf.
 (c) What is β_T where T is the full binary tree with 6 leaves and every left child is a leaf.
 (d) What is β_T where T is the complete binary tree with 8 leaves. \diamond

Exercise 4.9: Joe Smart suggested that we can slightly improve the compressed bit representation of full binary trees on n leaves as follows: since the first bit is always 0 and the last bit is always 1, we can use only $2n - 3$ bits instead of $2n - 1$. What are some issues that might arise with this improvement? \diamond

Exercise 4.10: The text gave a method to represent any full binary tree T on n leaves using a binary string β_T with $2n - 1$ bits. Clearly, not every binary string of length $2n - 1$ represents a full binary tree. For instance, the first and last bits must be 0 and 1, respectively. Give a necessary and sufficient condition for a binary string to be a valid representation. \diamond

Exercise 4.11: For any binary full tree T , we have given two representations: the array A_T and the bit string β_T . Give detailed algorithms for the following conversion problems:
 (a) To construct the string β_T from the array A_T .
 (b) To construct the array A_T from the string β_T . \diamond

Exercise 4.12: Let T be a full binary tree on n leaves. Give an algorithm to convert its compressed bit representation $\beta_T[1..2n - 1]$ to a $4n - 4$ array $B[1..4n - 4]$ representing the traversal of T . \diamond

Exercise 4.13: Suppose we want to represent an arbitrary binary tree, not necessarily full. HINT: there is a bijection between arbitrary binary trees and full binary trees. Exploit our compressed bit-representation of full binary trees. \diamond

Exercise 4.14: (a) Prove (25).

(b) It is important to note that we defined $COST(T_{f,C})$ to be the sum of $f(u)$ where u range over the *internal* nodes of $T_{f,C}$. That means that if $|\Sigma| = 1$ (or $T_{f,C}$ has only one node which is also the root) then $COST(T_{f,C}) = 0$. Why does Huffman code theory break down at this point?

(c) Suppose we (accidentally) defined $COST(T_{f,C})$ to be the sum of $f(u)$ where u range over the *all* nodes of $T_{f,C}$. Where in your proof in (a) would the argument fail? \diamond

Exercise 4.15: (Kraft Inequality) Let $0 \leq d_1 \leq d_2 \leq \dots \leq d_n$ be the depths of the leaves in a binary tree with n leaves. Prove that

$$1 \geq \sum_{i=1}^n 2^{-d_i}. \quad (33)$$

Moreover, if the binary tree is a full binary tree, then the inequality is an equality. \diamond

Exercise 4.16: (Elias) Let $bin(n)$ denote the standard binary encoding of $n \in \mathbb{N}$ and $len(n)$ be the length of this encoding. E.g., $bin(0, 1, 2, 3, 4) = (\epsilon, 1, 10, 11, 100)$ and $len(0, 1, 2, 3, 4) = (0, 1, 2, 2, 3)$. Here, ϵ is the empty string and we use the compact notation $bin(n_1, n_2, \dots, n_k) := (bin(n_1), bin(n_2), \dots, bin(n_k))$, etc.

We want a binary encoding of natural numbers, $rep : \mathbb{N} \rightarrow \{0, 1\}^*$, with the following property: if n_1, n_2, \dots is any sequence of natural numbers, the binary string $rep(n_1)rep(n_2)\dots$ is uniquely decodable. Such an encoding $rep(n)$ is **self-limiting** in the sense that whenever we know the start of $rep(n)$, we can also determine where it ends. Alternatively, the representation is prefix-free: if $n \neq n'$ then $rep(n)$ is not a prefix of $rep(n')$.

(a) Consider the following encoding scheme $rep_1 : \mathbb{N} \rightarrow \{0, 1\}^*$ for natural numbers: $rep_1(0) = 1$ and for $n \geq 1$, $rep_1(n) = 0^{len(n)}bin(n)$. E.g., $rep_1(0, 1, 2, 3, 4) = (1, 0'1, 00'10, 00'11, 000'100)$. Note we use the prime mark ($'$) for decoration. What is $rep_1(99)$? What is the length of $rep_1(n)$ as a function of n ?

(b) Now consider $rep_2 : \mathbb{N} \rightarrow \{0, 1\}^*$ where $rep_2(n) = rep_1(len(n))bin(n)$ for $n \geq 0$. E.g., $rep_2(0, 1, 2, 3, 4) = (1, 01'1, 0010'10, 0010'11, 0011'100)$. What is $rep_2(99)$? What is the length of $rep_2(n)$ as a function of n ?

(c) What is wrong with the suggestion to define $rep(n)$ recursively as follows:

$$rep(n) = rep(len(n))bin(n)$$

for all $n \geq n_0$ (for some n_0)? How can you fix this issue? How small can n_0 be, and what can you do for $n < n_0$? What is the length of your representation as a function of n ? What is your representation of 99? For what values of n will your representation be shorter than $rep_2(n)$? \diamond

Exercise 4.17: Here is an idea to improve the compressed bit representation of Huffman code trees given in our text. Say a full binary tree T is **canonical** if at every level, every leaf must appear before any internal node in the left-right listing of the level. A Huffman code tree is canonical if its underlying tree is canonical.

(a) Show that every full binary tree can transformed into a canonical one by a sequence of swaps between pairs of nodes at the same level. NOTE: when we swap two nodes, we are really swapping the entire subtrees rooted at these two nodes.

- (b) Show that swaps do not change the cost of a Huffman tree.
 (c) Give an **improved compressed bit representation** for canonical Huffman trees.
 (d) What is the range of sizes for your improved compressed bit representation for a canonical tree with n leaves? \diamond

Exercise 4.18: (Gashlin 2012) The previous exercise gave an “improved compressed bit representation” for canonical Huffman trees. We now give a further improvement, but using a radically different approach: instead of focusing on the leaves, we focus on the internal nodes. An integer sequence of the form

$$\mathbf{n} = (n_0, n_1, \dots, n_{d-1}) \quad (34)$$

is called the **profile** of a full binary tree of depth d and at level $i = 0, \dots, d-1$, it has n_i internal nodes.

- (a) Show that there is a bijection between canonical trees with N internal nodes and integer sequences of the form (34) that satisfies $N = \sum_{i=0}^{d-1} n_i$, $n_0 = 1$ and $1 \leq n_i \leq 2n_{i-1}$ ($i = 1, \dots, d-1$).
 (b) Let the **leaf profile** of a full binary tree of depth d be the sequence $(\ell_1, \ell_2, \dots, \ell_d)$ where ℓ_i is the number of leaves at level $i = 1, \dots, d$. Given the profile (34) of a full binary tree, what is the corresponding leaf profile?
 (c) Give a self-limiting encoding of profiles of canonical trees. HINT: use simple self-limiting encodings, but be sure we can detect the end of the encoding.
 (d) Give an upper bound on the number $T(N)$ of canonical Huffman trees with N internal nodes.
 (e) Give a lower bound on $T(N)$. \diamond

Exercise 4.19: Below is President Lincoln’s address at Gettysburg, Pennsylvania on November 19, 1863.

- (a) Give the Huffman code for the string S comprising the first two sentences of the address. Also state the length of the Huffman code for S , and the percentage of compression so obtained (assume that the original string uses 7 bits per character). View caps and small letters as distinct letters, and introduce symbols for space and punctuation marks. But ignore the newline characters.
 (b) The previous part was meant to be done by hand. Now write a program in your favorite programming language to compute the Huffman code for the entire Gettysburg address. What is the compression obtained?

Four score and seven years ago our fathers brought forth on this continent a new nation, conceived in liberty and dedicated to the proposition that all men are created equal. Now we are engaged in a great civil war, testing whether that nation or any nation so conceived and so dedicated can long endure. We are met on a great battlefield of that war. We have come to dedicate a portion of that field as a final resting-place for those who here gave their lives that that nation might live. It is altogether fitting and proper that we should do this. But in a larger sense, we cannot dedicate, we cannot consecrate, we cannot hallow this ground. The brave men, living and dead who struggled here have consecrated it far above our poor power to add or detract. The world will little note nor long remember what we say here, but it can never forget what they did here. It is for us the living rather to be dedicated here to the unfinished work which they who fought here have thus far so nobly advanced. It is rather for us to be here dedicated to the great task remaining before us -- that from these honored dead we take increased devotion to that cause for which they gave the last full measure of devotion -- that we here highly resolve that these dead

shall not have died in vain, that this nation under God shall have a new birth of freedom, and that government of the people, by the people, for the people shall not perish from the earth.

◇

Exercise 4.20: Let (f_0, f_1, \dots, f_n) be the frequencies of $n+1$ symbols (assuming $|\Sigma| = n+1$). Consider the Huffman code in which the symbol with frequency f_i is represented by the i th code word in the following sequence

$$1, 01, 001, 0001, \dots, \underbrace{00 \cdots 01}_{n-1}, \underbrace{00 \cdots 001}_n, \underbrace{00 \cdots 000}_n.$$

(a) Show that a sufficient condition for optimality of this code is

$$\begin{aligned} f_0 &\geq f_1 + f_2 + f_3 + \cdots + f_n, \\ f_1 &\geq f_2 + f_3 + \cdots + f_n, \\ f_2 &\geq f_3 + \cdots + f_n, \\ &\dots \\ f_{n-2} &\geq f_{n-1} + f_n. \end{aligned}$$

(b) Suppose the frequencies are distinct. Give a set of sufficient and necessary conditions.

◇

Exercise 4.21: Suppose you are given the frequencies f_i in sorted order. Show that you can construct the Huffman tree in linear time.

◇

Exercise 4.22: (Representation of Binary Trees) In the text, we showed that a full binary tree on n leaves can be represented using $2n - 1$ bits. Suppose T is an arbitrary binary tree, not necessarily full. With how many bits can you represent T ? HINT: by extending T into a full binary tree T' , then we could use the previous encoding on T' .

◇

Exercise 4.23: Huffman code is based on transmitting bits. Suppose we transmit in ‘trits’ (a base-3 digit). Then the corresponding 3-ary Huffman code $C : \Sigma \rightarrow \{0, 1, 2\}^*$ is represented by a 3-ary code tree T where each leaf is associated with a unique letter in Σ and each internal node has degree at most 3. If $f : \Sigma \rightarrow \mathbb{N}$ is a frequency function, this assigns a weight to each node of T : the leaf associated with $x \in \Sigma$ has weight $f(x)$, and each internal node has a weight equal to the sum of the weights of its children. The cost of T is defined as usual, as the sum of the weights of the internal nodes of T . We are interested in optimal trees T , i.e., whose cost is minimum.

(a) Show that in an optimal 3-ary code tree, there are no nodes of degree 1 and at most one node of degree 2. Furthermore, if a node has degree 2, it must have leaves as both of its children.

(b) Let T be a tree whose internal nodes have degrees 2 or 3. If there are d_i nodes of degree i ($i = 0, 2, 3$) in T show that $d_0 = 1 + d_2 + 2d_3$.

(c) Show that there are optimal 3-ary code trees with this property: if $|\Sigma|$ is odd, there are no degree 2 nodes, and if $|\Sigma|$ is even, there is one degree 2 node. Moreover, if the unique node u of degree 2 we may assume its children have minimum frequencies among all the leaves.

(d) Give an algorithm for constructing an optimal 3-ary code tree and prove its correctness.

◇

Exercise 4.24: Refer to the previous question on optimal Huffman ternary code (i.e., using trits instead of bits). Prove the correctness of the following algorithm for computing the optimal ternary Huffman code algorithm:

TERNARY HUFFMAN CODE ALGORITHM:

Input: Frequency function $f : \Sigma \rightarrow \mathbb{N}$.

Output: Optimal ternary code tree T^* for f .

1. Let Q be a priority queue containing weighted code trees.
Priority is determined by the weight of the root.
Initially, Q is the set of $n = |\Sigma|$ trivial trees,
each tree with one node representing a single character in Σ .
2. If n is even,
 $T \leftarrow Q.\text{deleteMin}()$, $T' \leftarrow Q.\text{deleteMin}()$.
 $Q.\text{enqueue}(\text{Merge}(T, T'))$
3. While Q has more than one tree,
3.1. $T \leftarrow Q.\text{deleteMin}()$, $T' \leftarrow Q.\text{deleteMin}()$, $T'' \leftarrow Q.\text{deleteMin}()$
3.2. $Q.\text{enqueue}(\text{Merge}(T, T', T''))$.
4. Now Q has only one tree T^* . Output T^* .

◇

Exercise 4.25: We want to compare the relative efficiency of bits versus trits in Huffman coding. Let $f : \Sigma \rightarrow \mathbb{N}$ be a frequency function. Suppose $H_2(f)$ is cost of f under the standard (binary) Huffman code; let $H_3(f)$ be the cost using a ternary Huffman code. We say “bits are better than trits” on f if $H_2(f) < H_3(f) \lg(3)$ (and conversely if the inequality goes the other way). Using the algorithm in the previous question to compute $H_3(f)$, answer the question whether “bits are better than trits” for the following f ’s:

- (a) Let $f(a) = f(b) = f(c) = 1$ (and f is zero on other characters).
- (b) Let f be the frequency function for `compress this please`.

◇

Exercise 4.26: We consider the 4-ary version of the previous question. Let T be an optimum 4-ary code tree for some frequency function $f : \Sigma \rightarrow \mathbb{N}$.

- (a) Give a short inductive proof of the following fact: Suppose T is *any* 4-ary tree on $n \geq 1$ leaves, and let N_d be the number of nodes with d children ($d = 0, 1, 2, 3, 4$). Thus, $n = N_0$. Give a short inductive proof for the following formula: $n = 1 + N_2 + 2N_3 + 3N_4$.
- (b) Show that if T is an optimal code tree, then $N_1 = 0$ and $3N_2 + 2N_3 \leq 4$, and every non-full internal node has only leaves as children and the depth of these leaves must equal the height of T .
- (c) Moreover, we can always transform T from part (b) into T' such that the corresponding degrees satisfy $N'_1 = 0$ and $N'_2 + N'_3 \leq 1$. Also, for any non-full internal node of T' , its children have weights no larger than any other leaves.
- (d) Suppose $r = (n - 1) \bmod 3$. So $r \in \{0, 1, 2\}$. Show how N'_2, N'_3 in part (b) is determined by r .
- (e) Describe an algorithm to construct an optimal code tree from a frequency function f .
- (f) Show the optimal 4-ary Huffman tree for the input string `hello world!`. Please state the cost of this optimal tree.

◇

Exercise 4.27: Further generalize the 3-ary Huffman tree construction to arbitrary k -ary codes for $k \geq 4$.

◇

Exercise 4.28: Suppose that the cost of a binary code word w is $z + 2o$ where z (resp. o) is the number of zeros (resp. ones) in w . Call this the **skew cost**. So ones are twice as

expensive as zeros (this cost model might be realistic if a code word is converted into a sequence of dots and dashes as in Morse code). We extend this definition to the **skew cost** of a code C or of a code tree. A code or code tree is **skew Huffman** if it is optimum with respect to this skew cost. For example, see Figure 5 for a skew Huffman tree for alphabet $\{a, b, c\}$ and $f(a) = 3$, $f(b) = 1$ and $f(c) = 6$.




Figure 5: A skew Huffman tree with skew cost of 21.

- (a) Argue that in some sense, there is no greedy solution that makes its greedy decisions based on a linear ordering of the frequencies.
- (b) Consider the special case where all letters of the alphabet has equal frequencies. Describe the shape of such code trees. For any n , is the skew Huffman tree unique?
- (c) Give an algorithm for the special case considered in (b). Be sure to argue its correctness and analyze its complexity. HINT: use an “incremental algorithm” in which you extend the solution for n letters to one for $n + 1$ letters. \diamond

Exercise 4.29: (Golin-Rote) Further generalize the problem in the previous exercise. Fix $0 < \alpha < \beta$ and let the cost of a code word w be $\alpha \cdot z + \beta \cdot o$. Suppose α/β is a rational number. Show a dynamic programming method that takes $O(n^{\beta+2})$ time. NOTE: The best result currently known gets rid of the “+2” in the exponent, at the cost of two non-trivial ideas. \diamond

Exercise 4.30: (Open) Give a non-trivial algorithm for the problem in the previous exercise where α/β is not rational. An algorithm is “trivial” here if it essentially checks all binary trees with n leaves. \diamond

Exercise 4.31: The range of the frequency function f was assumed to be natural numbers. If the range is arbitrary integers, is the Huffman theory still meaningful? Is there fix? What if the range is the set of non-negative real numbers? \diamond

Exercise 4.32: (Shift Key in Huffman Code) We want to encode small as well as capital letters in our alphabet. Thus ‘a’ and ‘A’ are to be distinguished. There are two methods to do this. (I) View the small and capital letters as distinct symbols. (II) Introduce a special “shift” symbol, and each letter is assumed to be small unless it is preceded by a shift symbol, in which case the following letter is capitalized. As input string for this problem, use the text of this question. Punctuation marks are part of this string, but there is only one SPACE character. Newlines and tabs are regarded as instances of SPACE. Two or more consecutive SPACE characters are replace by a single SPACE.

- (a) What is the length of the Huffman code for our input string using method (I). Note that the input string begins with “We want to en...” and ends with “...ngle SPACE.”.
- (b) Same as part (a) but using method (II).
- (c) Discuss the pros and cons of (I) and (II).

(d) There are clearly many generalizations of shift keys, as seen in modern computer keyboards. The general problem arises when our letters or characters are no longer indivisible units, but exhibit structure (as in Chinese characters). Give a general formulation of such extensions. \diamond

END EXERCISES

§5. Dynamic Huffman Code

The original Huffman coding formulation does not take the full context of its applications, which we now clarify.

¶25. **The Larger Context of Coding.** Here is the typical sequence of steps for compressing and transmitting a string s using the Huffman code algorithm:

- (i) First make a pass over the string s to compute its frequency function.
- (ii) Next compute a Huffman code tree T_C corresponding to some code C .
- (iii) Using T_C , compute the compressed string $C(s)$.
- (iv) Finally, transmit the tree T_C (Theorem 11), together with the compressed string $C(s)$, to the receiver.

The receiver receives T_C and $C(s)$, and hence can recover the string s . Since the sender must process the string s in two passes (steps (i) and (iii)), the original Huffman tree algorithm is sometimes called the “2-pass Huffman encoding algorithm”. There are two deficiencies with this 2-pass process: (a) Multiple passes over the input string s makes the algorithm unsuitable for realtime data transmissions. Note that if s is a large file, this require extra buffer space. (b) The Huffman code tree must be explicitly transmitted before the decoding can begin. We need some way to encode T_C . This calls for a separate algorithm to handle T_C in the encoding and decoding process.


An approach called “Dynamic Huffman coding” (or adaptive Huffman coding) overcomes these problems: there is no need to explicitly transmit the code tree, and it passes over the string s only once. In fact, it does not even have to pass over the entire string even once, but can transmit as much of the string as has been read! This property is important for transmitting continuous stream of data that has no apparent end (e.g., ticker tape, satellite signals). Two known algorithms for dynamic Huffman coding [10] are the **FGK Algorithm** (Faller 1973, Gallager 1978, Knuth 1985) and the **Lambda Algorithm** (Vitter 1987). The dynamic Huffman code algorithm can be used for data compression: for example, it is used¹² in the Unix utility `compress/uncompress`.

¹²This particular utility has been replaced by better compression schemes.

¶26. **Sibling Property.** In Dynamic Huffman Coding, the weighted code tree T must evolve as characters from the input string is read. It must evolve in two ways: not only does the frequency of letters in Σ increase over time, but Σ itself can grow as new letters are encountered. We need to update our representation of T as this happens. The key idea is the “sibling property” of Gallagher.

Assume T has $k \geq 0$ internal nodes. So it has $k + 1$ leaves or $2k + 1$ nodes in all. We say T has the **sibling property** if its nodes can be **ranked** from 0 to $2k$ satisfying:

- (S1) (Weights are non-decreasing with rank) If w_i is the weight of node with rank i , then $w_{i-1} \leq w_i$ for $i = 1, \dots, 2k$.
- (S2) (Siblings have consecutive ranks) The nodes with ranks $2j$ and $2j + 1$ are siblings (for $j = 0, \dots, k - 1$).



For example, the weighted code tree in Figure 3 has been given the rankings $0, 1, 2, \dots, 16$. We check that this ranking satisfies the sibling property. Note that the node with rank $2k$ is necessarily the root, and it has no siblings. In general, let $r(u)$ denote the rank of node u . If the weights of nodes are all distinct, then the rank $r(u)$ is uniquely determined by Property (S1).

Lemma 12 *Let T be weighted code tree. Then T is Huffman iff it has the sibling property.*

Proof. If T is Huffman then by definition, it is constructed by the Huffman code algorithm. We can rank the nodes in the order that nodes are extracted from the priority queue, and this ordering implies the sibling property. Conversely, the sibling property of T determines an obvious order for merging pairs of nodes to form a Huffman tree. **Q.E.D.**

¶27. **Sibling Representation of Huffman Tree.** We provide an array representation of Huffman trees that exploits the sibling property. Let T be a Huffman tree with $k + 1 \geq 1$ leaves. Each of its $2k + 1$ nodes may be identified by its rank, *i.e.*, a number from 0 to $2k$. Hence node i has rank i . We use two arrays

$$\text{Wt}[0..2k], \quad \text{Lc}[0..2k]$$

of length $2k + 1$ where $\text{Wt}[i]$ is the *weight* of node i , and $\text{Lc}[i]$ is an even integer indicating the *left child* of node i . In case node i is a leaf, we may let $\text{Lc}[i] = -1$. Alternatively, as will be done in these notes, we let $\text{Lc}[i]$ store a character from the alphabet Σ ; of course, this convention assumes that one can distinguish between elements in the set $0, \dots, 2k$ and the characters in Σ . Thus, for a non-leaf i , its right child is given by $\text{Lc}[i] + 1$, and $\text{Lc}[i]$ is always an even integer. This, the left and right child of any node is a pair of the form $(2j, 2j + 1)$ (for some j). We ensure that the root is node $2k$.

We stress that storing elements of Σ in Lc is not essential, but it makes the examples more transparent. What *is* essential for our algorithms is the *inverse* representation that tells us, for each letter $x \in \Sigma$, which leaf in T contains x . Moreover, because of the dynamic nature of Σ , we need a more general mapping,

$$\text{Cm} : \Sigma_0 \rightarrow \{-1, 0, 1, 2, \dots, 2k\}$$


such that $\mathbf{Cm}[x] = i$ iff $Lc[i] = x \in \Sigma$, and $\mathbf{Cm}[x] = -1$ if $x \notin \Sigma$. Call \mathbf{Cm} the **character map array**. Initially, $\mathbf{Cm}[x] = -1$ for all $x \in \Sigma_0$ (i.e., initially $\Sigma = \emptyset$). As new letters in Σ_0 are encountered, they are added to Σ and the corresponding entry $\mathbf{Cm}[x]$ updated.

In summary, our Huffman tree is represented by three arrays $\mathbf{Lc}, \mathbf{Wt}, \mathbf{Cm}$. For example, the Huffman tree in Figure 3 is¹³ illustrated by the arrays in Table 2: Two of these arrays, \mathbf{Lc} and

Rank	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16
\mathbf{Lc}	h	e	□	w	r	d	!	o	0	2	4	6	ℓ	8	10	12	14
\mathbf{Wt}	1	1	1	1	1	1	1	2	2	2	2	3	3	4	5	7	12

Table 2: Compact representation of Huffman tree in Figure 3

\mathbf{Wt} , are explicitly shown. But the array $\mathbf{Cm}[x \in \Sigma_0]$ is easily inferred from the leaf entries of \mathbf{Lc} . E.g., $\mathbf{Cm}[h] = 0$, $\mathbf{Cm}[e] = 1$ and $\mathbf{Cm}[a] = -1$. There is no “Rank” array in this representation because, trivially, $\mathbf{Rank}[v] = v$ for all $v \in \{0, \dots, 2k\}$.



Here is a simple application of the Sibling representation. Suppose we are given a letter $x \in \Sigma$, and we want to determine the corresponding Huffman code $C(x)$. We need to first go to the leaf u of T corresponding to x . This is of course given by $u = \mathbf{Cm}[x]$. The last bit of $C(x)$ is therefore equal to the “parity” of u . (The parity of a natural number u is equal to 0 if u is even, and equal to 1 otherwise.) Then we replace u by $\mathit{parent}(u)$, and thereby determine the next bit of $C(x)$. Iterating this process, we stop when u eventually becomes the root. So the macro to compute the bits of $C(x)$ (in reverse order) is given by

```

C(x):
  u ← Cm[x]
  Output(parity(u))
  While u < 2k
    u ← parent(u)
    Output(parity(u))

```

The parent of u is computed by a simple for-loop:

```

parent(u):
  ℓ ← 2 ⌊u/2⌋
  for p ← u + 1 to 2k
    if (Lc[p] = ℓ), Return(p).

```

The for-loop is sure to terminate when u is non-root. Moreover, the number of times that the p variable is updated (over repeated calls to $\mathit{parent}(u)$ by $C(x)$) is at most $2k$ values since the value of p is strictly increasing with each assignment. This ensures the overall complexity of $C(x)$ is $O(k)$. Ideally, we would like to generate $C(x)$ in time $O(|C(x)|)$ (Exercise).

¹³If you compare Table 2 with the Huffman tree in Figure 3, you might be surprised that the left child of 16 is 14 and not 15. Until now, we have not taken the oriented-ness of Huffman trees seriously (since the length of the compressed string did not depend on the ordering of the 2 children of an internal node). According to the displayed binary tree in Figure 3, node 15 is the left child of the root. But the sibling property requires node 14 to be the left child.

¶28. **The Restoration Problem.** The key problem of dynamic Huffman tree is how to restore Huffman-ness under a particular kind of perturbation: let T be Huffman and suppose the weight of a leaf u is incremented by 1. So weight of each node along the path from u to the root must be similarly incremented. The result is a weighted code tree T' , but it may no longer be Huffman. *Informally, our problem is to restore Huffman-ness in such a tree T' .*

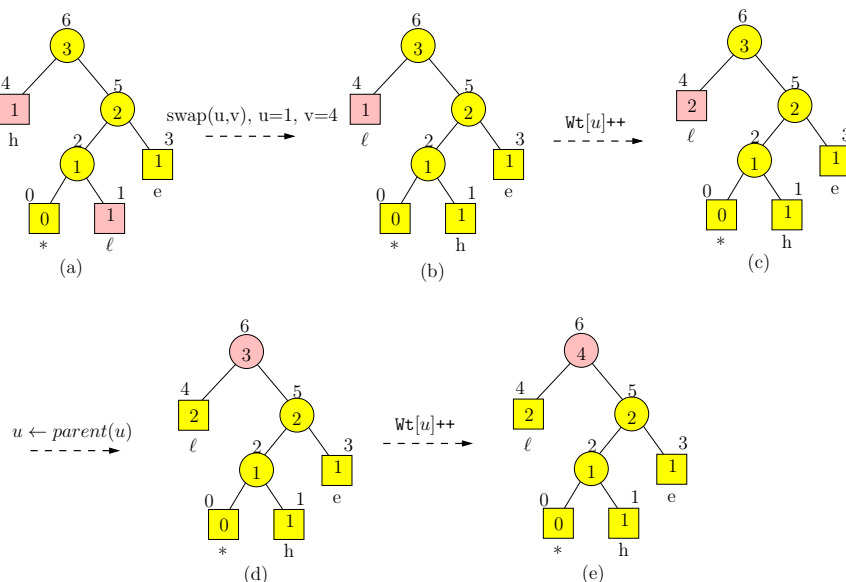


Figure 6: Restoring Huffmanness after incrementing the frequency of letter ℓ

Let us first give some intuition of what has to be done, using our example of **hello world!**. Begin with the Huffman tree after having transmitted the prefix **hel**. Assume that, somehow, we managed to construct a Huffman tree for this string as shown in Figure 6(a). The letters h , e and ℓ are stored in nodes 4, 3 and 1 (respectively). Note that there is a leaf with weight 0, but we ignore this for now. Each letter has frequency (=weight) of 1. The next transmitted letter is ℓ , and if we simply increase the frequency of node 1 (which represents ℓ) to $2 = 1 + 1$, we would violate the ranking property (S1) of ¶26. This is because the weight of a node of rank 1 would now be greater than the weights of nodes with greater rank (3 and 4). The key idea is to first *swap node 1 with node 4*. This is shown in Figure 6(b). Now, the letter ℓ is represented by node 4, and incrementing its weight by 1 is no longer a problem. The result is seen in Figure 6(c). We must next increment the weight of the parent of node 4, namely node 6. So the focus moves to node 6, as indicated by Figure 6(d). We can simply increment the weight of node 6 because it is the root. But if it is not the root, we may have to do a swap first. The result is Figure 6(e). The process stops since we have reached the root of the tree.

Consider the following algorithm for restoring Huffman-ness in T . For each node v in T , let $R(v)$ denote its rank in the original tree T . But our usual convention is that v is identified with its rank, i.e., $R(v) = v$. Let u be the current node. Initially, u is the leaf whose weight was incremented. We use the following iterative process:

RESTORE (u)

▷ u is a node whose weight is to be incremented

While (u is not the root) do

1. ▷ Find the node v with the largest rank $R(v)$ subject
 ▷ to the constraint $\text{Wt}[v] = \text{Wt}[u]$. Specifically:
 $v \leftarrow u$
 While ($\text{Wt}[v + 1] = \text{Wt}[u]$)
 $v++$
2. If ($v \neq u$)
3. Swap(u, v). ◁ This swaps the subtrees rooted at u and v .
4. $\text{Wt}[u]++$. ◁ Increment the weight of u
5. $u \leftarrow \text{parent}(u)$. ◁ Reset u
6. $\text{Wt}[u]++$. ◁ Now, u is the root

We need to explain one detail in the RESTORE routine. The swap operation in Line 3 needs to be explained: conceptually, swapping u and v means the subtree rooted at u and the subtree rooted at v exchange places. This can be confusing since our encoding identifies the nodes u and v with their rank. So for the moment, imagine that u is a node in a tree where nodes have parent, left child and right child pointers, etc. Suppose u' and v' were the parents (respectively) of u and v before the swap. Then after the swap, v' (resp., u') becomes the parent of u (resp., v). Coming back to our representation using the Lc array, we only have to exchange the values in the array entries Lc[u] and Lc[v]. But note that this swap may involve leaves, in which case we have to update the character map Cm:

SWAP(u, v)

(($(tmp \leftarrow \text{Lc}[u]) \leftarrow \text{Lc}[v]) \leftarrow tmp$)

If ($\text{Lc}[u] \in \Sigma_0$) then Cm[Lc[u]] $\leftarrow u$

If ($\text{Lc}[v] \in \Sigma_0$) then Cm[Lc[v]] $\leftarrow v$

Note that we write $((a \leftarrow b) \leftarrow c)$ to mean we first assign b to a and then assign c to b . We do not have to exchange $\text{Wt}[u]$ and $\text{Wt}[v]$ since these have the same values! A swap is done only if $v > u$ (Line 2). Thus the rank of the current node u is strictly increased by such swaps. After swapping u and v , their siblings will change (recall that rank $2j$ and rank $2j + 1$ nodes are siblings).

The reader may verify that the informal example of Figure 6 is really an operation of the RESTORE routine.

But let us walk through an example of the operations of RESTORE, this time seeing its transformation on the Lc, Wt arrays. Suppose we have just completely processed our famous string “hello world!”, and assume that the resulting Huffman tree T is given by Figure 3. Let the next character to be transmitted be \sqcup (space character), and set u to the node corresponding to \sqcup . So $\text{Wt}[u]$ is to be incremented, and we call RESTORE(u). We use the representation of T by the arrays Lc, Wt above: in this case u is the node (whose rank is) 2 (or v_2 , for clarity). It has weight $\text{Wt}[v_2] = 1$ and so we must find the largest ranked node with weight 1, namely node v_6 . Swapping v_2 with v_6 , and then incrementing the weight of v_6 , we get:

Rank	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16
Lc	h	e	!	w	r	d	□	o	0	2	4	6	ℓ	8	10	12	15
Wt	1	1	1	1	1	1	1+1	2	2	2	2	3	3	4	5	7	12
After first swap			v				u										

Next, u is set to the parent of node of rank 6, namely v_{11} . This has weight 3, and so we must swap it with the element v_{12} which is the highest ranked node with weight 3. After swapping v_{11} and v_{12} , we increment the new v_{12} . The following table illustrates the remaining changes:


Rank	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16
Lc	h	e	!	w	r	d	□	o	0	2	4	ℓ	6	8	10	12	15
Wt	1	1	1	1	1	1	1+1	2	2	2	2	3	3+1	4	5	7	12
After second swap												v	u				
Lc	h	e	!	w	r	d	□	o	0	2	4	ℓ	6	8	10	12	15
Wt	1	1	1	1	1	1	1+1	2	2	2	2	3	3+1	4	5	7+1	12
No third swap																u = v	
Lc	h	e	!	w	r	d	□	o	0	2	4	ℓ	6	8	10	12	15
Wt	1	1	1	1	1	1	1+1	2	2	2	2	3	3+1	4	5	7+1	12+1
No final swap																	u = v

¶29. **How to add a new letter: the 0-Node.** Our dynamic Huffman code tree T must be capable of expanding its alphabet. E.g., if the current alphabet is $\Sigma = \{h, e\}$ and we next encounter the letter l , we want to expand the alphabet to $\Sigma = \{h, e, l\}$. For this purpose, we introduce in T a special leaf with weight 0. Call this the **0-node**. This node does not represent any letters of the alphabet, but in another sense, it represents all the yet unseen letters. We might say that the 0-node represents the character ‘*’. Upon seeing a new letter like l , we take three steps to update T :

1. First, we “expand” the 0-node by giving it two children. Its left child is the new 0-node, and its right child u is a new leaf representing the letter l .
2. Next, we must give ranks to all the nodes: the new 0-node has rank 0, the new leaf u has rank 1, and all the previous nodes have their ranks incremented by 2. In particular, the original 0-node will have rank 2.
3. Finally, we must update the weights. The weight of the new 0-node is 0, and the weight of u is 1. We must now increase the weight of all the nodes along the path from the old 0-node to the root: this is done by calling `RESTORE` on the old 0-node.

Observe that our algorithm for expanding the 0-node maintains the following useful invariant: *the 0-node has rank 0 and its parent has rank 2*. The operations of the restore function using this 0-node convention is illustrated in Figure 7. Here, we begin with an initial Huffman tree containing just the 0-node, and show successive Huffman trees on inserting the first five letters of our hello example.

Note that the transition from `hel` to `hell` is already described in detail in Figure 6.

Figure 7: Evolving Huffman tree on inserting the string `hello`

The existence of the 0-node actually causes a wrinkle in the RESTORE routine. We leave the fix to an Exercise.

¶30. Interface between Huffman Code and Canonical Encoding. Let Σ denote the set of characters in the current Huffman code. We view Σ as a subset of a fixed universal set U where $U \subseteq \{0,1\}^N$. Call U the **canonical encoding**. In reality, U might be the set of ASCII characters with $N = 8$. A more complicated example is where U is some unicode set. We assume the transmitter and receiver both know this global parameter N and the set U . In the encoding process, we assume that each character of the string comes from U . Upon seeing a letter x , we must decide whether $x \in \Sigma$ (i.e., in our current Huffman tree), and if so, what is its current Huffman code. If $|U|$ is not too large (e.g., $|U| = 2^8$), we can provide an array $C[1..2^N]$ such that $C[x]$ maps to a leaf of the Huffman tree. To be specific, suppose $|\Sigma| = k$ and the current Huffman tree T is represented by the arrays $\text{Lc}[0..2k], \text{Wt}[0..2k]$. If $x \in \{0,1\}^N$, let $C[x] = i$ if node i (of rank i) is the leaf of T representing the letter x . Initially, let $C[x] = -1$ for all x . Hence, the array C is a representation of the alphabet Σ .

For instance, if $C[x]$ is the 0-node, this means x is not in Σ . If $|U|$ is large, we can use hashing techniques.

Even though we know the leaf, it requires some work to obtain the corresponding Huffman code. [This is the encoding problem – but the Huffman code tree is specially designed for the inverse problem, i.e., decoding problem.] One way to solve this encoding problem is assume that our Huffman tree has parent pointer. In terms of our Lc, Wt array representation, we now add another array $P[0..2k]$ for parent pointers.

Here now is the dynamic Huffman coding method for transmitting a string s :

DYNAMIC HUFFMAN TRANSMISSION ALGORITHM:

Input: A string s of indefinite length.

Output: The dynamically encoded sequence representing s .

▷ *Initialization*

Initialize T to contain just the 0-node.

▷ *Main Loop*

While s is non-empty

1. Remove the next character x from the front of string s .
 2. Let $u = C[x]$ be the leaf of T that corresponds to x .
 3. Using u , transmit the code word for x .
 4. If u is the 0-node ◁ *x is a new character*
 5. Expand the 0-node to have two children, both with weight 0;
 6. Let u be the right sibling, representing the character x
 and the left sibling represent the new 0-node.
 7. Call $\text{RESTORE}(u)$.
- Signal termination, using some convention.

Decoding is also relatively straightforward. We are processing a continuous binary sequence, but we know where the implicit “breaks” are in this continuous sequence. Call the binary sequence between these breaks a **word**. We know how to recognize these words by maintaining the same dynamic Huffman code tree T as the transmission algorithm. For each received word, we know whether it is (a) a code word for some character, (b) signal to add a new letter to the alphabet Σ , or (c) the canonical representation of a letter. Thus the receiver can faithfully reproduce the original string s .

Another practical issue is that whenever we insert a new node, the rank of each existing node implicitly increases by 2. A literal implementation requires updating the entire array for Lc and Wt . There is a simple solution to this. Let us store the array in reverse order. All invocations of $\text{Lc}[i]$ is really an invocation of $\text{Lc}[2k - i]$. Similarly for $\text{Wt}[i]$. We leave it to the student to work out this detail.

REMARKS: It can be shown that the FGK Algorithm transmits at most $2H_2(s) + 4|s|$ bits. The Lambda Algorithm of Vitter ensures that the transmitted string length is $\leq H_2(s) + |s| - 1$ where $H_2(s)$ is the number of bits transmitted by the 2-pass algorithm for s , independent of alphabet size. In Chapter VI, we will show another approach to dynamic compression of strings based on the move-to-front heuristic and splay trees [1].

¶31. **Beyond Huffman Coding: Lempel-Ziv Coding.** Although Huffman coding is optimal, we can go beyond its character-by-character encoding assumption. In other words, we can look for blocks of characters that occur with high frequency. For example, in English, certain digraphs like **in**, **at**, **th**, **ng**, etc, occur with high frequency. Certain trigraphs like **the**, **ing**, **ion**, **and**, etc, are also very common. But more generally, there is no need to restrict attention to blocks of any fixed size, but to let the input string determine this. This is the basis of another highly successful encoding scheme due to J.Ziv and A.Lempel [4].

¶32. **Notes on the ASCII Character Set.** ASCII stands for *The American Standard Code for Information Interchange*, and refers to the 7-bit encoding of 128 characters. See Figure 8(a) for

	0	1	2	3	4	5	6	7
0	NUL	DLE	space	0	@	P		p
1	SOH	DC1	XON	1	A	Q	a	q
2	STX	DC2	"	2	B	R	b	r
3	ETX	DC3	XOFF	3	C	S	c	s
4	EOT	DC4	\$	4	D	T	d	t
5	ENQ	NAK	%	5	E	U	e	u
6	ACK	SYN	&	6	F	V	f	v
7	BEL	ETB	'	7	G	W	g	w
8	BS	CAN	(8	H	X	h	x
9	HT	EM)	9	I	Y	i	y
A	LF	SUB	*	:	J	Z	j	z
B	VT	ESC	+	:	K	[k	{
C	FF	FS	,	<	L	\	l	
D	CR	GS	-	=	M]	m	}
E	SO	RS	.	>	N	^	n	~
F	SI	US	/	?	O	_	o	del

(a) ASCII set

	8	9	A	B	C	D	E	F
0	€			°	À	É	à	â
1		'	ı	±	Á	Ê	á	ã
2	,	'	¢	²	Â	Ë	â	ä
3	f	"	£	³	Ã	Ì	ã	å
4	"	"	¤	´	Ä	Ó	ä	ö
5	...	*	¥	µ	Å	Ô	å	ø
6	†	—	ı	¶	Æ	Õ	æ	ù
7	‡	—	§	·	Ç	×	ç	÷
8	^	~	¨	¸	È	Ø	è	ø
9	%	™	©	¹	É	Ù	é	ú
A	Š	š	ª	º	Ê	Ú	ê	û
B	<	>	«	»	Ë	Û	ë	ü
C	CE	œ	¬	¼	İ	Ü	ı	ü
D				½	Í	Ý	í	ý
E	Ž	ž	®	¾	Î	Þ	î	þ
F		ÿ	—	¿	Ï	ß	ï	ÿ

(b) ANSI extension

Figure 8: Extended ASCII

this character set. 95 of these are **printable characters** such as 0, 1, ..., 9, a, b, c, ..., x, y, z, and A, B, C, ..., X, Y, Z, including the space character which we denote by '␣'. The remaining 33 are non-printing or **control characters** such as backspace (BS), carriage-return (CR), bell (BEL), etc. They include many obsolete ones associated with technology such as teletype from a bygone era. There is a natural sorting order associated with these characters, so we may speak of the first, second, and last characters. For instance, the table in Figure 8(a) shows that the first and last characters of these first 128 characters are called NUL and DEL, respectively. Indeed, the first 32 characters (i.e., the first two columns in Figure 8(a)) are control characters, as is the last DEL.

The basic ASCII character set has been extended into a set of 256 characters; Figure 8(b) shows this so-called ANSI extension (ANSI stands for *American National Standards Institute*). For example, two special symbols we use throughout in this book comes from the extension: the section symbol '§' and paragraph symbol '¶'. The 256 characters is naturally associated with an 8-bit binary string called its **ASCII code**. We shall write $\text{ASCII}(x)$ to denote the ASCII code of a character x . The original 128 characters in the ASCII character set, naturally, occupy the first 128 positions; thus a character x belongs to this original set iff the most significant bit in $\text{ASCII}(x)$ is 0. Of course, the 8-bits can be broken up into two groups of 4-bits which are then interpreted as hexadecimal digits. This much more compact notation is preferred. The 16 hexadecimal digits are conventionally written as 0, 1, 2, ..., 9, A, B, C, ..., F. Because of the overlap between standard decimal digits and hexadecimal digits, a sequence like '10' would be ambiguous. To indicate that the hexadecimal digits are meant, we typically prefix the digits by '0x'. Thus, the ASCII code for the character A is $(01000001)_2$ in binary or 0x41 in hexadecimal notation. Here are some ASCII codes:

$$\begin{aligned}\text{ASCII}(\text{A}) &= 0x41, & \text{ASCII}(\text{Z}) &= 0x5A, \\ \text{ASCII}(\text{a}) &= 0x61, & \text{ASCII}(\text{z}) &= 0x71, \\ \text{ASCII}(\text{0}) &= 0x30, & \text{ASCII}(\text{9}) &= 0x39, \\ \text{ASCII}(\text{␣}) &= 0x20, & \text{ASCII}(\text{*}) &= 0x2A.\end{aligned}$$

Note that ␣ (space) is considered a printing character. These are all part of the original ASCII set since the first hexadecimal digit are all in the range 0 to 7. In contrast, the extended characters begin with a digit in the range 8 to F. E.g., $\text{ASCII}(\text{§}) = 0xA7$ and $\text{ASCII}(\text{¶}) = 0xB6$. Actually, not all of the available codes in the extension has been assigned to characters. A recent assignment is $\text{ASCII}(\text{Euro}) = 0x80$ for the Euro symbol.

There are many variants of the ASCII encoding, as many countries adapted ASCII to their unique requirements. Today, the ASCII encoding is re-interpreted as the first 128 characters of the UTF-8

A tele... wha??

You know the joke: there are 10 kinds of people in the world — those who count in binary and those who count in decimal.

encoding. The latter is part of the **Unicode**, a character encoding system of amazing scope and generality. This system will be briefly described next.

The original ASCII set together with the ANSI extension is often called the **extended ASCII set**. However, unless otherwise noted in this book, *we will use “ASCII set” to refer to the extended ASCII set..*

¶33. Notes on Unicode. The Unicode is an evolving standard for encoding the character sets of most human languages, including dead ones like Egyptian hieroglyphs. Here, we must make a basic distinction between **characters** (or graphemes) and their many **glyphs** (or graphical renderings). The idea is to assign a unique number, called a **code point**, to each character. Typically, we write such a number as U+XXXXXX where the X’s are hexadecimal. As usual, leading zeros are insignificant. For instance the first 128 code points in Unicode, U+0000 to U+007F, correspond to the ASCII code. The code points below U+0020 are control characters in ASCII code. But there are many subtle points because human languages and writing are remarkably diverse. Characters are not always atomic objects, but may have internal structure. Thus, should we regard “**é**” as a single Unicode character, or as the character “**e**” with a combining acute “**ˆ**”? Answer: *both solutions are provided in unicode*. If combined, what kinds of combinations do we allow? Coupled with this, we must meet the needs of computer applications: computers use unprintable or control characters, but should these be characters for Unicode? Answer: *of course, this is already part of ASCII*.

There are other international standards (ISO) and these have some compatibility with Unicode. For instance, the first 256 code points corresponds to ISO 8859-1. There are two methods for encoding in Unicode called Unicode Transformation Format (UTF) and Universal Character Set (UCS). These leads to UTF-*n*, UCS-*n* for various values of *n*. Let us just focus on one of these, UTF-8. This was created by K.Thompson and R.Pike, which is a de facto standard in many applications (e.g., electronic mail). It has a basic 8-bit format with variable length extensions that uses up to 4 bytes (32 bits). It is particularly compact for ASCII characters: only 1 byte suffices for the 127 US-ASCII characters. A major advantage of UTF-8 is that a plain ASCII string is also a valid UTF-8 string (with the same meaning of course). Here is UTF-8 in brief:

1. Any code point below U+0080 is encoded by a single byte. Of course, 080 in hex is just 128 in decimal. Thus, U+00XY where $X < 8$ can be represented by the single byte XY that has a leading 0-bit.
2. Code points between U+0080 to U+07FF uses two bytes. The first byte begins with 110, second byte begins with 10.
3. Code points between U+0800 to U+FFFF uses three bytes. The first byte begins with 1110, remaining two bytes begin with 10.
4. Code points between U+10000 to U+10FFFF uses four bytes. The first byte begins with 11110, remaining three bytes begin with 10.

Observe that each code point is self-limiting, i.e., you can tell when you have reached the end of a code point.

EXERCISES

Exercise 5.1: In this question, we are asking for three numbers. But you must summarize to show intermediate results of your computations. Assume that the alphabet Σ is a subset of $\{0, 1\}^8$ (i.e., ASCII code).

- (a) What is the length of the (static) Huffman code of the string “hello, world!”? The

quotation marks are not part of the string, but the space and punctuation marks are.

(b) How many bits does it take to transmit the Huffman code for the string of (a)?

(c) How many bits would be transmitted by the Dynamic Huffman code algorithm in sending the string “hello, world!”? Compare this number with (a)+(b). ◇

Exercise 5.2: (Gashlin, 2012) Our description of the `RESTORE(u)` routine is correct only if every leaf has positive weight. In the presence of the 0-node, the routine is incorrect.

(a) Describe the precise situation where the routine fails.

(b) Provide the correct solution (be sure to justify it). ◇

Exercise 5.3: What binary string would you transmit in order to send the string “now is the time”, under the dynamic Huffman algorithm? Show your working. Note: you would have to transmit ASCII codes for the letters n, o, w, etc. Just write `ASCII(n)`, `ASCII(o)`, `ASCII(w)`, etc. ◇

Exercise 5.4: Natural languages are highly redundant (for good reasons). Here is one way to test this.

(a) Please transmit the following string using dynamic Huffman coding, and state the bit length of your transmission.

Aoccdnrng to a rscheearch at Cmabrigde Uinervtisy, it deosn't mtttaer
in waht oredr the ltteers in a wrod are, the olny iprmoetnt tihng is
taht the frist and lsat ltteer be at the rghit pclae. The rset can be
a total mses and you can sitll raed it wouthit porbelm. Tihs is
bcuseae the huamn mnid deos not raed ervey lteter by istlef, but the
wrod as a wlohe.

(b) Now repeat part (a) but using similar string in which the words are now properly spelled. Should we expect a drop in the number of transmitted bits? Note that this experiment could not be done using standard Huffman coding since the frequency function in (a) and (b) are identical. ◇

Exercise 5.5:

(a) Please reconstruct the Huffman code tree T from the following representation:

$$r(T) = 0000, 1111, 0011, 011d, mrit, yo$$

CONVENTIONS: the commas in $r(T)$ are just decorative, and meant to help you parse the string. Other than 0/1 symbols, the letters d, m, i , etc, stands for 8-bit ASCII codes. The leftmost leaf in the tree is the 0-node, and its label (namely ‘*’) is implicit. NOTE that the ‘*’ letter actually has ASCII code `0x2A` in hex, so it is not literally the smallest letter. The smallest letter in the ASCII code is `0x00` and is usually called the NUL character. The remaining leaves are labeled by 8-bit ASCII codes for d, m, r, i, t, y, o , in left-to-right order.

(b) Here is a string encoded using this Huffman code:

$$0001, 1110, 1001, 1001, 0111, 1011, 10$$

Decode the string.

(c) Assume that the leaves of the Huffman tree in (a) has the following frequencies (or weights):

$$f(*) = 0, \quad f(d) = f(m) = f(i) = f(t) = d(y) = 1, \quad f(r) = f(o) = 2.$$

Assign a rank (i.e., numbers from $0, 1, \dots, 14$) to the nodes of the tree in (a) so that the sibling property is obeyed. Redraw this tree with the ranking listed next to each node. Also, write the arrays $\text{Lc}[0..14]$ and $\text{Wt}[0..14]$ which encodes this ranking of the Huffman tree. Recall that these arrays encode the left-child relation and weights (frequencies), respectively.

(d) Suppose that we now insert a new letter \sqcup (blank space) into the weighted Huffman code tree of (c). Draw the new Huffman tree with updated ranking. Also, show the updated arrays $\text{Lc}[0..16]$ and $\text{Wt}[0..16]$.

(e) Give the Huffman code for the string “dirty room” (this string has is a blank character \sqcup , but the quotes are not part of the string). What is the relation between this string and the one in (d)? \diamond

Exercise 5.6: Give the dynamic Huffman coding for the following anagrams: (1) the morse code (2) here come dots \diamond

Exercise 5.7: Assume the Sibling Representation of the Huffman code $C : \Sigma \rightarrow \{0, 1\}^*$. Give the routine to compute the code word $C(x) \in \{0, 1\}^*$ of any given $x \in \Sigma$. \diamond

Exercise 5.8: Give a careful and efficient implementation of the dynamic Huffman code. Assume the compact representation of Huffman tree using the arrays Wt and Lc described in the text. \diamond

Exercise 5.9: Consider 3-ary Huffman tree code. State and prove the Sibling property for this code. \diamond

Exercise 5.10: A previous Exercise asks you to construct the standard Huffman code of Lincoln’s speech at Gettysburg.

(a) Construct the optimal Huffman code tree for this speech. Please give the length of Lincoln’s coded speech, and also the size of the code tree.

(b) Please give the length of the dynamic Huffman code for this speech. How does it compare to part (a)? Also, compare the code tree at the end of the dynamic coding process with the one in part (a). \diamond

Exercise 5.11: In the text, we have represented the Huffman code tree ways: as an binary code tree and as the arrays Wt , Lc , Cm . The former gives $O(|C(x)|)$ complexity for finding the code of $x \in \Sigma$, but the latter is only $O(|\Sigma|)$. Show how to improve the latter complexity in the setting of dynamic Huffman coding. \diamond

Exercise 5.12: The correctness of the dynamic Huffman code depends on the fact that the weight at the leaves are integral and the change is $+1$.

(a) Suppose the leave weights can be any positive real number, and the change in weight is also by an arbitrary positive number. Modify the algorithm.

(b) What if the weight change can be negative? \diamond

Exercise 5.13: Programming Project. Suppose we are given s_1 , a string of (extended) ASCII characters. Let C_1 be the static Huffman code for s_1 . Consider the binary string

$$\alpha_{C_1} C_1(s_1)$$

where the first part α_{C_1} is the encoding of C_1 described in ¶23 above, and the second part $C_1(s_1)$ is just the application of C_1 to s_1 . Break this up into 8-bit blocks and so reinterpret it as an ASCII string denoted s_2 . If necessary, we pad this string with 0's so that the last block still has 8 bits. This transformation $s_1 \rightarrow s_2$ can be repeated: $s_2 \rightarrow s_3$, etc. What is the limit of this process? \diamond

Exercise 5.14: Programming Project. Compare the use of a fixed Huffman code of a large document versus a dynamic Huffman encoding. For the fixed Huffman code, use the following statistics¹⁴ on the frequency English letters based on a sample of 40,000 words:

Letter:	E	T	A	O	I	N	S	R	H	D	L	U	C
Count:	21912	16587	14810	14003	13318	12666	11450	10977	10795	7874	7253	5246	4943
Frequency:	12.02	9.10	8.12	7.68	7.31	6.95	6.28	6.02	5.92	4.32	3.98	2.88	2.71

Letter:	M	F	Y	W	G	P	B	V	K	X	Q	J	Z
Count:	4761	4200	3853	3819	3693	3316	2715	2019	1257	315	205	188	128
Frequency:	2.61	2.30	2.11	2.09	2.03	1.82	1.49	1.11	0.69	0.17	0.11	0.10	0.07

END EXERCISES

§6. Minimum Spanning Tree

¶34. **Minimizing over Maximal Sets.** In the **minimum spanning forest problem** we are given a costed bigraph

$$G = (V, E; C)$$

where $C : E \rightarrow \mathbb{R}$. An acyclic set $T \subseteq E$ of maximum cardinality is called a **spanning forest**; in this case, $|T| = |V| - c$ where G has $c \geq 1$ components. The **cost** $C(T)$ of any subset $T \subseteq E$ is given by $C(T) = \sum_{e \in T} C(e)$. An acyclic set is **minimum** if its cost is minimum. It is conventional to make the following simplification:

The input bigraph G is connected.

With this assumption, a spanning forest T is actually a tree, and the problem is known as the **minimum spanning tree (MST) problem**. The simplification is not too severe: if our graph is not connected, we can first compute its connected components (we saw efficient solutions to this basic graph problem in Chapter IV). Then we apply the MST algorithm to each component. Alternatively, it is not hard to modify most MST algorithms so that they apply to non-connected graphs.

An important point to note is that we are minimizing over spanning trees that are maximal acyclic edge sets. In general, minimax problems can have high complexity. But for MST, these maximal sets has a great deal of structure – one formalization is the concept of matroid structure in the next section.

Consider the costed bigraph in Figure 9 with vertices $V = \{a, b, c, d, e\}$. One such MST is $\{a-b, b-c, c-d, d-e\}$, with cost 6 and shown in Figure ??(a). It is easy to verify that there are a total five MST's for this bigraph, as shown in Figure ??.

¹⁴Source: [Math Explorer's Club, Cornell Math Department](#).




Figure 9: A bigraph with edge costs.

¶35. **Some Global Properties of MSTs.** Fix a connected bigraph $G = (V, E; C)$ and let $MST(G)$ denote the set of all MST's of G . Each $T \in MST(G)$ has size $n - 1$ where $n = |V|$. If $T, T' \in MST(G)$, notice that their symmetric difference $T \oplus T' = (T \setminus T') \cup (T' \setminus T)$ has an even cardinality since T and T' have the same size implies $|T \setminus T'| = |T' \setminus T|$.

Define the **(MST) exchange graph** $Exch(G)$ to be the bigraph $(MST(G), H)$ whose nodes are MST's and edges $T - T' \in H$ are characterized by the property that $T \oplus T'$ has size 2.

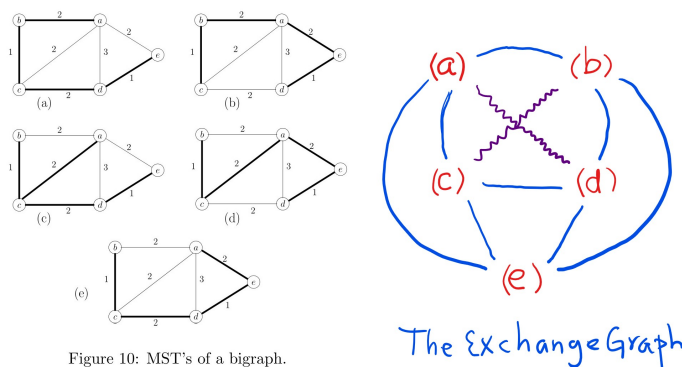


Figure 10: MST's of a bigraph.

Figure 10: Exchange graph of G_5

Lemma 13 (Exchange) If $T \neq T' \in MST(G)$, then there exists $e \in T \setminus T'$ and $e' \in T' \setminus T$ such that $T + e' - e \in MST(G)$.

Proof. Pick $e' = \operatorname{argmin}_{e \in T' \setminus T} C(e)$. Then $T + e'$ has a unique cycle Z . In this cycle, there exists some $e \in T \setminus T'$ (otherwise $Z \subseteq T'$, contradiction). Now $T - e + e'$ is a tree. Also $C(e) \leq C(e')$ since T is MST. If $C(e) = C(e')$ then $T - e + e'$ is our desired MST, proving our claim.

Otherwise, $C(e) < C(e')$ and we obtain a contradiction: by the same argument, $T' + e$ has a unique cycle Z' and there exists in this cycle some $e'' \in T' \setminus T$. Also $C(e'') \leq C(e)$ since T' is a MST. This implies $C(e'') \leq C(e) < C(e')$, contradicting our initial choice of e' as an argmin. **Q.E.D.**

Note that if all the edges of G have unit cost, then $MST(G)$ is just the set of all spanning trees.

Lemma 14 Let T, T' be two MST's of a connected bigraph on $n \geq 2$ vertices. Then $|T \cap T'|$ is odd iff n is odd.

Proof. We use induction on n . Result is true for $n = 2, 3$. Suppose $n > 3$. If $T \neq T'$ then the exchange lemma, we can find another MST $S_1 = T - e + e'$ such that $|S_1 \cap T'| = |T \cap T'| + 2$. We can repeat this to find S_2, S_3, \dots such that $|S_i \cap T'| = |T \cap T'| + 2i$. We stop when $|S_i \cap T'|$ is $n - 1$ or $n - 2$. If $|S_i \cap T'| = n - 1$, this proves that... **Q.E.D.**

Lemma 15 $\text{Exch}(G)$ is a connected graph. of diameter at most $\lfloor n - 1/2 \rfloor$ where $n = |V|$.

May 2020: NEED TO RE-WRITE THIS WHOLE SECTION!!! IDEA: Say S is **feasible** if it is a subset of an MST. Call T a **feasible extension** of S if $S \subseteq T$. If $T = S + e$, we call T a **unit extension** of S . A unit-extension $S + e$ is called a Boruvka-extension if $e = (u-v)$ and $C(e)$ has the minimum cost among all the feasible edges that grow the S -component of u or of v . Not all unit extensions are Boruvka: E.g., in the graph G_6 , if $S = \dots$ then ...

¶36. Generic MST Algorithm. There several distinct algorithms for MST. They all fit into the following framework:

GENERIC GREEDY MST ALGORITHM
 Input: $G = (V, E; C)$ a connected bigraph with edge costs.
 Output: $S \subseteq E$, a MST for G .
 $S \leftarrow \emptyset$.
 for $i = 1$ to $|V| - 1$ do
 1. Greedy Step: find an $e \in E \setminus S$ that is “good for S ”.
 2. $S \leftarrow S + e$.
 Output S as the minimum spanning tree.

NOTATION: as illustrated in Line 2, we shall write “ $S + e$ ” for “ $S \cup \{e\}$ ”. Likewise, “ $S - e$ ” shall denote the set “ $S \setminus \{e\}$ ”.

What does it mean for “ e to be good for S ”? This will be made specific next.

¶37. Some Greedy MST Criteria. Let us say that e is a **candidate** for S if $S + e$ is acyclic. Consider the graph $G|S = (V, S)$, the restriction of G to S . The connected components of $G|S$ are called **S -components**. For instance if S is the empty set, then each vertex forms its own S -component (such components are said to be **trivial**). For any candidate $e = u-v$, let the S -component that contains u and v (resp.) be denoted C_u and C_v . Clearly $C_u \cap C_v$ is empty (otherwise we get a cycle). Moreover, $C_u \cup C_v$ is a $(S + e)$ -component. We say that e **extends** the component C_u (and also C_v by symmetry). Among the candidates for S , only some are “good”. Here are 4 notions of what “good” means:

- (Minimal) $S + e$ is extendible to some MST.

- (Kruskal) Edge e has the least cost among all the candidates.
- (Boruvka) If $e = (u-v)$ then either (i) e has the least cost among all the candidates that extend the component C_u , or (ii) e has the least cost among all the candidates that extend the component C_v . In case (i), we call C_u a **Boruvka witness** for e ; similarly for C_v in case (ii). It is possible that both C_u and C_v are witnesses for e .
- (Prim) This can be viewed adding an extra requirement to Boruvka's condition: assume we are given a vertex $s \in V$ called the source. The edge e must have least cost among candidates that extends the unique component $U_s \subseteq V$ containing s .

A set $S \subseteq E$ that may arise during the execution of the generic MST algorithm is said to be **X-good** where $X \in \{\text{minimally, Boruvka, Kruskal, Prim}\}$ depending on the criteria used. By assumption, the empty set S is X-good for any X. The correctness of these algorithms amounts to showing that “*X-good implies minimally-good*” where $X = \text{Kruskal, Boruvka or Prim}$. We now prove this:

Lemma 16 (Correctness of Algorithm X)

Let $S \subseteq E$.

- (a) If S is Prim-good then S is Boruvka-good.
- (b) If S is Kruskal-good then S is Boruvka-good.
- (c) If S is Boruvka-good then S is minimally-good.

Proof. In each case, we want to show that if S is X-good, then it is Y-good, for appropriate X and Y. We use induction on $|S|$. When $|S| = 0$ and the lemma holds trivially. Suppose S and $S + e$ are X-good. By induction, S is Y-good. It remains to show that $S + e$ is Y-good.

- (a) X=Prim, Y=Boruvka: We must show that $S + e$ is Boruvka-good. Since $S + e$ is assumed to be Prim-good, we know e has least cost among the edges that extend the S -component C_s containing the source s . So C_s is a Boruvka-witness for e . Thus $S + e$ is Boruvka-good.
- (b) X=Kruskal, Y=Boruvka: We must show that $S + e$ is Boruvka-good. Since $S + e$ is assumed to be Kruskal-good, we know e has least cost among the edges that extend *any* S -component. The S -component for which e has least cost serves as a Boruvka-witness for e . Thus $S + e$ is Boruvka-good.
- (c) X=Boruvka, Y=minimally: We need to prove that $S + e$ is minimally-good. By the Boruvka-goodness of S , there is a S -component U which is the Boruvka-witness for e . By induction hypothesis, S is minimally-good. Hence there is a MST T that contains S . If $e \in T$, then we are done (as T is witness that $S + e$ is minimally-good). So assume $e \notin T$. This means that T_e contains a closed path Z .


Let $e = u-v$ where $u \in U$ and $v \notin U$. The closed path Z has the form

$$Z := (u-v-v_1-v_2-\cdots-v_k-u).$$

Clearly, there exists some $i = 0, \dots, k$ such that $v_i \notin U$ and $v_{i+1} \in U$ where $v = v_0$ and $u = v_{k+1}$ in this notation. Rewriting,

$$Z = (u-v-v_1-\cdots-v_i-v_{i+1}-\cdots-u).$$

Let $e' := (v_i-v_{i+1})$. Note that $T' := (T - e') + e$ is acyclic (since we broke the unique cycle Z by omitting e'). It is also a spanning tree. Looking at costs, we have $C(e) \leq C(e')$, by our choice of e as least cost edge out of U . Hence $C(T') \leq C(T)$. Since T is MST, this means T' must also be an MST. Thus S is minimally-good because S is contained in the MST T .

Figure 11: Extending a component U by $e = (u, v)$.

Q.E.D.

This minimally-good criterion is computational ineffective. The remaining three criteria are effective and are named after the inventors of three well-known MST algorithms. We next discuss the algorithmic techniques needed to make these criteria effective:

- (Kruskal) Kruskal tells us to sort the edges first, and then consider each edge e in order of increasing cost. How can we quickly tell if $S + e$ is acyclic? If e is $u-v$, this amounts to checking if u, v are in the same connected component of the graph $G' = (V, S)$. A simple method is to do this is have a linked list for each connected component of $G' = (V, S)$, with the nodes of the linked list representing vertices of the component. Given a vertex u , assume we have a pointer from u to the representative for u in such a linked list. To decide if two vertices u, v are in the same connected component, we go to the linked lists nodes that represent u and v , and follow the links till the end of their respective linked lists. *The ends of these two linked list are equal iff $S + e$ has a cycle.*

The elaborations of this linked list idea will ultimately lead us to the union-find data structure which is studied in Chapter XIII. An Exercise below will explore some of these ideas.

- (Boruvka) REWRITE: suppose S is a Boruvka-good set. We want to talk about a set of edges T where each $e \in T$ is a “ S -extension”... But current language is confusing!

If S is not yet a spanning tree of G , the number of S -components must at least 2. Let these S -components be U_1, \dots, U_k ($k \geq 2$). For each U_i , there is at least one e_i that extends U_i with least cost. These e_i ’s need not be distinct (it is possible that $e_i = e_j$ with $i \neq j$). Nevertheless, there are at least $\lceil k/2 \rceil$ distinct choices for e_i . Algorithmically, we want maintain the least cost edge that extends each S -component. Again we can exploit the union-find data structure. The key feature of Boruvka’s algorithm is that we can select the good edges in “phases” where each phase calls for a pass through the set of remaining edges. This feature can be exploited in parallel algorithms.

- (Prim) Because of its focus on one component, Prim’s algorithm is somewhat easier to implement than Boruvka’s. The ultimate version of Prim’s algorithm can only be taken up in Chapter VI (amortization techniques).

¶38. **Good sets of vertices.** Let us extend the notion of “goodness” to sets of vertices. For any set $S \subseteq E$ of edges, let $V(S)$ denote the set of vertices that are incident on some edge of S . We say a set $U \subseteq V$ is **X -good** if there exists an X -good set $S \subseteq E$ such that $U = V(S)$. Here, X is equal to ‘minimally’, ‘Prim’, ‘Kruskal’ or ‘Boruvka’. We also declare any singleton set with only one vertex to be X -good.

¶39. **Hand Simulation of MST Algorithms.** Students are expected to understand those aspects of Kruskal’s and Prim’s algorithms that are independent of their ultimate realizations via efficient data structures. That is, you must do “hand simulations” where you act as the oracle for queries to the data structures. For Kruskal’s algorithm, this is easy – we just list the edges by increasing weight order and indicate the acceptance/rejection of successive edges.

For Prim’s algorithm, we maintain an array $d[1..n]$ assuming the vertex set is $V = \{1, \dots, n\}$. We shall maintain a subset $S \subseteq V$ representing the set of vertices which we know how to connect to the source node 1 in a MST. The set S is “Prim good”. Initially, let $S = \emptyset$ and $d[1] = 0$ and $d[v] = \infty$ for $v = 2, \dots, n$. In general, the entry $d[v]$ ($v \in V \setminus S$) represents the “cheapest” cost to connect vertex v to the MST on the set S . Our simulation consists in building up a matrix M which is a $n \times n$ matrix, where the 0th row representing the initial array d . Each time the array d is updated, we rewrite it as a new row of a matrix M .

At stage $i \geq 1$, suppose we pick a node $v_i \in V \setminus S$ where $d[v_i] = \min\{d[j] : j \in V \setminus S\}$. We add v_i to S , and update all the values $d[u]$ for each $u \in V \setminus S$ that is adjacent to v_i . The update rule is this:

$$d[u] = \min\{d[u], C[v_i, u]\}.$$

The resulting array is written as row i in our matrix.




Figure 12: The house graph: cost of edge v_i-v_j is $C(v_i) + C(v_j)$, where $C(v)$ is the value next to v . E.g. $C(v_1-v_4) = 1 + 6 = 7$.

Let us illustrate the process on the graph of Figure 12. The vertex set is $V = \{v_1, v_2, \dots, v_{11}, v_{12}\}$. The cost of an edge is the sum of the costs associated to each vertex. E.g., $C(v_1, v_4) = C(v_1) + C(v_4) = 1 + 6 = 7$. The final matrix is the following:

Stage	1	2	3	4	5	6	7	8	9	10	11	12
0	0	∞	∞	∞	∞	∞	∞	∞	∞	∞	∞	∞
1	X	3	1	7	∞	∞	∞	∞	∞	∞	∞	∞
2			X	6				3				
3		X			6							
4							X		8			
5				X		7						
6					X		6					6
7						3	X			3	2	
8						1				1	X	2
9						X						
10									6	X		
11											X	
12									X			

Conventions in this matrix: We mark the newly picked node in each stage with an ‘X’. Also, any value that is unchanged from the previous row may be left blank. Thus, in stage 2, the node 3 is picked and we update $d[v_4]$ using $d[v_4] = \min\{d[v_4], C[v_3, v_4]\} = \min\{7, 6\} = 6$.

The final cost of the MST is 37. To see this, each X corresponds to a vertex v that was picked, and the last value of $d[v]$ contributes to the cost of the MST. E.g., the X corresponding to vertex 1 has cost 0, the X corresponding to vertex 2 has cost 3, etc. Summing up over all X’s, we get 37.

¶40. Boruvka’s MST Algorithm. Let $S \subseteq E$ be Boruvka-good, and $G_S = (V, S)$ is the subgraph of G with S as edge set. If C is a connected component of G_S , an edge $e = (u-v)$ where $u \in C$ and $v \notin C$ is called an **outgoing edge** of C . Recall that if e has minimum cost among all outgoing edges of C , then the extension $S + e$ is Boruvka-good. If G_S has k connected components, we pick one such minimum outgoing edge per component. Let e_i be the picked edge from component C_i ($i = 1, \dots, k$). There are two issues. First, it may happen that the same edge may be picked by two components: $e_i = e_j$ for $i \neq j$. But no more than two components may share the same minimum outgoing edge. Thus, among the k picked edges, there are at least $k/2$ (hence $\lceil k/2 \rceil$) distinct edges. Second, we want to ensure that $S \cup \{e_1, e_2, \dots, e_{k'}\}$ is acyclic. To do this, we must break ties carefully in choosing the e_i ’s. Assume there is a global ordering “ $<^*$ ” on the edge set E such that among the outgoing edges in a component with the same minimum weight, e_i is the minimum under this $<^*$ -order. It is now easy to see that there are no cycles. In particular, k' cannot be equal to k (otherwise $S \cup \{e_1, e_2, \dots, e_k\}$ has a cycle). In summary:

$$\left. \begin{array}{l} \text{the number } k' \text{ of distinct edge in the set } \{e_1, e_2, \dots, e_k\} \\ \text{is between } k/2 \text{ and } k-1, \text{ and } S \cup \{e_1, \dots, e_{k'}\} \text{ is acyclic.} \end{array} \right\} \quad (35)$$

This extension of S by simultaneously adding $k' \geq k/2$ edges is called a “phase”. These k' edges could be chosen “in parallel” if we were using parallel computers. But since we only consider on sequential algorithms, here is the conventional form of this algorithm.

```

Boruvka’s MST Algorithm
Input:  $G = (V, E; C)$  a connected graph
Output: MST  $S \subseteq E$  of  $G$ 
 $S \leftarrow \emptyset$      $\triangleleft$  Initialize a Boruvka-good set of edges
While ( $|S| < n - 1$ )
    for each connected component  $C$  of  $G_S$ ,     $\triangleleft$  Do Phase
         $e \leftarrow \operatorname{argmin} \{C(v, u) : v \in C, u \notin C\}$ 
         $S \leftarrow S + \{e\}$ 

```

How many phases can there be? If G_S has k components, and we form S' by adding k' distinct edges to S , then $G_{S'}$ has $k - k'$ connected components. From (35), we conclude that $1 \leq k - k' \leq k/2$. In other words, the number of connected components it at least halved. Since we started with n components, the number of phases is at most $\lg n$. We next show that each phase can be computed in time $O(m + n)$, and so the overall complexity of Boruvka’s algorithm is $O((m + n) \log n)$.

¶41. **Implementation of a Phase.** We now show how to implement a phase of Boruvka's algorithm in time $O(m + n)$. The key subroutine is a method to compute the connected components of a bigraph. In ¶IV.23, we have provided such an algorithm based on BFS.

- The main data structure is an array $CC[1..n]$ to keep track of the connected components of G_S . For each $i \in \{1, \dots, n\} = V$, let $CC[i]$ refer to an arbitrary but fixed vertex j in the connected component of i . This j is called the **representative** of that connected component. In particular, we always have $CC[j] = j$. Initially, $S = \emptyset$ and hence we have $CC[i] = i$ for all i . Assume inductively that the array $CC[1..n]$ is available at the beginning of the phase. If $CC[i] = j$, then we say i belongs to component j .
- Next, we shall compute the minimum outgoing edge for each connected component with the help of an array, $A[1..n]$. Let $A[j]$ store the current minimum outgoing edge from component j . We initialize $A[j] \leftarrow \text{nil}$, reflecting the fact that no outgoing edges are initially known. Note that if there are k components, then only k entries of the arrays A are in use. We incrementally update A as follows:

For each edge $(i-j) \in E$,
 If $CC[i] \neq CC[j]$ \triangleleft *$(i-j)$ is an outgoing edge of component of i*
 If the cost of $A[CC[i]]$ is greater than $C(i, j)$
 $A[CC[i]] \leftarrow (i-j)$.

At the end of this for-loop, it is clear that the array A has the desired information.

- We must now extend the set S to S' : this amounts to updating the array $CC[1..n]$ for the next phase. Imagine the reduced graph $G^r = (V^r, E^r)$ whose vertices are the set of representatives of connected components in $G|S$, and whose edges are the edges in $A[1..n]$. Viewed as a bigraph, our goal is to compute the connected components of G^r . For each $i \in V$, we retrieve the edge $(j-k) \leftarrow A[CC[i]]$. We add $CC[i]$ to V^r and $(CC[j]-CC[k])$ to E^r . Then we use a BFS Driver to compute the connected components of G^r . Assume (see ¶IV.23) that the BFS Driver returns an array CC^r such that for each $j \in V^r$, we have $CC^r[j]$ is a representative vertex in the connected component of j . Now we update CC with the help of CC^r :

For each $i \in V$,
 $CC[i] \leftarrow CC^r[CC[i]]$

This concludes the phase.

Remarks: Boruvka (1926) has the first MST algorithm; his algorithm was rediscovered by Sollin (1961). The algorithm attributed to Prim (1957) was discovered earlier by Jarník (1930). These algorithms have been rediscovered many times. See [9] for further references. Both Boruvka and Jarník's work are in Czech. The Prim-Jarník algorithm is very similar in structure to Dijkstra's algorithm which we will encounter in the chapter on minimum cost paths.

EXERCISES

Exercise 6.1: We consider minimum spanning trees (MST's) in an undirected graph $G = (V, E)$ where each vertex $v \in V$ is given a numerical value $C(v) \geq 0$. The **cost** $C(u, v)$ of an edge $(u-v) \in E$ is defined to be $C(u) + C(v)$.

(a) Let G be the graph in Figure 12. Compute an MST of G using Boruvka's algorithm.

Please organize your computation so that we can verify intermediate results. Also state the cost of your minimum spanning tree.

(b) Can you design an MST algorithm that takes advantage of the fact that edge costs has the special form $C(u, v) = C(u) + C(v)$? \diamond

Exercise 6.2: Redo the previous problem with a different cost function, where $C(u-v) = C(u)C(v)$ (the product instead of the sum). Is the result the same? \diamond

Exercise 6.3: Suppose G is the complete bipartite graph $G_{m,n}$. That is, the vertices V are partitioned into two subsets V_0 and V_1 where $|V_0| = m$ and $|V_1| = n$ and $E = V_0 \times V_1$. Give a simple description of an MST of $G_{m,n}$. Argue that your description is indeed an MST. HINT: transform an arbitrary MST into your description by modifying one edge at a time. \diamond

Exercise 6.4: Let G_n be the bigraph whose vertices are $V = \{1, 2, \dots, n\}$. It has two kinds of edges:

(i) **Prime edges:** for each $i \in V$, if i is prime, then $(1, i) \in E$ with weight i . [Recall that 1 is not considered prime, so 2 is the smallest prime.]

(ii) **Divisibility edges:** For $1 < i < j$, if i divides j then we add (i, j) to E with weight j/i .

(a) Draw the graph G_{10} .

(b) Compute the MST of G_{10} using Prim's algorithm, using node 1 as the source vertex. State the cost of the MST. Organize your working as outlined in ¶39.

(c) Are there special properties of the graphs G_n that can be exploited? \diamond

Exercise 6.5: Let $G = (V, E; W)$ be a connected bigraph with edge weight function W . Fix a constant M and define the weight function W' where $W'(e) = M - W(e)$ for each $e \in E$. Let $G' = (V, E; W')$. Show that T is a maximum spanning tree of G iff T is a minimum spanning tree of G' . NOTE: Thus we say that the concepts of maximum spanning tree and minimum spanning tree are “cryptomorphic versions” of each other. \diamond

Exercise 6.6: Describe the rule for reconstructing the MST from the matrix M using in our hand-simulation of Prim's Algorithm. \diamond

Exercise 6.7: Hand simulation of Kruskal's Algorithm on the graph of Figure 12. This exercise suggests a method for carry out the steps of this algorithm. We consider each edge in their sorted order, maintaining a partition of $V = \{1, \dots, 12\}$ into disjoint sets. Let $L(i)$ denote the set containing vertex i . Initially, each node is in its own set, i.e., $L(i) = \{i\}$. Whenever an edge $i-j$ is added to the MST, we merge the corresponding sets $L(i) \cup L(j)$. E.g., in the first step, we add edge 1-3. Thus the lists $L(1) = \{1\}$ and $L(3) = \{1\}$ are merged, and we get $L(1) = L(3) = \{1, 3\}$. To show the computation of Kruskal's algorithm, for each edge, if the edge is “rejected”, we mark it with an “X”. Otherwise, we indicate the merged list resulting from the union of $L(i)$ and $L(j)$: Please fill in the last two columns of the table (we have filled in the first 4 rows for you).

Sorting Order	Edge	Weight	Merged List	Cumulative Weight
1	1-3:	1	{1, 3}	1
2	6-11:	1	{6, 11}	2
3	10-11:	1	{6, 10, 11}	3
4	6-10:	2	X	3
5	7-11:	2		
6	11-12:	2		
7	1-2:	3		
8	3-8:	3		
9	6-7:	3		
10	7-10:	3		
11	2-5:	6		
12	3-4:	6		
13	5-7:	6		
14	5-12:	6		
15	9-10:	6		
16	1-4:	7		
17	4-6:	7		
18	8-9:	8		
19	4-5:	10		
20	4-9:	11		

◇

Exercise 6.8: This question considers two concrete ways to implement Kruskal's algorithm.

Let $V = \{1, 2, \dots, n\}$ and $D[1..n]$ be an array of size n that represents a **forest** $G(D)$ with vertex set V and edge set $E = \{(i, D[i]) : i \in V\}$. More precisely, $G(D)$ is an directed graph that has no cycles except for self-loops (i.e., edges of the form (i, i)). A vertex i such that $D[i] = i$ is called a **root**. The set V is thereby partitioned into disjoint subsets $V = V_1 \cup V_2 \cup \dots \cup V_k$ (for some $k \geq 1$) such that each V_i has a unique root r_i , and from every $j \in V_i$ there is a path from j to r_i . For example, with $n = 7$, $D[1] = D[2] = D[3] = 3$, $D[4] = 4$, $D[5] = D[6] = 5$ and $D[7] = 6$ (see Figure 13). We call V_i a **component** of the graph $G(D)$ (this terminology is justified because V_i is a component in the usual sense if we view $G(D)$ as an *undirected* graph).




Figure 13: Directed graph $G(D)$ with three components (V_1, V_2, V_3)

(i) Consider two restrictions on our data structure: Say D is **list type** if each component is a linear list. Say D is **star type** if each component is a star (i.e., each vertex in the component points to the root). E.g., in Figure 13, V_2 and V_3 are linear lists, while V_1 and V_2 are stars. Let $\text{ROOT}(i)$ denote the root r of the component containing i . Give a pseudo-code for computing $\text{ROOT}(i)$, and give its complexity in the 2 cases: (1) D is list type, (2) D is star type.

(ii) Let $\text{COMP}(i) \subseteq V$ denote the component that contains i . Define the operation $\text{MERGE}(i, j)$ that transforms D so that $\text{COMP}(i)$ and $\text{COMP}(j)$ are combined into a new component (but all the other components are unchanged). E.g., the components in Figure 13 are $\{1, 2, 3\}$, $\{4\}$ and $\{5, 6, 7\}$. After $\text{MERGE}(1, 4)$, we have two components, $\{1, 2, 3, 4\}$ and $\{5, 6, 7\}$. Give a pseudo-code that implements $\text{MERGE}(i, j)$ under the assumption that i, j are roots and D is list type which you must preserve. Your algorithm *must* have complexity $O(1)$. To achieve this complexity, you need to maintain some additional information (perhaps by a simple modification of D).

- (iii) Similarly to part (ii), implement $MERGE(i, j)$ when D is star type. Give the complexity of your algorithm.
- (iv) Describe how to use $ROOT(i)$ and $MERGE(i, j)$ to implement Kruskal's algorithm for computing the minimum spanning tree (MST) of a weighted connected undirected graph H .
- (v) What is the complexity of Kruskal's in part (iv) if (1) D is list type, and if (2) D is star type. Assume H has n vertices and m edges. \diamond

Exercise 6.9: Give two alternative proofs that the suggested algorithm for computing minimum base is correct:

- (a) By verifying the analogue of the Correctness Lemma.
- (b) By replacing the cost $C(e)$ (for each $e \in E$) by the cost $c_0 - C(e)$. Choose c_0 large enough so that $c_0 - C(e) > 0$. \diamond

Exercise 6.10: Let G be a bigraph G with distinct weights. Give a direct argument for the

- (a) and (b).
- (a) Prove that the MST of G must contain that the edge of smallest weight.
- (b) Prove that the MST of G must contain that the edge of second smallest weight.
- (c) Must it contain the edge of third smallest weight? \diamond

Exercise 6.11: Show that every MST can be obtained from Kruskal's algorithm by a suitable re-ordering of the edges which have identical weights. Conclude that when the edge weights are unique, then the MST is unique. \diamond

Exercise 6.12: Student Joe wants to reduce the minimum base problem for a costed matroid $(S, I; C)$ to the MIS problem for $(S, I; C')$ where C' is a suitable transformation of C . See next section for matroid definitions.

- (a) Student Joe considers the modified cost function $C'(e) = 1/C(e)$ for each e . Construct an example to show that the MIS solution for C' need not be the same as the minimum base solution for C .
- (b) Next, student Joe considers another variation: he now defines $C'(e) = -C(e)$ for each e . Again, provide a counter example. \diamond

Exercise 6.13: Extend the algorithm to finding MIS in contracted matroids. \diamond

Exercise 6.14: If $S \subseteq E$ is Prim-good, then clearly $G' = (V(S), S)$ is clearly a tree. Prove that S is actually an MST of the restricted graph $G|V(S)$. \diamond

Exercise 6.15: Given a costed bigraph $G = (V, E; C)$, let $MST(G)$ denote the set of all MST's of G . Define **MST exchange graph** $Exch(G)$ of G to be the bigraph $(MST(G), H)$ whose nodes are MST's and whose edges $T-T' \in H$ are characterized by this property: the symmetric difference $T \oplus T' = (T \setminus T') \cup (T' \setminus T)$ has size 2.

- (a) Let G_5 be bigraph in Figure ?? . Draw the MST exchange graph $Exch(G_5)$, assuming that $MST(G_5)$ is comprised of the 5 listed MST's.
- (b) Prove that if $T \neq T' \in MST(G)$, then there exists $e \in T \setminus T'$ and $e' \in T' \setminus T$ such that $T + e' - e \in MST(G)$. Conclude that $Exch(G)$ is a connected graph. \diamond

Exercise 6.16: Refer to the graph in Figure ??.

- (a) List all the X -good sets of size 1 for the graph in this figure, where X =Kruska, Boruvka or Prim. For Prim-goodness, assume that node a is the source in Figure 10.
 (b) Repeat part(a) but for sets of size 2. \diamond

Exercise 6.17: (a) Let G_5 be the bigraph in Figure ??. Prove that there are no other MST's other than the 5 which are shown.

- (b) Draw the exchange graph if all the edges in Figure ?? are unit costs? \diamond

Exercise 6.18:

- (a) Enumerate the X -good sets of vertices in Figure 9. Here, X is 'minimally', 'Kruskal', 'Boruvka' or 'Prim'.
 (b) Characterize the good singletons (relative to any of the three notions of goodness). \diamond

Exercise 6.19: This question will develop Boruvka's approach to MST: for each vertex v , pick the edge $(v-u)$ that has the least cost among all the nodes u that are adjacent to v . Let P be the set of edges so picked.

- (a) Show that $n/2 \leq P \leq n-1$. Give general examples to show that these two extreme bounds are achieved for each n .
 (b) Show that if the costs are unique, P cannot contain a cycle. What kinds of cycles can form if weights are not unique?
 (c) Assume edges in P are picked with the tie breaking rule: among the edges $v-u_i$ ($i = 1, 2, \dots$) adjacent to v that have minimum cost, pick the u_i that is the smallest numbered vertex (assume vertices are numbered from 1 to n). Prove that P is acyclic and has the following property: if adding an edge e to P creates a cycle Z in $P+e$, then e has the maximum cost among the edges in Z .
 (d) For any costed bigraph $G = (V, E; C)$, and $P \subseteq E$, define a new costed bigraph denoted G/P as follows. First, two vertices of V are said to be equivalent modulo P if they are connected by a sequence of edges in P . For $v \in V$, let $[v]$ denote the equivalence class of v . The vertex set of G/P is $\{[v] : v \in V\}$. The edge set of G/P comprises those $([u]-[v])$ such that there exists an edge $(u'-v') \in E$ where $u' \in [u]$ and $v' \in [v]$. The cost of $([u]-[v])$ is defined as $\min\{C(u', v') : u' \in [u], v' \in [v], (u'-v') \in E\}$. Note that G/P has at most $n/2$ vertices. Moreover, we can pick another set P' of edges in G/P using the same rules as before. This gives us another graph $(G/P)/P'$ with at most $n/4$ vertices. We can continue this until V has 1 vertex. Please convert this informal description into an algorithm to compute the cost of the MST. (You need not show how to compute the MST.)
 (e) Determine the complexity of your algorithm. You will need to specify suitable data structures for carrying out the operations of the algorithm. (Please use data structures that you know up to this point.) \diamond

Exercise 6.20: (Tarjan) Consider the following **generic accept/reject algorithm** for MST.

- This consists of steps that either **accept** or **reject** edges. In our generic MST algorithm, we only explicitly accept edges. However, we may be implicitly rejecting edges as well, as in the case of Kruskal's algorithm. Let S, R be the sets of accepted and rejected edges (so far). We say that (S, R) is **minimally-good** if there is an MST that contains S but not containing any edge of R . Note that this extends our original definition of "minimally good". Prove that the following extensions of S and R will maintain minimal goodness:
 (a) Let $U \subseteq V$ be any subset of vertices. The set of edges of the form (u, v) where $u \in U$

and $v \notin U$ is called a **U -cut**. If e is the minimum cost edge of a U -cut and there are no accepted edges in the U -cut, then we may extend S by e .

(b) If e is the maximum cost edge in a cycle C and there are no rejected edges in C then we may extend R by e . \diamond

Exercise 6.21: With respect to the generic accept/reject version of MST:

(a) Give a counter example to the following rejection rule: let e and e' be two edges in a U -cut. If $C(e) \geq C(e')$ then we may reject e' .

(b) Can the rule in part (a) be fixed by some additional properties that we can maintain?

(c) Can you make the criterion for rejection in the previous exercise (part (b)) computationally effective? Try to invent the “inverses” of Prim’s and Boruvka’s algorithm in which we solely reject edges.

(d) Is it always a bad idea to *only* reject edges? Suppose that we alternatively accept and reject edges. Is there some situation where this can be a win? \diamond

Exercise 6.22: Consider the following recursive “MST algorithm” on input $G = (V, E; C)$:

(I) Subdivide $V = V_1 \uplus V_2$.

(II) Recursive find a “MST” T_i of $G|V_i$ ($i = 1, 2$).

(III) Find e in the V_1 -cut of minimum cost. Return $T_1 + e + T_2$.

Give a small counter example to this algorithm. Can you fix this algorithm? \diamond

Exercise 6.23: Is there an analogue of Prim and Boruvka’s algorithm for the MIS problem for matroids? \diamond

Exercise 6.24: Let $G = (V, E; C)$ be the complete graph in which each vertex $v \in V$ is a point in the Euclidean plane and $C(u, v)$ is just the Euclidean distance between the points u and v . Give efficient methods to compute the MST for G . \diamond

Exercise 6.25: Fix a connected undirected graph $G = (V, E)$. Let $T \subseteq E$ be any spanning tree of G . A pair (e, e') of edges is called a **swappable pair for T** if

(i) $e \in T$ and $e' \in E \setminus T$ (Notation: for sets A, B , their difference is denoted $A \setminus B = \{a \in A : a \notin B\}$)

(ii) The set $(T \setminus \{e\}) \cup \{e'\}$ is a spanning tree.

Let $T(e, e')$ denote the spanning tree $(T \setminus \{e\}) \cup \{e'\}$ obtained from T by **swapping e and e'** (see illustration in Figure 14(a), (b)).

(a) Suppose (e, e') is a swappable pair for T and $e' = (u, v)$. Prove that e lies on the unique path, denoted by $P(u, v)$, of T from u to v . In Figure 14(a), $e' = (1-5) = (5-1)$. So the path is either $P(1, 5) = (1-2-3-5)$ or $P(5, 1) = (5-3-2-1)$.

(b) Let $n = |V|$. Relative to T , we define a $n \times n$ matrix *First* indexed by pairs of vertices u, v , where $First[u, v] = w$ means that the first edge in the unique path $P(u, v)$ is (u, w) . (In the special case of $u = v$, let $First[u, u] = u$.) In Figure 14(a), $First[1, 5] = 2$ and $First[5, 1] = 3$. Show the matrix *First* for the tree T in Figure 14(a). Similarly, give the matrix *First* for the tree $T(e, e')$ in Figure 14(b).

(c) Describe an $O(n^2)$ algorithm called *Update(First, e, e')* which updates the matrix *First* after we transform T to $T(e, e')$. HINT: For which pair of vertices (x, y) does the value of $First[x, y]$ have to change? Suppose $e' = (u', v')$ and $P(u', v') = (u_0, u_1, \dots, u_\ell)$ is as illustrated in Figure 14(c). Then $u' = u_0, v' = u_\ell$, and also $e = (u_k, u_{k+1})$ for some $0 \leq k < \ell$. Then, originally $First[u_0, u_\ell] = u_1$ but after the swap, $First[u_0, u_\ell] = u_\ell$. What else must change?




Figure 14: (a) A swappable pair (e, e') for spanning tree T . (b) The new spanning tree $T(e, e')$
 [NOTE: tree edges are indicated by thick lines]

(d) Analyze your algorithm to show that that it is $O(n^2)$. Be sure that your description in (c) is clear enough to support this analysis. \diamond

END EXERCISES

§7. Matroids

An abstract structure that supports greedy algorithms is matroids. Indeed, we will see that Kruskal's algorithm for MST is an instance of a general greedy method to solve a matroid problem. We first illustrate the idea of matroids.

¶42. Graphic matroids. Let $G = (V, S)$ be a bigraph. A subset $A \subseteq S$ is **acyclic** if it does not contain any cycle. The I comprising all acyclic subsets of S is called the **graphic matroid** of G . Let us prove 2 critical proper We note two properties of I :

Hereditary property: If $A \subseteq B$ and $B \in I$ then $A \in I$. In particular, the empty set belongs to I .

Exchange property: If $A, B \in I$ and $|A| < |B|$ then there is an edge $e \in B - A$ such that $A \cup \{e\} \in I$.

The hereditary property is obvious. To prove the exchange property, note that the subgraph $G|A := (V, A)$ has $|V| - |A|$ (connected) components; similarly the subgraph $G|B := (V, B)$ has $|V| - |B|$ components. CLAIM: *There exists a component $U \subseteq V$ of $G|B$ that is not contained in any component of $G|A$.* [Pf: If every component $U \subseteq V$ of $G|B$ is contained in some component of $G|A$, then $|V| - |B| < |V| - |A|$ implies that some component of $G|A$ contains no vertices, contradiction.] Let $T := B|U$ be the restriction of B to U . Since U is a component, the graph (U, T) is a spanning tree. Hence there exists an edge $e = (u-v) \in T$ such that u and v belongs to different components of $G|A$. This e will serve for the exchange property: that is, $A + e$ contains no cycle and hence $A + e \in I$.

For example, in Figure 9 the sets $A = \{a-b, a-c, a-d\}$ and $B = \{b-c, c-a, a-d, d-e\}$ are acyclic. Then the exchange property between A and B is witnessed by the edge $d-e \in B \setminus A$, since adding $d-e$ to A will result in an acyclic set.

¶43. **Matroids.** The above system (V, I) is called the **graphic matroid** corresponding to graph $G = (V, S)$. In general, a **matroid** is a hypergraph

$$M = (S, I)$$

with special properties: S and $I \subseteq 2^S$ are both non-empty sets such that I has both the hereditary and exchange properties. The set S is called the **ground set**. Elements of I are called **independent sets**; other subsets of S are called **dependent sets**. Note that the empty set \emptyset is always a member of I .

Another example of matroids arise with numerical matrices: for any matrix M , let S be its set of columns, and I be the family of linearly independent subsets of columns. Call this the **matrix matroid** of M . The terminology of independence comes from this setting. This was the motivation of Whitney, who coined the term ‘matroid’.

The explicit enumeration of the set I is usually out of the question. So, in computational problems whose input is a matroid (S, I) , the matroid is usually implicitly represented. The above examples illustrate this: a graphic matroid is represented by a graph G , and the matrix matroid is represented by a matrix M . The size of the input is then taken to be the size of G or M , not of $|I|$ which can exponentially larger.

¶44. **Submatroids.** Given matroids $M = (S, I)$ and $M' = (S', I')$, we call M' a **submatroid** of M if $S' \subseteq S$ and $I' \subseteq I$. There are two general methods to obtain submatroids, starting from a non-empty subset $R \subseteq S$:

(i) Induced submatroids. The **R -induced submatroid** of M is

$$M|R := (R, I \cap 2^R).$$

(ii) Contracted¹⁵ submatroids. The **R -contracted submatroid** of M is

$$M \wedge R := (R, I \wedge R)$$

where $I \wedge R := \{A \cap R : A \in I, S - R \subseteq A\}$. Thus, there is a bijective correspondence between the independent sets A' of $M \wedge R$ and those independent sets A of M which contain $S - R$. Indeed, $A' = A \cap R$. Of course, if $S - R$ is dependent, then $I \wedge R$ is empty.

We leave it to an exercise to show that $M|R$ and $M \wedge R$ are matroids. Special cases of induced and contracted submatroids arise when $R = S - \{e\}$ for some $e \in S$. In this case, we say that $M|R$ is obtained by **deleting** e and $M \wedge R$ is obtained by **contracting** e .

¶45. **Bases.** Let $M = (S, I)$ be a matroid. If $A \subseteq B$ and $B \in I$ then we call B an **extension** of A ; if $A = B$, the extension is **improper** and otherwise it is **proper**. A **base** of M (alternatively: a **maximal independent set**) is an independent set with no proper extensions. If $A \cup \{e\}$ is independent and $e \notin A$, we call $A \cup \{e\}$ a **simple extension** of A .

¹⁵Contracted submatroids are introduced here for completeness. They are not used in the subsequent development (but the exercises refer to them).

and say that e **extends** A . If $R \subseteq S$, we may relativize these concepts to R : we may speak of “ $A \subseteq R$ being a base of R ”, “ e extends A in R ”, etc. This is the same as viewing A as a set of the induced submatroid $M|_R$.

¶46. **Ranks.** We note a simple property: *all bases of a matroid have the same size*. If A, B are bases and $|A| > |B|$ then there is an $e \in A - B$ such that $B \cup \{e\}$ is a simple extension of B . This is a contradiction. Note that this property is true even if S has infinite cardinality. Thus we may define the **rank** of a matroid M to be the size of its bases. More generally, we may define the rank of any $R \subseteq S$ to be the size of the bases of R (this size is just the rank of $M|_R$). The **rank function**

$$r_M : 2^S \rightarrow \mathbb{N}$$

simply assigns the rank of $R \subseteq S$ to $r_M(R)$.

¶47. **Problems on Matroids.** A **costed matroid** is given by $M = (S, I; C)$ where (S, I) is a matroid and $C : S \rightarrow \mathbb{R}$ is a cost¹⁶ function. The cost of a set $A \subseteq S$ is just the sum $\sum_{x \in A} C(x)$. The **maximum independent set problem** (abbreviated, MIS) is this: given a costed matroid $(S, I; C)$, find an independent set $A \subseteq S$ with maximum cost. A closely related problem is the **maximum base problem** where, given $(S, I; C)$, we want to find a base $B \subseteq S$ of maximum cost. If the costs are non-negative, then it is easy to see the MIS problem and the maximum base problem are identical. The following algorithm solves the maximum base problem:

GREEDY ALGORITHM FOR MAXIMUM BASE:

Input: matroid $M = (S, I; C)$ with cost function C .

Output: a base $A \in I$ with maximum cost.

1. Sort $S = \{x_1, \dots, x_n\}$ by cost.
Suppose $C(x_1) \geq C(x_2) \geq \dots \geq C(x_n)$.
2. Initialize $A \leftarrow \emptyset$.
3. For $i = 1$ to n ,
 put x_i into A provided this does not make A dependent.
4. Return A .

The steps in this abstract algorithm needs to be instantiated for particular representations of matroids. In particular, testing if a set A is independent is usually non-trivial (recall that matroids are usually given implicitly in terms of other combinatorial structures). We discuss this issue for graphic matroids below. It is interesting to note that the usual Gaussian algorithm for computing the rank of a matrix is an instance of this algorithm where the cost $C(x)$ of each element x is unit.

Let us see why the greedy algorithm is correct.

Lemma 17 (Correctness) *Suppose the elements of A are put into A in this order:*

$$z_1, z_2, \dots, z_m,$$

¹⁶Recall our convention that costs may be negative. If the costs are non-negative, we call C a “weight function”.

where $m = |A|$. Let $A_i = \{z_1, z_2, \dots, z_i\}$, $i = 1, \dots, m$. Then:

1. A is a base.

2. If $x \in S$ extends A_i then $i < m$ and $C(x) \leq C(z_{i+1})$.

3. Let $B = \{u_1, \dots, u_k\}$ be an independent set where $C(u_1) \geq C(u_2) \geq \dots \geq C(u_k)$. Then $k \leq m$ and $C(u_i) \leq C(z_i)$ for all i .

Proof. 1. By way of contradiction, suppose $x \in S$ extends A . Then $x \notin A$ and we must have decided not to place x into the set A at some point in the algorithm. That is, for some $j \leq m$, $A_j \cup \{x\}$ is dependent. This contradicts the hereditary property because $A_j \cup \{x\}$ is a subset of the independent set $A \cup \{x\}$.

2. Suppose x extends A_i . By part 1, $i < m$. If $C(x) > C(z_{i+1})$ then for some $j \leq i$, we must have decided not to place x into A_j . This means $A_j \cup \{x\}$ is dependent, which contradicts the hereditary property since $A_j \cup \{x\} \subseteq A_i \cup \{x\}$ and $A_i \cup \{x\}$ is independent.

3. Since all bases are independent sets with the maximum cardinality, we have $k \leq m$. The result is clearly true for $k = 1$ and assume the result holds inductively for $k - 1$. So $C(u_j) \leq C(z_j)$ for $j \leq k - 1$. We only need to show $C(u_k) \leq C(z_k)$. Since $|B| > |A_{k-1}|$, the exchange property says that there is an $x \in B - A_{k-1}$ that extends A_{k-1} . By part 2, $C(z_k) \geq C(x)$. But $C(x) \geq C(u_k)$, since u_k is the lightest element in B by assumption. Thus $C(u_k) \leq C(z_k)$, as desired. **Q.E.D.**

From this lemma, it is not hard to see that an algorithm for the MIS problem is obtained by replacing the for-loop (“for $i = 1$ to n ”) in the above Greedy algorithm by “for $i = 1$ to m ” where x_m is the last positive element in the list $(x_1, \dots, x_m, \dots, x_n)$.

¶48. Greedoids. While the matroid structure allows the Greedy Algorithm to work, it turns out that a more general abstract structure called **greedoids** is tailor-fitted to the greedy approach. To see what this structure looks like, consider the set system (S, F) where S is a non-empty finite set, and $F \subseteq 2^S$. In this context, each $A \in F$ is called a **feasible set**. We call (S, F) a **greedoid** if

Accessibility property If A is a non-empty feasible set, then there is some $e \in A$ such that $A \setminus \{e\}$ is feasible.

Exchange property: If A, B are feasible and $|A| < |B|$ then there is some $e \in B \setminus A$ such that $A \cup \{e\}$ is feasible.

EXERCISES

Exercise 7.1: Consider the graphic matroid in Figure 9. Determine its rank function. \diamond

Exercise 7.2: The text described a modification of the Greedy Maximum Base Algorithm so that it will solve the MIS problem. Verify its correctness. \diamond

Exercise 7.3:

(a) Interpret the induced and contracted submatroids $M|R$ and $M \wedge R$ in the bigraph of Figure 9, for various choices of the edge set R . When is $M|R = M \wedge R$?

(b) Show that $M|R$ and $M \wedge R$ are matroids in general. \diamond

Exercise 7.4: Show that $r_M(A \cup B) + r_M(A \cap B) \leq r_M(A) + r_M(B)$. This is called the **submodularity property** of the rank function. It is the basis of further generalizations of matroid theory. \diamond

Exercise 7.5: In Gavril's activities selection problem, we have a set A of intervals of the form $[s, f)$. Recall that a subset $S \subseteq A$ is said to be compatible if S is pairwise disjoint. Does the set of compatible subsets of A form a matroid? If yes, prove it. If no, give a counter example. \diamond

END EXERCISES

§8. Generating Permutations

In §1, we saw how the general bin packing problem can be reduced to linear bin packing. This reduction depends on the ability to generate all permutations of n elements efficiently. Since there are many uses for such permutation generators, we take a small detour from greedy methods in order to address this interesting topic. A survey of this classic problem is given by Sedgewick [8]. Perhaps the oldest incarnation of this problem is the “change ringing problem” of bell-ringers in early 17th Century English churches [7]. This calls for ringing a sequence of n bells in all $n!$ permutations.

¶49. Combinatorial Enumeration Problems. The problem of generating all permutations efficiently is representative of the important class of **combinatorial enumeration problems**. For instance, we might want to generate all size k subsets of a set, all graphs of size n , all convex polytopes based some given set of n vertices, etc. Such an enumerations would be considered optimal if the algorithm takes $O(1)$ time to generate each member.

It is good to fix some terminology. A **n -permutation** of a finite set X is a surjective function $p : \{1, \dots, n\} \rightarrow X$. Surjectivity of p implies $n \geq |X|$. The function p may be represented by a sequence $(p(1), p(2), \dots, p(n))$. Here we are interested in the case $n = |X|$, i.e., permutation of *distinct* elements. We use a path-like notation for permutations, writing “ $(p(1)-p(2)-\dots-p(n))$ ” for the permutation $(p(1), p(2), \dots, p(n))$.

Example: let $X = \{a, b, c\}$ and $n = 4$. Two 4-permutations of X are $(a-a-b-c)$ and $(b-c-a-b)$. However, $(a-b-c)$ and $(a-b-a-b)$ are not 4-permutations of X .

The set of all n -permutations of $X = \{1, 2, \dots, n\}$ is rather special, and is denoted S_n . Each element of S_n is simply called an **n -permutation** (i.e., X is implicitly $\{1, \dots, n\}$). Note that $|S_n| = n!$. Here is an enumeration of S_3 :

$$(1-2-3), \quad (1-3-2), \quad (3-1-2); \quad (3-2-1), \quad (2-3-1), \quad (2-1-3). \quad (36)$$

Two n -permutations $\pi = (x_1 - \dots - x_n)$ and $\pi' = (x'_1 - \dots - x'_n)$ are **adjacent** (to each other) if there is some $i = 2, \dots, n$ such that

$$x_i = \begin{cases} x'_{i-1} & \text{if } i = j \\ x'_{i+1} & \text{if } i = j - 1 \\ x'_i & \text{else.} \end{cases}$$

We may write $\pi' = \text{Exch}_i(\pi)$ in this case, and call π' an **exchange** of π (and vice-versa). E.g., $\pi = (1-2-4-3)$ and $\pi' = (1-4-2-3)$ are adjacent since $\pi' = \text{Exch}_3(\pi)$. An **adjacency ordering** of a set S of permutations is a listing of the elements of S such that every two consecutive permutations in this listing are adjacent. For instance, the listing of S_3 in (36) is an adjacency ordering.

The **adjacency graph** of S_n is the bigraph with vertex edge S_n and edges given by adjacency permutations. For example, the adjacency graph of S_3 consists of one cycle given by (36): consecutive permutations in (36) are edges of this graph, and also the first $(1-2-3)$ and last $(2-1-3)$ permutations. Since each permutation in S_3 is adjacent to exactly two other permutations, there are no other edges.

We need another concept: if $\pi = (x_1 - \dots - x_{n-1})$ is an $(n-1)$ -permutation, and π' is obtained from π by inserting the letter n into π , then we call π' an **extension** of π . Indeed, if n is inserted just before the i th letter in π , then we write $\pi' = \text{Ext}_i(\pi)$ for $i = 1, \dots, n$. The meaning of “ $\text{Ext}_n(\pi)$ ” should be clear: it is obtained by appending ‘ n ’ to the end of the sequence π . Note that there are n extensions of an $(n-1)$ -permutation. E.g., if $\pi = (1-2)$ then the three extensions of π are $(3-1-2), (1-3-2), (1-2-3)$.

¶50. **The Johnson-Trotter Ordering.** Among the several known methods to generate n -permutations, we will describe one that is independently discovered by S.M. Johnson and H.F. Trotter (1962), and apparently known to 17th Century English bell-ringers [7]. Hugo Steinhaus (1958) describes the problem of generating permutations by n particles moving along a line moving at variable speeds. The two main ideas in the Johnson-Trotter algorithm are (1) the n -permutations are generated as an adjacency ordering, and (2) the n -permutations are generated recursively. Suppose π is an $(n-1)$ -permutation that has been recursively generated. Then we note that the n extensions of π can give one of two adjacency orderings. It is either

$$UP(\pi) : \text{Ext}_1(\pi), \text{Ext}_2(\pi), \dots, \text{Ext}_n(\pi)$$

or the reverse sequence

$$DOWN(\pi) : \text{Ext}_n(\pi), \text{Ext}_{n-1}(\pi), \dots, \text{Ext}_1(\pi).$$

E.g., $UP(1-2-3)$ is equal to

$$(4-1-2-3), (1-4-2-3), (1-2-4-3), (1-2-3-4).$$

Note that if π' is another $(n-1)$ -permutation that is adjacent to π , then the concatenated sequences

$$UP(\pi); DOWN(\pi')$$

and

$$DOWN(\pi); UP(\pi')$$

are both adjacency orderings. We have thus shown:

Lemma 18 (Johnson-Trotter ordering) *If $\pi_1, \dots, \pi_{(n-1)!}$ is an adjacency ordering of S_{n-1} , then the concatenation of alternating DOWN/UP sequences*

$$DOWN(\pi_1); UP(\pi_2); DOWN(\pi_3); \dots; DOWN(\pi_{(n-1)!})$$

is an adjacency ordering of S_n .

For example, starting from the adjacency ordering of 2-permutations ($\pi_1 = (1-2), \pi_2 = (2-1)$), our above lemma says that $DOWN(\pi_1), UP(\pi_2)$ is an adjacency ordering. Indeed, this is the ordering shown in (36).

Let us define the **permutation graph** G_n to be the bigraph whose vertex set is S_n and whose edges comprise those pairs of vertices that are adjacent in the sense defined for permutations. We note that the adjacency ordering produced by Lemma 18 is actually a cycle in the graph G_n . In other words, the adjacency ordering has the additional property that the first and the last permutations of the ordering are themselves adjacent. A cycle that goes through every vertex of a graph is said to be **Hamiltonian**. If $(\pi_1 - \pi_2 - \cdots - \pi_m)$ (for $m = (n-1)!$) is a Hamiltonian cycle for G_{n-1} , then it is easy to see that

$$(DOWN(\pi_1); UP(\pi_2); \cdots; UP(\pi_m))$$

is a Hamiltonian cycle for G_n .

¶51. **The Permutation Generator.** We proceed to derive an efficient means to generate successive permutations in the Johnson-Trotter ordering. We need an appropriate high level view of this generator. The generated permutations are to be used by some “permutation consumer” such as our greedy linear bin packing algorithm. There are two alternative views of the relation between the “permutation generator” and the “permutation consumer”. We may view the consumer as calling¹⁷ the generator repeatedly, where each call to the generator returns the next permutation. Alternatively, we view the generator as a skeleton program with the consumer program as a (shell) subroutine. We prefer the latter view, since this fits the established paradigm of BFS and DFS as skeleton programs (see Chapter 4). Indeed, we may view the permutation generator as a bigraph traversal: the implicit bigraph here is the permutation graph G_n .

In the following, an n -permutation is represented by the array $per[1..n]$. We will transform per by exchange of two adjacent values, indicated by

$$per[i] \Leftrightarrow per[i-1] \tag{37}$$

for some $i = 2, \dots, n$, or

$$per[i] \Leftrightarrow per[i+1]$$

where $i = 1, \dots, n-1$.

¶52. **A Counter for n factorial.** To keep track of the successive exchanges in Johnson-Trotter generator, we introduce an array of n counters

$$C[1..n]$$

where each $C[i]$ is initialized to 1 but always satisfying the relation $1 \leq C[i] \leq i$. Of course, $C[1]$ may be omitted since its value cannot change under our restrictions. The array counter C has $n!$ distinct values. We say the i -th counter is **full** iff $C[i] = i$. The **level** of the C is the largest index ℓ such that the ℓ -th counter is not full. If all the counters are full, the level of C is defined to be 1. E.g., $C[1..5] = [1, 2, 2, 1, 5]$ has level 4. We define the **increment** of this counter array as follows: if the level of the counter is ℓ , then (1) we increment $C[\ell]$ provided $\ell > 1$, and (2) we set $C[i] = 1$ for all $i > \ell$. E.g., the increment of $C[1..5] = [1, 2, 2, 1, 5]$ gives $[1, 2, 2, 2, 1]$. In code:

¹⁷The generator in this viewpoint is a **co-routine**. It has to remember its state from the previous call.

```

INC(C)
   $\ell \leftarrow n.$ 
  While ( $C[\ell] = \ell \wedge (\ell > 1)$ )
     $C[\ell--] \leftarrow 1.$ 
  If ( $\ell > 1$ )
     $C[\ell]++.$ 
  Return( $\ell$ )

```

Note that INC returns the level of the original counter value. This macro is a generalization of the usual increment of binary counters (Chapter 6.1). For instance, for $n = 4$, starting with the initial value of $[1, 1, 1]$, successive increments of this array produce the following cyclic sequence:

$$\begin{aligned}
 C[2, 3, 4] &= [1, 1, 1] \rightarrow [1, 1, 2] \rightarrow [1, 1, 3] \rightarrow [1, 1, 4] \rightarrow [1, 2, 1] \\
 &\rightarrow [1, 2, 2] \rightarrow [1, 2, 3] \rightarrow [1, 2, 4] \rightarrow [1, 3, 1] \rightarrow \cdots \\
 &\rightarrow [2, 3, 3] \rightarrow [2, 3, 4] \rightarrow [1, 1, 1] \rightarrow \cdots
 \end{aligned} \tag{38}$$

Let the cost of incrementing the counter array be equal to $n+1-\ell$ where ℓ is the level. CLAIM: the cost to increment the counter array from $[1, 1, \dots, 1]$ to $[2, 3, \dots, n]$ is $< 2(n!)$. In proof, note that $C[\ell]$ is updated after every $n!/\ell!$ steps, so that the overall, $C[\ell]$ is updated $\ell!$ times. Hence the total number of updates for the $n-1$ counters is

$$n! + (n-1)! + \cdots + 2! < 2(n!),$$

which proves our Claim.

This gives us the top level structure for our permutation generator:

```

JOHNSON-TROTTER GENERATOR (SKETCH)
Input: natural number  $n \geq 2$ 
  ▷ Initialization
     $per[1..n] \leftarrow [1, 2, \dots, n].$     ◁ Initial permutation
     $C[2..n] \leftarrow [1, 1, \dots, 1].$     ◁ Initial counter value
  ▷ Main Loop
    do
       $\ell \leftarrow Inc(C)$ 
      UPDATE( $\ell$ )    ◁ The permutation is updated
      CONSUME( $per$ )  ◁ Permutation is consumed
    While ( $\ell > 1$ )

```

The shell macro CONSUME is application-dependent. As default, we simply use it to print the current permutation.

¶53. **How to update the permutation.** We now describe the UPDATE macro. It uses the previous counter level ℓ to transform the current permutation to the next permutation. For example, the successive counter values in (38) correspond to the following sequence of

permutations:

$$\begin{aligned}
 & \xrightarrow{[1,1,1]} (1-2-3-4) \xrightarrow{[1,1,2]} (1-2-4-3) \xrightarrow{[1,1,3]} (1-4-2-3) \xrightarrow{[1,1,4]} (4-1-2-3) \xrightarrow{[1,2,1]} (4-1-3-2) \quad (39) \\
 & \xrightarrow{[1,2,2]} (1-4-3-2) \xrightarrow{[1,2,3]} (1-3-4-2) \xrightarrow{[1,2,4]} (1-3-2-4) \xrightarrow{[1,3,1]} (3-1-2-4) \xrightarrow{[1,3,2]} \dots \\
 & \xrightarrow{[2,3,3]} (1-4-2-3) \xrightarrow{[2,3,4]} (1-2-4-3) \xrightarrow{[1,1,1]} (1-2-3-4) \rightarrow \dots
 \end{aligned}$$

To interpret the above, consider a general step of the form

$$\dots \xrightarrow{[c_2, c_3, c_4]} (x_1 - x_2 - x_3 - x_4) \xrightarrow{[c'_2, c'_3, c'_4]} (x'_1 - x'_2 - x'_3 - x'_4) \dots$$

We start with the counter value $[c_2, c_3, c_4]$ and permutation $(x_1 - x_2 - x_3 - x_4)$. After calling **Inc**, the counter is updated to $[c'_2, c'_3, c'_4]$, and it returns the level ℓ of $[c_2, c_3, c_4]$. If $\ell = 1$, we may terminate; otherwise, $\ell \in \{2, 3, 4\}$. We find the index i such that $x_i = \ell$ (for some $i = 1, 2, 3, 4$). UPDATE will then exchange x_i with its neighbor x_{i+1} or x_{i-1} . The resulting permutation is $(x'_1 - x'_2 - x'_3 - x'_4)$.

In (39), we indicate x_i by an underscore, “ \underline{x}_i ”. The choice of which neighbor (x_{i-1} or x_{i+1}) depends on whether we are in the “UP” phase or “DOWN” phase of level ℓ . Let $UP[1..n]$ be a Boolean array where $UP[\ell]$ is true in the UP phase, and false in the DOWN phase when we are incrementing a counter at level ℓ . Moreover, the value of $UP[\ell]$ is changed (flipped) each time $C[\ell]$ is reinitialized to 1. For instance, in the first row of (39), $UP[4] = \text{false}$ and so the entry 4 is moving down with each swap involving 4. In the next row, $UP[4] = \text{true}$ and so the entry 4 is moving up with each swap.

Hence we modify our previous INC macro to include this update:

```

INCREMENT(C)
Output: Increments C, updates UP, and returns the previous level of C
   $\ell \leftarrow n$ .
  While ( $C[\ell] = \ell$ )  $\wedge$  ( $\ell > 1$ )     $\triangleleft$  Loop to find the counter level
     $C[\ell] \leftarrow 1$ ;
     $UP[\ell] \leftarrow \neg UP[\ell]$ ;     $\triangleleft$  Flips the boolean value  $UP[\ell]$ 
     $\ell--$ .
  If ( $\ell > 1$ )
     $C[\ell]++$ .
  Return( $\ell$ ).

```

For a given level ℓ , the UPDATE macro need to find the “position” i where $per[i] = \ell$ ($i = 1, \dots, n$). We could search for this position in $O(n)$ time, but it is more efficient to maintain this information directly: let $pos[\ell]$ denote the current position of ℓ . Thus the $pos[1..n]$ is just the inverse of the array $per[1..n]$ in the sense that

$$per[pos[\ell]] = \ell \quad (\ell = 1, \dots, n).$$

We can now specify the UPDATE macro to update both pos and per :

¹⁸In case we want to continue, the case $\ell = 1$ is treated as if $\ell = n$. E.g., in (39), the case $\ell = 1$ is treated as $\ell = 4$.

```

UPDATE( $\ell$ )
  if ( $UP[\ell]$ )
     $per[pos[\ell]] \leftrightarrow per[pos[\ell] + 1];$      $\triangleleft$  modify permutation
     $pos[per[pos[\ell]]] \leftarrow pos[\ell];$      $\triangleleft$  update position array
     $pos[\ell]++;$      $\triangleleft$  update position array
  else
     $per[pos[\ell]] \leftrightarrow per[pos[\ell] - 1];$ 
     $pos[per[pos[\ell]]] \leftarrow pos[\ell];$ 
     $pos[\ell]--;$ 

```

Thus, the final algorithm is:

```

JOHNSON-TROTTER GENERATOR
Input: natural number  $n \geq 2$ 
   $\triangleright$  Initialization
     $per[1..n] \leftarrow [1, 2, \dots, n].$      $\triangleleft$  Initial permutation
     $pos[1..n] \leftarrow [1, 2, \dots, n].$      $\triangleleft$  Initial positions
     $C[2..n] \leftarrow [1, 1, \dots, 1].$      $\triangleleft$  Initial counter value
   $\triangleright$  Main Loop
    do
       $\ell \leftarrow Increment(C);$ 
      UPDATE( $\ell$ );     $\triangleleft$  The permutation is updated
      CONSUME( $per$ );     $\triangleleft$  Permutation is consumed
    While( $\ell > 1$ )

```

The Java code for the Johnson-Trotter Algorithm is presented in an appendix of this chapter.

Remarks:

1. We can introduce early termination criteria into our permutation generator. For instance, in the bin packing application, there is a trivial lower bound on the number of bins, namely $b_0 = \lceil (\sum_{i=1}^n w_i) / M \rceil$. We can stop when we found a solution with b_0 bins. If we want only an approximate optimal, say within a factor of 2, we may exit when we achieve $\leq 2b_0$ bins.
2. We have focused on permutations of distinct objects. In case the objects may be identical, more efficient techniques may be devised. For more information about permutation generation, see the book of Paige and Wilson [6]. Knuth's much anticipated 4th volume will treat permutations; this will no doubt become a principle reference for the subject.
3. The Johnson-Trotter enumeration of permutations can be extended to the enumeration of all strings of length n over a finite alphabet Σ . Now, we define two strings to be **adjacent** if they differ at only one position. For $\Sigma = \{0, 1\}$, the most famous such enumerations is the Gray code (or reflected binary code): if $G(n)$ is the Gray code for strings of length n , then $G(n+1) = 0 \cdot G(n); 1 \cdot G^R(n)$ where $G^R(n)$ is the reverse (reflected) enumeration. E.g., $G(2) = (00, 01, 11, 10)$ and $G(3) = 0 \cdot (00, 01, 11, 10); 1 \cdot (10, 11, 01, 00)(000, 001, 011, 010; 110, 111, 101, 100)$. An application is physical implementation of counters. To count to 2^n , we want to cycle through all the binary strings of length n . Using the standard binary representation, a counter increment may cause more than one bit to change. If each bit is controlled by some physical switch, we might produce intermediate states due to the switches not changing simultaneously. This problem does not arise if we use an enumerations like the Gray code.

—EXERCISES

Exercise 8.1:

- (a) Draw the adjacency bigraph corresponding to 4-permutations. HINT: first draw the adjacency graph for 3-permutations and view 4-permutations as extension of 3-permutations.
- (b) How many edges are there in the adjacency bigraph of n -permutations?
- (c) What is the radius and diameter of the bigraph in part (b)? [See definition of radius and diameter in Exercise 4.8 (Chapter 4).] \diamond

Exercise 8.2: Another way to list all the n -permutations in S_n is lexicographic ordering: $(x_1 - \cdots - x_n) < (x'_1 - \cdots - x'_n)$ if the first index i such that $x_i \neq x'_i$ satisfies $x_i < x'_i$. Thus the lexicographic smallest n -permutation is $(1-2-\cdots-n)$. Give an algorithm to generate n -permutations in lexicographic ordering. Compare this algorithm to the Johnson-Trotter algorithm. \diamond

Exercise 8.3: All adjacency orderings of 2- and 3-permutations are cyclic. Is it true of 4-permutations? \diamond

Exercise 8.4: Two n -permutations π, π' are **cyclic equivalent** if $\pi = (x_1 - x_2 - \cdots - x_n)$ and $\pi' = (x_i - x_{i+1} - \cdots - x_n - x_1 - x_2 - \cdots - x_{i-1})$ for some $i = 1, \dots, n$. A **cyclic n -permutation** is an equivalence class of the cyclic equivalence relation. Note that there are exactly n permutations in each cyclic n -permutation. Let S'_n denote the set of cyclic n -permutations. So $|S'_n| = (n-1)!$. Again, we can define the cyclic permutation graph G'_n whose vertex set is S'_n , and edges determined by adjacent pairs of cyclic permutations. Give an efficient algorithm to generate a Hamiltonian cycle of G'_n . \diamond

Exercise 8.5: Suppose you are given a set S of n points in the plane. Give an efficient method to generate all the convex polygons whose vertices are from S . Give the complexity of your algorithm as a function of n . \diamond

END EXERCISES

§9 APPENDIX: Java Code for Permutations

```

2214
2215 /*****
2216  * Per(mutations)
2217  *     This generates the Johnson-Trotter permutation order.
2218  *     By n-permutation, we mean a permutation of the symbols {1,2,...,n}.
2219  *
2220  * Usage:
2221  *     % javac Per.java
2222  *     % java Per [n=3] [m=0]
2223  *
2224  *     will print all n-permutations. Default values n=3 and m=0.
2225  *     If m=1, output in verbose mode.
2226  *     Thus "java Per" will print
2227  *         (1,2,3), (1,3,2), (3,1,2), (3,2,1), (2,3,1), (2,1,3).
2228  *     See Lecture Notes for details of this algorithm.
2229  *
2230  *****/
2231
2232 public class Per {
2233
2234     // Global variables
2235     //////////////////////////////////////
2236     static int n;                // n-permutations are being considered
2237                                // Quirk: Following arrays are indexed from 1 to n
2238     static int[] per;            // represents the current n-permutation
2239     static int[] pos;            // inverse of per: per[pos[i]]=i (for i=1..n)
2240     static int[] C;              // Counter array: 1 <= C[i] <= i (for i=1..n)
2241     static boolean[] UP;         // UP[i]=true iff pos[i] is increasing
2242                                //      (going up) in the current phase
2243
2244     // Display permutation or position arrays
2245     //////////////////////////////////////
2246     static void showArray(int[] myArray, String message){
2247         System.out.print(message);
2248         System.out.print("(" + myArray[1]);
2249         for (int i=2; i<=n; i++)
2250             System.out.print(", " + myArray[i]);
2251         System.out.println(")");
2252     }
2253
2254     // Print counter
2255     //////////////////////////////////////
2256     static void showC(String m){
2257         System.out.print(m);
2258         System.out.print("(" + C[2]);
2259         for (int i=3; i<=n; i++)
2260             System.out.print(", " + C[i]);
2261         System.out.println(")");
2262     }
2263
2264     // Increment counter
2265     //////////////////////////////////////
2266     static int inc(){

```



```

2267     int ell=n;
2268     while ((C[ell]==ell) && (ell>1)){
2269         UP[ell] = !(UP[ell]);    // flip Boolean flag
2270         C[ell--]=1;
2271     }
2272     if (ell>1)
2273         C[ell]++;
2274     return ell;                // level of previous counter value
2275 }
2276
2277 // Update per and pos arrays
2278 ///////////////////////////////////////////////////
2279 static void update(int ell){
2280     int tmpSymbol;            // this is not necessary, but for clarity
2281     if (UP[ell]) {
2282         tmpSymbol = per[pos[ell]+1]; // Assert: pos[ell]+1 makes sense!
2283         per[pos[ell]] = tmpSymbol;
2284         per[pos[ell]+1] = ell;
2285         pos[ell]++;
2286         pos[tmpSymbol]--;
2287     } else {
2288         tmpSymbol = per[pos[ell]-1]; // Assert: pos[ell]-1 makes sense!
2289         per[pos[ell]] = tmpSymbol;
2290         per[pos[ell]-1] = ell;
2291         pos[ell]--;
2292         pos[tmpSymbol]++;
2293     }
2294 }
2295
2296 // Main program
2297 ///////////////////////////////////////////////////
2298 public static void main (String[] args)
2299     throws java.io.IOException
2300 {
2301     //Command line Processing
2302     n=3;                // default value of n
2303     boolean verbose=false; // default is false (corresponds to second argument = 0)
2304     if (args.length>0)
2305         n = Integer.parseInt(args[0]);
2306     if ((args.length>1) && (Integer.parseInt(args[1]) != 0))
2307         verbose = true;
2308
2309     //Initialize
2310     per = new int[n+1];
2311     pos = new int[n+1];
2312     C   = new int[n+1];
2313     UP  = new boolean[n+1];
2314     for (int i=0; i<=n; i++) {
2315         per[i]=i;
2316         pos[i]=i;
2317         C[i]=1;
2318         UP[i]=false;
2319     }
2320
2321     //Setup For Loop
2322     int count=0;        // only used in verbose mode
2323     int ell=1;
2324     System.out.println("Johnson-Trotter ordering of "+ n + "-permutations");
2325     if (verbose)

```

```

2326         showArray(per, count + ", level="+ ell + " :\t" );
2327     else
2328         showArray(per, "");
2329
2330     //Main Loop
2331     do {
2332         ell = inc();
2333         update(ell);
2334         if (verbose)
2335             count++;
2336         showArray(per, count + ", level="+ ell + " :\t" );
2337     else
2338         showArray(per, "");
2339     } while (ell>1);
2340
2341 }//main
2342 }//class Per

```

References

- 2344 [1] J. L. Bentley, D. D. Sleator, R. E. Tarjan, and V. K. Wei. A locally adaptive data
2345 compression scheme. *Comm. of the ACM*, 29(4):320–330, 1986.
- 2346 [2] X. Cai and Y. Zheng. Canonical coin systems for change-making problems.
2347 *arXiv:0809.0400v1 [cs.DS]*, 2009. 14 pages.
- 2348 [3] S. K. Chang and A. Gill. Algorithmic solution of the change-making problem. *J. ACM*,
2349 17(1):113–122, 1970.
- 2350 [4] J. Ziv and A. Lempel. Compression of individual sequences via variable-rate coding. *IEEE*
2351 *Trans. Inform. Theory*, IT-24:530–536, 1978.
- 2352 [5] D. Kozen and S. Zaks. Optimal bounds for the change-making problem. *Theor. Computer*
2353 *Science*, 123(2):377–388, 1994.
- 2354 [6] E. Page and L. Wilson. *An Introduction to Computational Combinatorics*. Cambridge
2355 Computer Science Texts, No. 9. Cambridge University Press, 1979.
- 2356 [7] T. W. Parsons. Letter: A forgotten generation of permutations, 1977.
- 2357 [8] R. Sedgewick. Permutation generation methods. *Computing Surveys*, 9(2):137–164, 1977.
- 2358 [9] R. E. Tarjan. *Data Structures and Network Algorithms*. SIAM, Philadelphia, PA, 1974.
- 2359 [10] J. S. Vitter. The design and analysis of dynamic huffman codes. *J. ACM*, 34(4):825–845,
2360 1987.