

## Parallel Computing Lab Assignment 1

In this lab you will write MPI code to find numbers divisible by a given number  $x$  in the range between  $A$  and  $B$  (inclusive) and test scalability and performance. For example: find numbers divisible by 3 in the range between 2 and 10000,

### General notes:

- The name of the source code file is: `netID.c` where `netID` is your NetID.
- You compile with `mpicc -std=c99 -Wall -o checkdiv netID.c`
- To execute it: `mpiexec -n p ./checkdiv A B x`  
Where  $A$  and  $B$  are positive numbers bigger than 1 and less than or equal to 100,000,000 (100 millions); and  $x$  is a positive number between 2 and 1000,000 (inclusive).
- The output of your program is a text file `B.txt`  
For example, if I type: `mpiexec -n 4 ./checkdiv 2 10 3`  
The output must be a text file with the name `10.txt` and that file contains:  
3  
6  
9  
one number per line.
- **The numbers must be written in the file in ascending order and one number per line** like this example above.
- Your program must also print on screen:  
times of part1 = num1 s  
times part2 = num2 s  
 $\text{num1}$  and  $\text{num2}$  are double precision floating points representing the time taken in seconds. To know what are `part1`, etc, continue reading. They are defined next page.
- You can assume that we will not do any tricks with the input (i.e. We will not deliberately test your program with wrong values of  $N$ , negative, float, non-numeric, etc).
- We are providing you with a file: `skeleton_lab1.c` that can help you start. You can use it and just fill-in the gaps.

### The parallel code:

Assume you have  $p$  processes.

- All processes will get the arguments of the `main()` function right away, without the need to communicate them.
- Part 1:
  - The range  $A$  to  $B$  (i.e.  $B-A+1$  numbers) will be divided among the processes with the last process (in terms of rank) taking slightly extra/less work if the range is not divisible by the range. So, if  $A = 2$  and  $B = 10$  and we have two processes, process 0 will work on range 2 to 6 and process 1 will work on the range 7 to 10, getting one less item than proc.
  - Then, each process works on its range and generates its own list.
  - Each process sends its list to process 0.

- Part 2:
  - Finally, process zero creates the file N.txt and writes all the numbers there.

## How to measure the performance and scalability of your code?

To see how efficient your implementation is, you need to compare against a version with one process. Therefore, we will use several methods.

The program is divided into two main parts, as we saw above.

**part 1** where each process generates its list of divisible numbers and sends it to process 0.

**part 2** where process 0 write the data to the disk.

To measure the time of a piece of code, we will use `MPI_Wtime()` as follows (Note: The following is a pseudo-code. You need to write proper one with correct declarations and header files):

```
start1 = MPI_Wtime();
part1
end1 = MPI_Wtime();
reduction operation, MAX to get largest end1-start1
start2 = MPI_Wtime();
part2
end2 = MPI_Wtime();
```

When executing your full program, use the Linux *time* command. It will be used in graph 2 as will be shown in the section “The report” below. For example

```
time mpirun -n 4 ./checdiv 1000000 5
```

Your code must show some speedup, in part 2, relative to running with only one process. Finally, you will report these numbers in the report as discussed below.

## The report

For that report, to generate the graphs, assume  $x = 91$ .

We may test your program with several different numbers though.

Write a report that contains the following graphs.

#### Graph1:

- X-axis with values (100, 1000, 10000, 100000, 1 million, 10 millions, and 100 millions). These values represent B. Assume  $A = 2$ .
- The speedup (y-axis) (**time of part 1 with 1 process / time part 1 with p process**).
- For each number in the x-axis, draw three bars: speedup using two processes, four processes, and eight processes.

#### Graph 2:

- Same values in x-axis and same bars (2, 4, and 8 processes) as graph 1 but the y-axis is the overall speedup generated by the *time* command as described above. That command generates three numbers: user, system, and real. We want the *real* one. The y-axis will then be the real time for one process divided by the real time of p process where  $p = 2, 4, \text{ and } 8$ .

#### **What to submit:**

A single zip file. The file name is your **netID.zip** where netID is your, well, NetID. Inside that zip file you need to have:

- **netID.c**
- pdf file containing the two graphs the file name must be **netID\_report.pdf**

Submit the zip file through Brightspace.

**Enjoy!**