



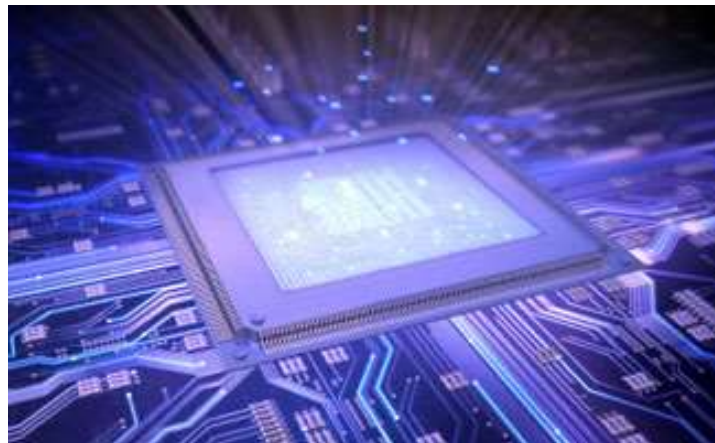
Parallel Computing

CUDA - I

Mohamed Zahran (aka Z)

mzahran@cs.nyu.edu

<http://www.mzahran.com>



GPU Computing Applications						
Libraries and Middleware						
cuDNN TensorRT	cuFFT, cuBLAS, cuRAND, cuSPARSE	CULA MAGMA	Thrust NPP	VSIPL, SVM, OpenCurrent	PhysX, OptiX, iRay	MATLAB Mathematica
Programming Languages						
C	C++	Fortran	Java, Python, Wrappers	DirectCompute	Directives (e.g., OpenACC)	
CUDA-enabled NVIDIA GPUs						
Turing Architecture (Compute capabilities 7.x)	DRIVE/JETSON AGX Xavier	GeForce 2000 Series		Quadro RTX Series	Tesla T Series	
Volta Architecture (Compute capabilities 7.x)	DRIVE/JETSON AGX Xavier				Tesla V Series	
Pascal Architecture (Compute capabilities 6.x)	Tegra X2	GeForce 1000 Series		Quadro P Series	Tesla P Series	
Maxwell Architecture (Compute capabilities 5.x)	Tegra X1	GeForce 900 Series		Quadro M Series	Tesla M Series	
Kepler Architecture (Compute capabilities 3.x)	Tegra K1	GeForce 700 Series GeForce 600 Series		Quadro K Series	Tesla K Series	
	EMBEDDED	CONSUMER DESKTOP, LAPTOP		PROFESSIONAL WORKSTATION	DATA CENTER	

Source: CUDA Toolkit Documentation by NVIDIA

Parallel Computing on a GPU

- GPUs deliver up to 13,800+ GFLOPS (FP32)
 - Available in laptops, clusters, etc.
- GPU parallelism is doubling almost every year
- Programming model scales transparently
 - **Data parallelism**
- Programmable in C/C++ (and other languages) with CUDA tools
- Multithreaded SPMD model uses application data parallelism and thread parallelism.
[**SPMD** = Single Program Multiple Data]



GeForce RTX 3080

Compute Capability

[illegible]

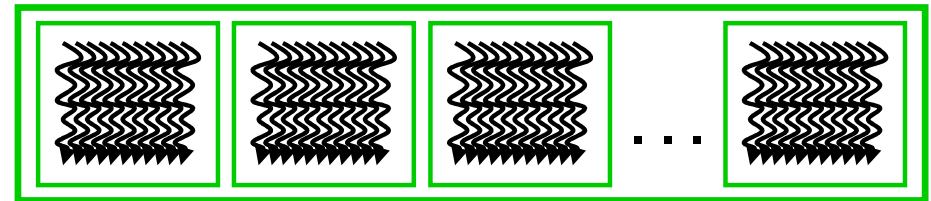
CUDA

- **Compute Unified Device Architecture**
- Integrated host+device app C program
 - Serial or modestly parallel parts in **host C code**
 - Highly parallel parts in **device SPMD kernel C code**

Serial Code (host)

Parallel **Kernel** (device)

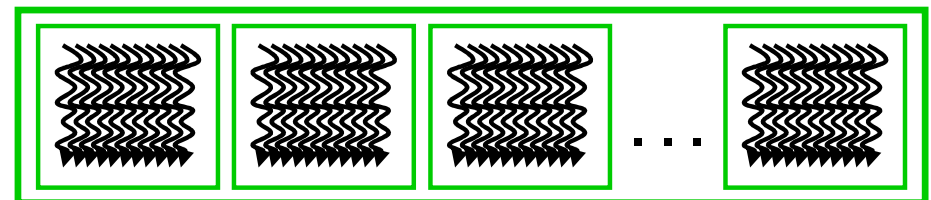
KernelA<<< nBlk, nTid >>>(args);



Serial Code (host)

Parallel **Kernel** (device)

KernelB<<< nBlk, nTid >>>(args);



```

1  #include <stdio.h>
2  #include <cuda.h>    /* Header file for CUDA */
3
4  /* Device code: runs on GPU */
5  __global__ void Hello(void) {
6
7      printf("Hello from thread %d!\n", threadIdx.x);
8  } /* Hello */
9
10
11 /* Host code: Runs on CPU */
12 int main(int argc, char* argv[]) {
13     int thread_count;    /* Number of threads to run on GPU */
14
15     thread_count = strtol(argv[1], NULL, 10);
16                     /* Get thread_count from command line */
17
18     Hello <<<1, thread_count>>>();
19                     /* Start thread_count threads on GPU, */
20
21     cudaDeviceSynchronize();    /* Wait for GPU to finish */
22
23     return 0;
24 } /* main */

```

```

1 #include <stdio.h>
2 #include <cuda.h>    /* Header file for CUDA */
3
4 /* Device code: runs on GPU */
5 __global__ void Hello(void) {
6
7     printf("Hello from thread %d!\n", threadIdx.x);
8 } /* Hello */
9
10
11 /* Host code: Runs on CPU */
12 int main(int argc, char* argv[]) {
13     int thread_count;    /* Number of threads to run on GPU */
14
15     thread_count = strtol(argv[1], NULL, 10);
16                     /* Get thread_count from command line */
17
18     Hello <<<1, thread_count>>>();
19                     /* Start thread_count threads on GPU, */
20
21     cudaDeviceSynchronize();    /* Wait for GPU to finish */
22
23     return 0;
24 } /* main */

```

\$./cuda_hello 10

and the output of will be

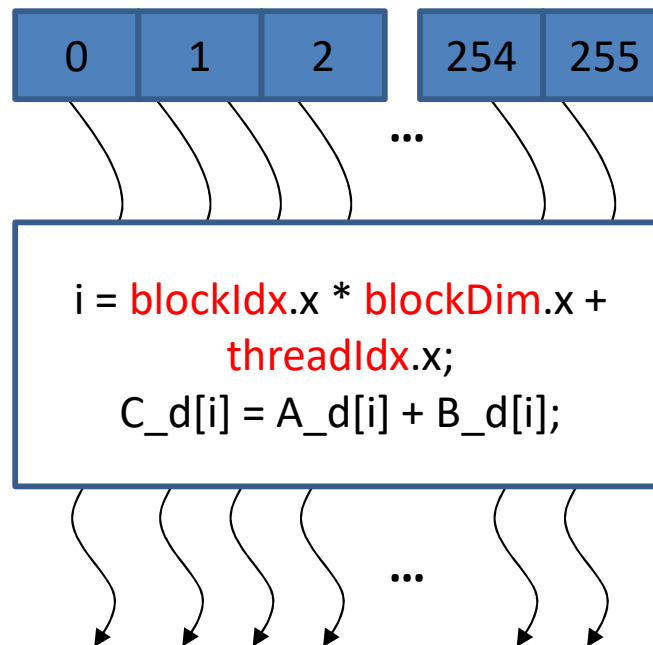
```

Hello from thread 0!
Hello from thread 1!
Hello from thread 2!
Hello from thread 3!
Hello from thread 4!
Hello from thread 5!
Hello from thread 6!
Hello from thread 7!
Hello from thread 8!
Hello from thread 9!

```

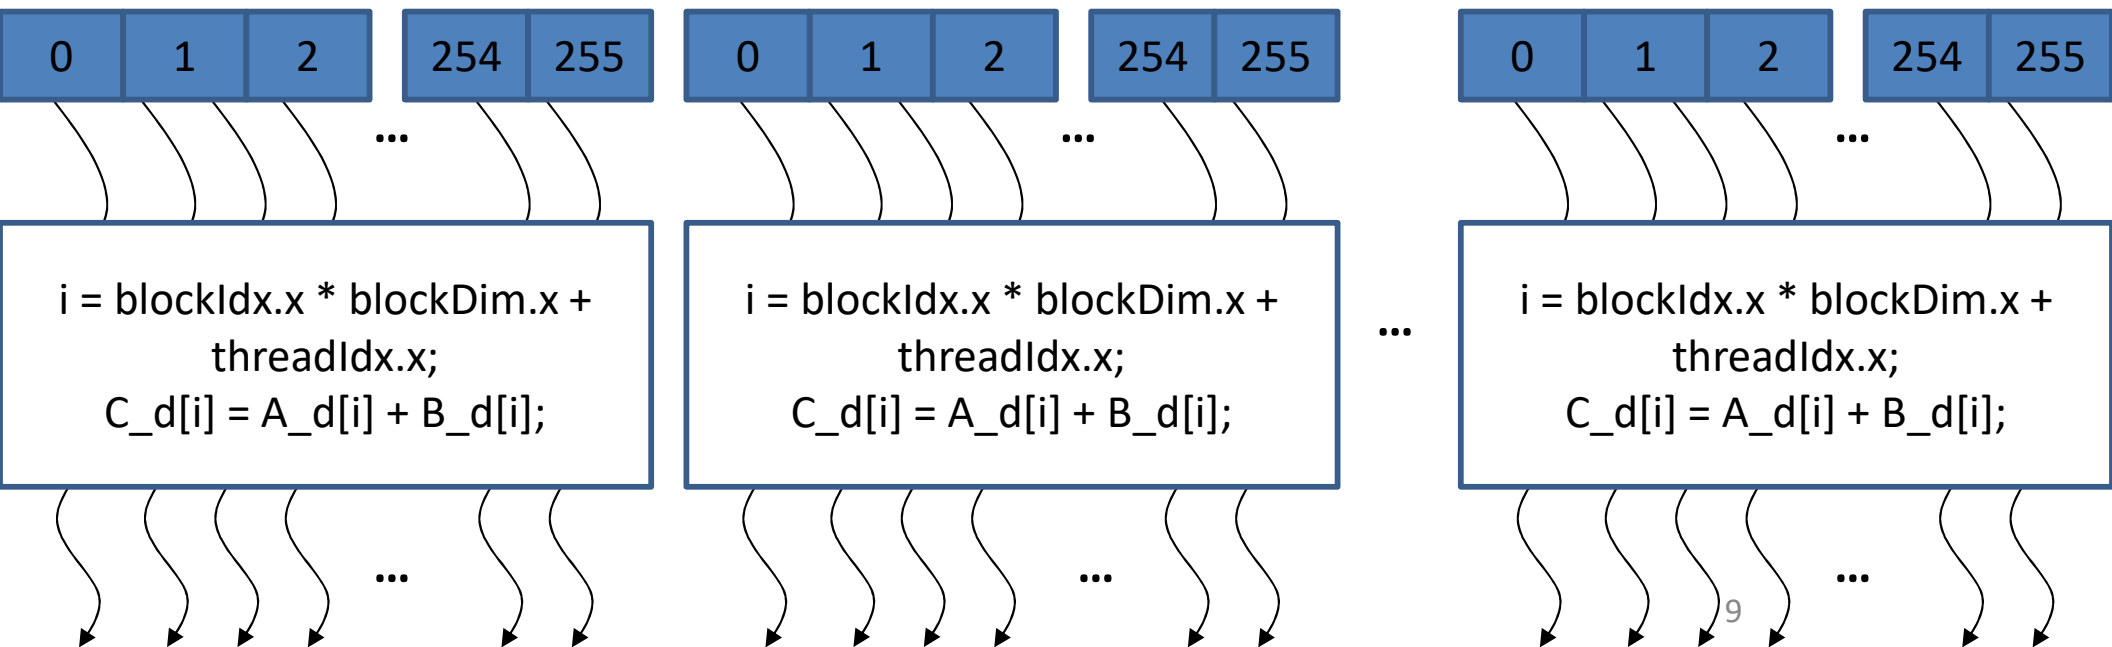
Parallel Threads

- A CUDA kernel is executed by an array of threads
 - All threads run the same code (the SP in SPMD)
 - Each thread has an ID that it uses to compute memory addresses and make control decisions



Thread Blocks

- Divide monolithic thread array into multiple blocks
 - Threads within a block cooperate via **shared memory, atomic operations** and **barrier synchronization, ...**
 - Threads in different blocks cannot cooperate



Kernel

- Launched by the host
- Very similar to a C function
- To be executed on device
- All threads will execute that same code in the kernel.



Grid

- 1D, 2D, or 3D organization of a Grid
- `gridDim.x`, `gridDim.y`, `gridDim.z` are the size of the grid in number of blocks



Block

- 1D, 2D, or 3D organization of a block
- Block is assigned to an SM
- `blockDim.x`, `blockDim.y`, `blockDim.z` are block dimensions counted as number of threads
- `blockIdx.x`, `blockIdx.y`, `blockIdx.z` are indices of the block within a GRID.



Thread

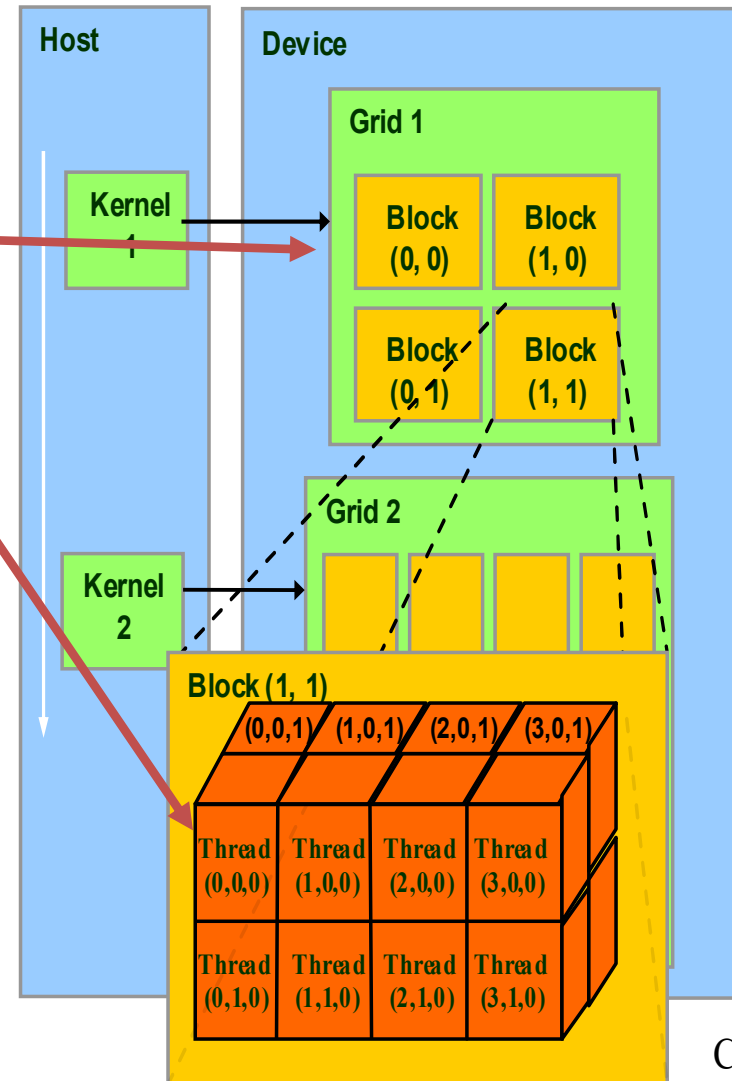
- `threadIdx.x`, `threadIdx.y`, `threadIdx.z` are the index *within* a block

Decisions you have to make as a GPU programmers:

1. Which part(s) of the program will be executed on the GPU?
2. How many total threads will you spawn?
3. How many blocks? That is: how many threads per block?
4. What will be the geometry of the block (1D, 2D, or 3D)?
5. What will be the geometry of the grid (1D, 2D, or 3D)?

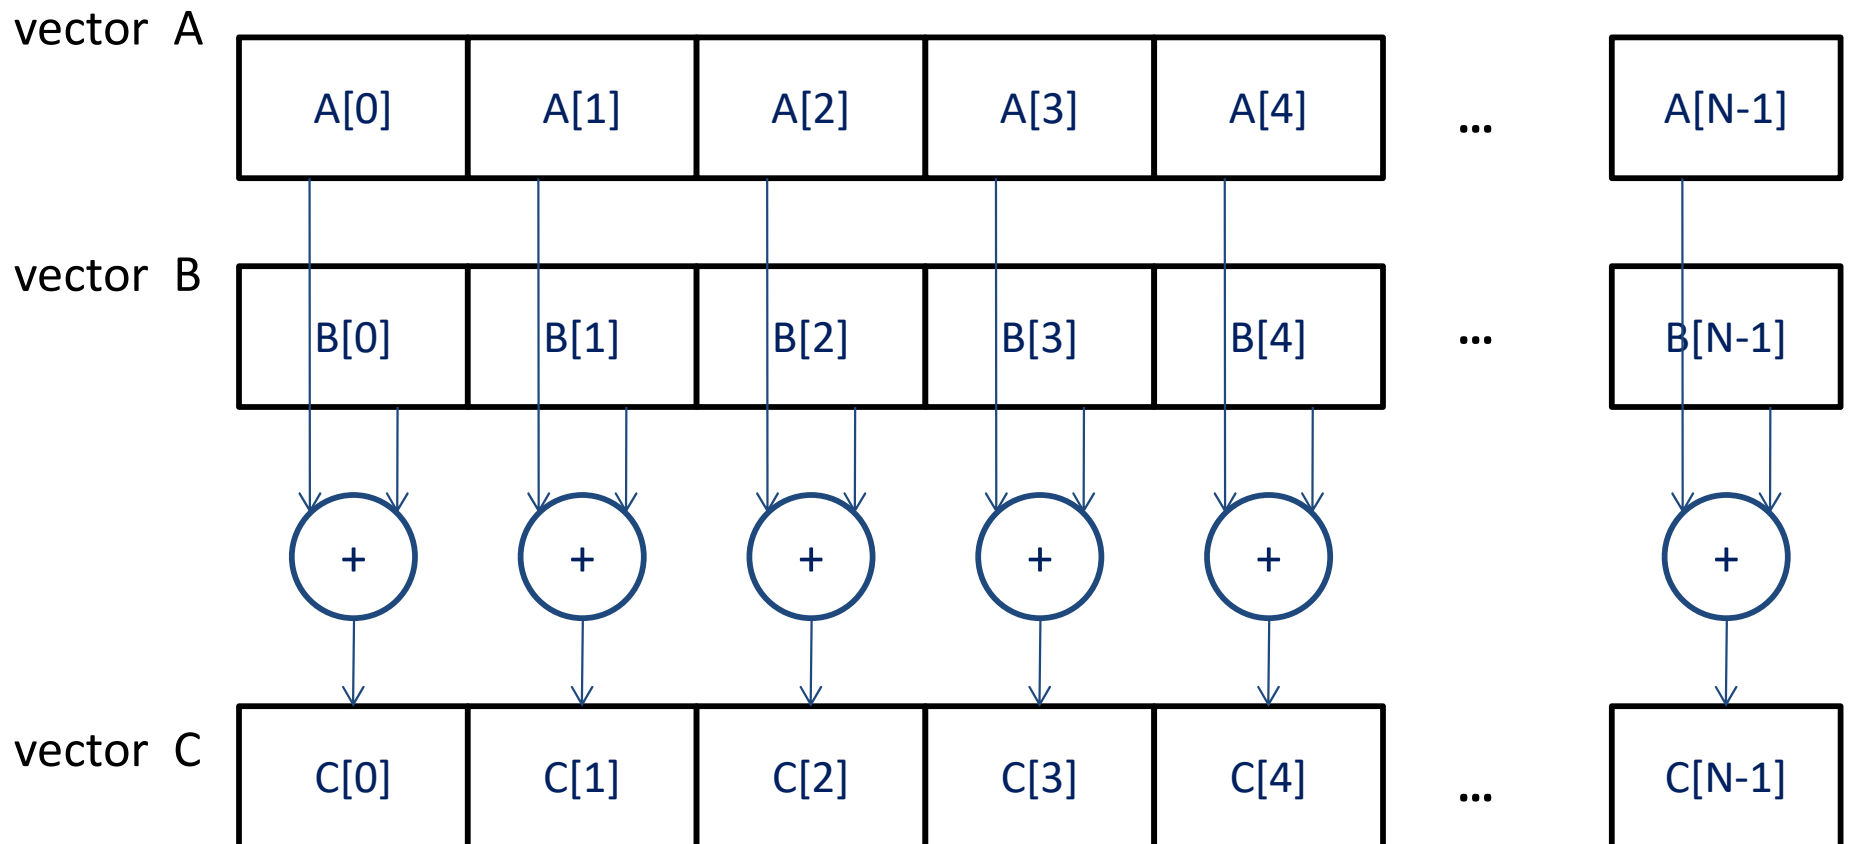
IDs

- Each thread uses IDs to decide what data to work on
 - Block ID: 1D or 2D or 3D
 - Thread ID: 1D, 2D, or 3D
- Simplifies memory addressing when processing multidimensional data
 - Image processing
 - Solving PDEs on volumes
 - ...



Courtesy: NDVIA

A Simple Example: Vector Addition



A Simple Example: Vector Addition

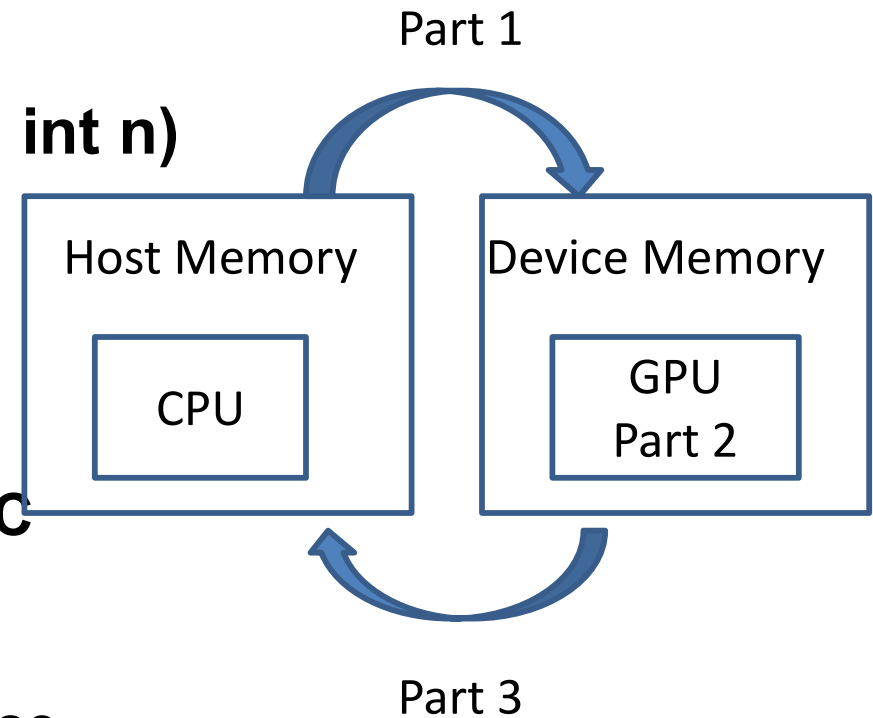
```
// Compute vector sum C = A+B
void vecAdd(float* A, float* B, float* C, int n)
{
    for (i = 0, i < n, i++)
        C[i] = A[i] + B[i];
}
```

GPU friendly!

```
int main()
{
    // Memory allocation for A_h, B_h, and C_h
    // I/O to read A_h and B_h, N elements
    ...
    vecAdd(A_h, B_h, C_h, N);
}
```

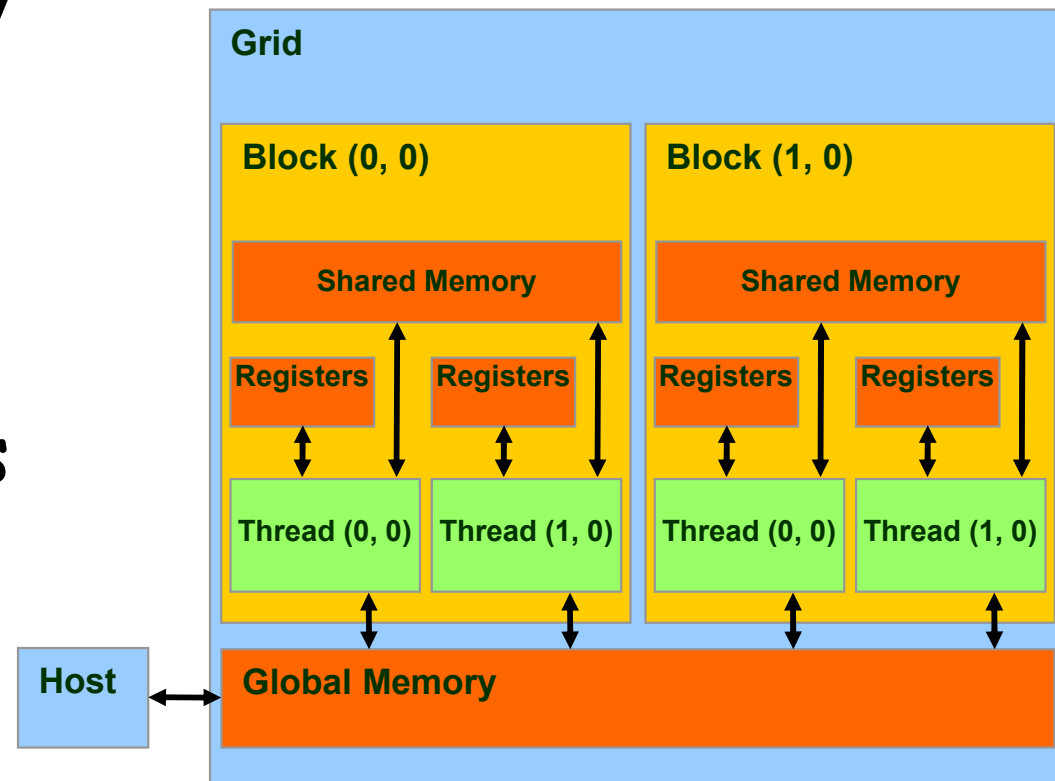
A Simple Example: Vector Addition

```
#include <cuda.h>  
void vecAdd(float* A, float* B, float* C, int n)  
{  
    int size = n* sizeof(float);  
    float* A_d, B_d, C_d;  
    ...  
1. // Allocate device memory for A, B, and C  
    // copy A and B to device memory  
  
2. // Kernel launch code – to have the device  
    // to perform the actual vector addition  
  
3. // copy C from the device memory  
    // Free device vectors  
}
```



CUDA Memory Model

- Global memory
 - Main means of communicating R/W Data between **host** and **device**
 - Contents visible to all threads
 - Long latency access
- Shared memory:
 - Per SM
 - Shared by all threads in a block

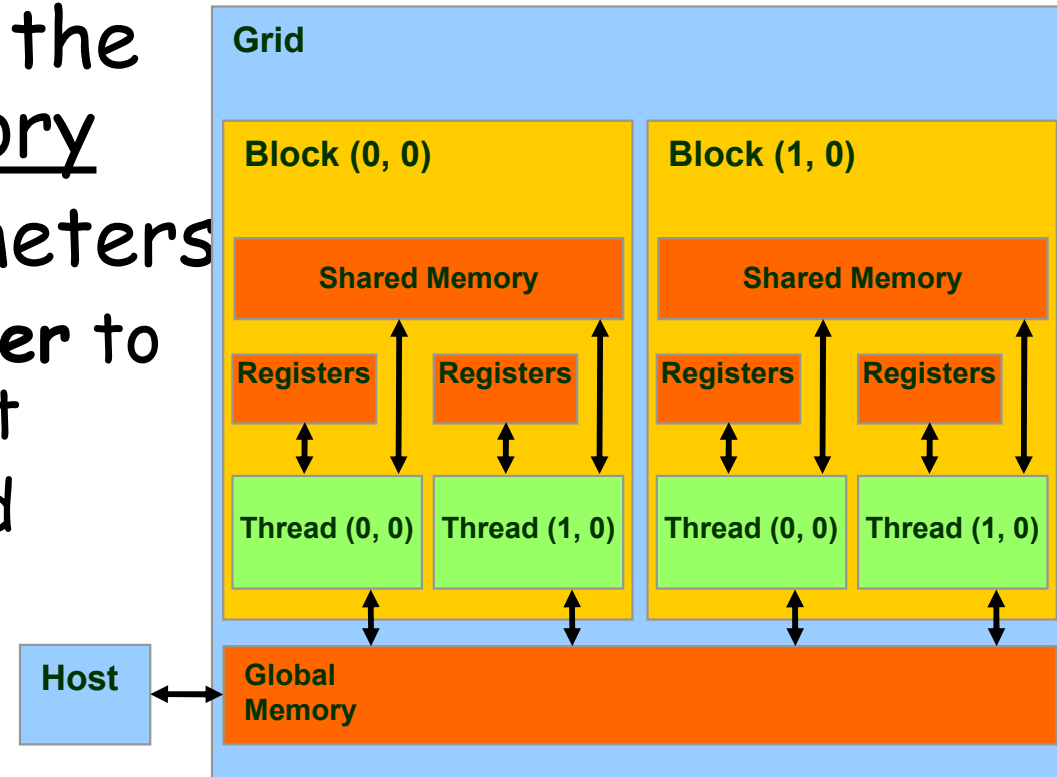


CPU & GPU Memory

- In CUDA, host and devices have separate memory spaces.
 - But in recent GPUs we have Unified Memory Access
- If GPU and CPU are on the same chip, then they share memory space → fusion

CUDA Device Memory Allocation

- **cudaMalloc()**
 - Allocates object in the device Global Memory
 - Requires two parameters
 - Address of a pointer to the allocated object
 - Size of of allocated object
- **cudaFree()**
 - Frees object from device Global Memory
 - Pointer to freed object

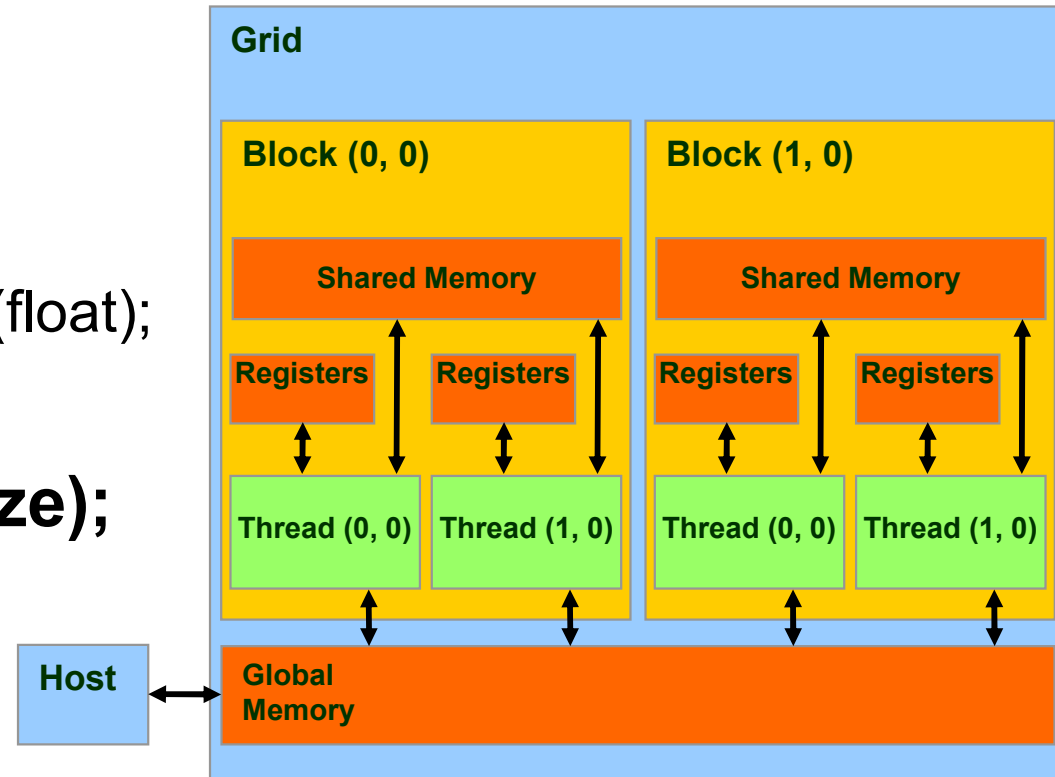


CUDA Device Memory Allocation

Example:

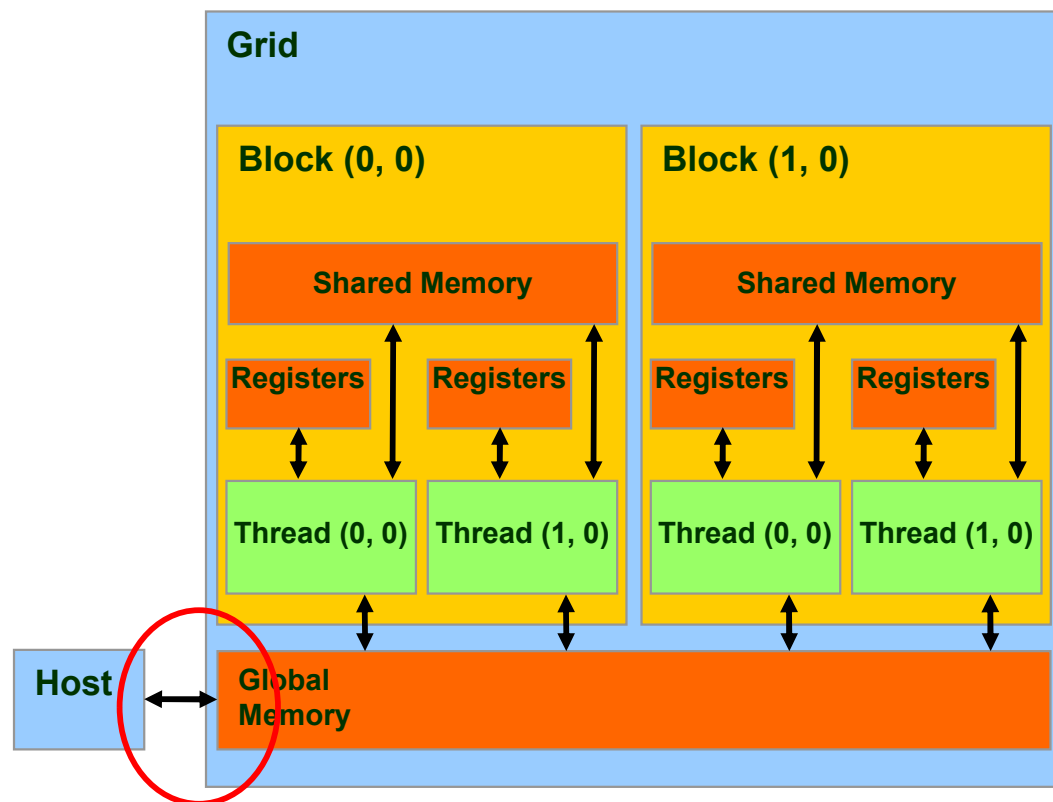
```
WIDTH = 64;  
float* Md  
int size = WIDTH * WIDTH * sizeof(float);
```

```
cudaMalloc((void**)&Md, size);  
cudaFree(Md);
```



CUDA Device Memory Allocation

- `cudaMemcpy()`
 - memory data transfer
 - Requires four parameters
 - Pointer to destination
 - Pointer to source
 - Number of bytes copied
 - Type of transfer
 - Host to Host
 - Host to Device
 - Device to Host
 - Device to Device
- Asynchronous transfer



Important!

`cudaMemcpy()` **cannot** be used to copy between different GPUs in multi-GPUs system

CUDA Device Memory Allocation

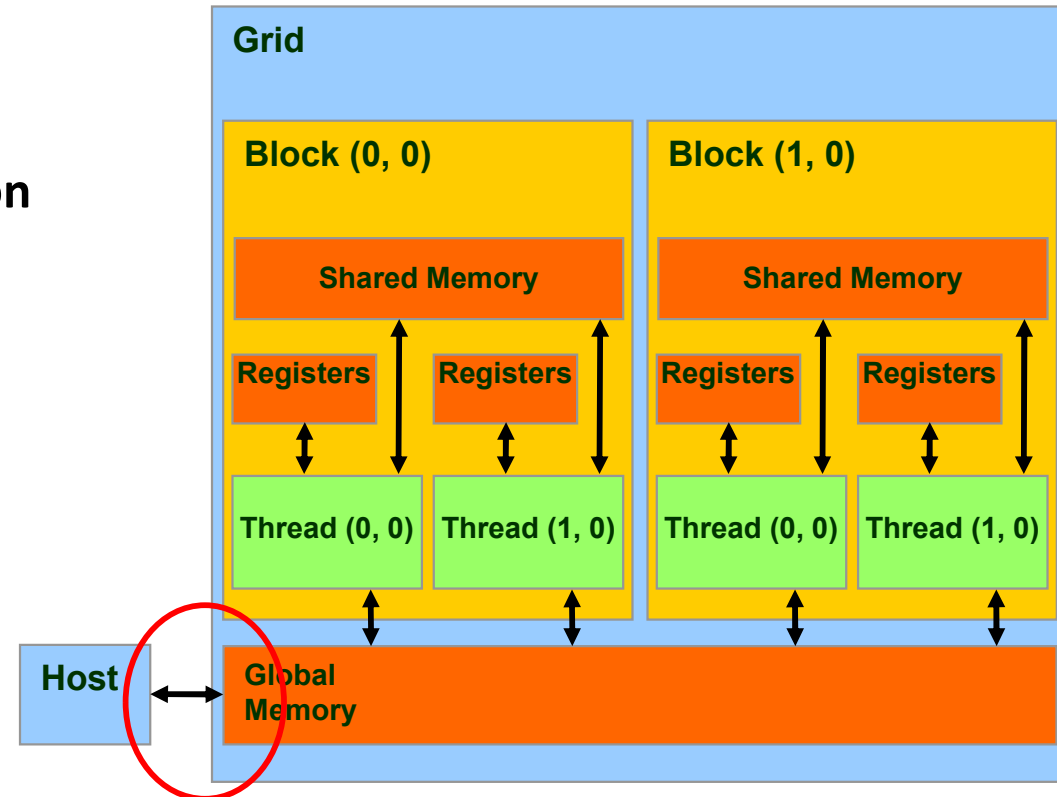
Example:

Destination
pointer

Source
pointer

Size
in bytes

Direction



`cudaMemcpy(Md, M, size, cudaMemcpyHostToDevice);`

`cudaMemcpy(M, Md, size, cudaMemcpyDeviceToHost);`

Note About Error Handling

- Almost all API calls return success or failure.
- The type of the outcome is: `cudaError_t`
- Success → `cudaSuccess`
- Translate the error code to an error message:

```
char * cudaGetErrorString (cudaError_t error)
```

A Simple Example: Vector Addition

```
void vecAdd(float* A, float* B, float* C, int n)
{
```

```
    int size = n * sizeof(float);
    float* A_d, * B_d, * C_d;
```

```
1. // Transfer A and B to device memory
```

```
    cudaMalloc((void **) &A_d, size);
    cudaMemcpy(A_d, A, size, cudaMemcpyHostToDevice);
    cudaMalloc((void **) &B_d, size);
    cudaMemcpy(B_d, B, size, cudaMemcpyHostToDevice);
```

```
    // Allocate device memory for
    cudaMalloc((void **) &C_d, size);
```

```
2. // Kernel invocation code – to be shown later
```

How to launch a kernel?

```
...
3. // Transfer C from device to host
```

```
    cudaMemcpy(C, C_d, size, cudaMemcpyDeviceToHost);
    // Free device memory for A, B, C
    cudaFree(A_d); cudaFree(B_d); cudaFree (C_d);
```


```
}
```



```

int vecAdd(float* A, float* B, float* C, int n)
{
    // A_d, B_d, C_d allocations and copies omitted
    // Run ceil(n/256) blocks of 256 threads each
    vecAddKernel<<ceil(n/256),256>>>(A_d, B_d, C_d, n);
}

```



// Each thread performs one pair-wise addition

__global__

```

void vecAddkernel(float* A_d, float* B_d, float* C_d, int n)
{
    int i = threadIdx.x + blockDim.x * blockIdx.x;
    if(i<n) C_d[i] = A_d[i] + B_d[i];
}

```

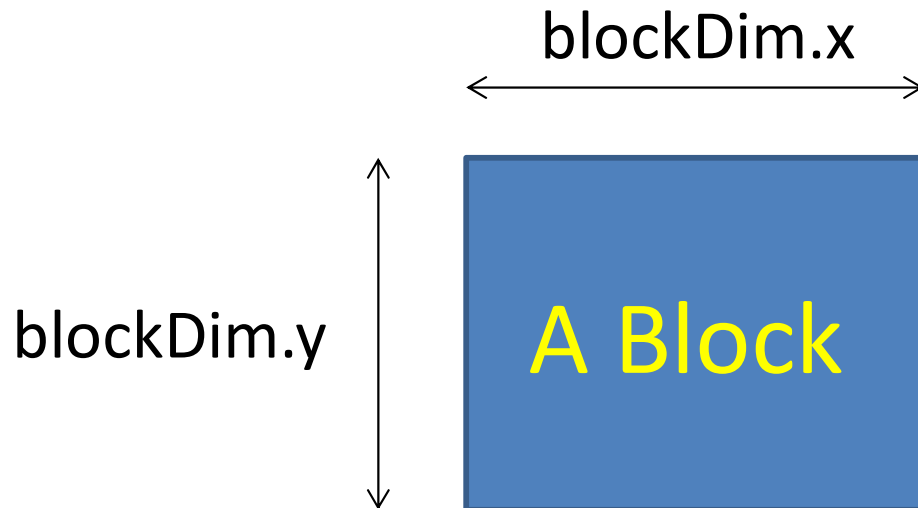
Unique ID
1D grid of 1D blocks

$\text{blockIdx.x} * \text{blockDim.x} + \text{threadIdx.x};$

Unique ID

1D grid of 2D blocks

$\text{blockIdx.x} * \text{blockDim.x} * \text{blockDim.y} +$
 $\text{threadIdx.y} * \text{blockDim.x} +$
 $\text{threadIdx.x};$



Unique ID

Example: 1D grid of 3D blocks

```
blockIdx.x * blockDim.x * blockDim.y *  
blockDim.z +  
threadIdx.z * blockDim.y * blockDim.x +  
threadIdx.y * blockDim.x +  
threadIdx.x;
```

Unique ID

Example: 2D grid of 1D blocks

```
int blockId  = blockIdx.y * gridDim.x +  
blockIdx.x;
```

```
int threadId = blockId * blockDim.x +  
threadIdx.x;
```


Unique ID

It is all a question of how to index data using thread and block IDs.

You can generate unique IDs for any combination of block and grid dimensions.

Kernels

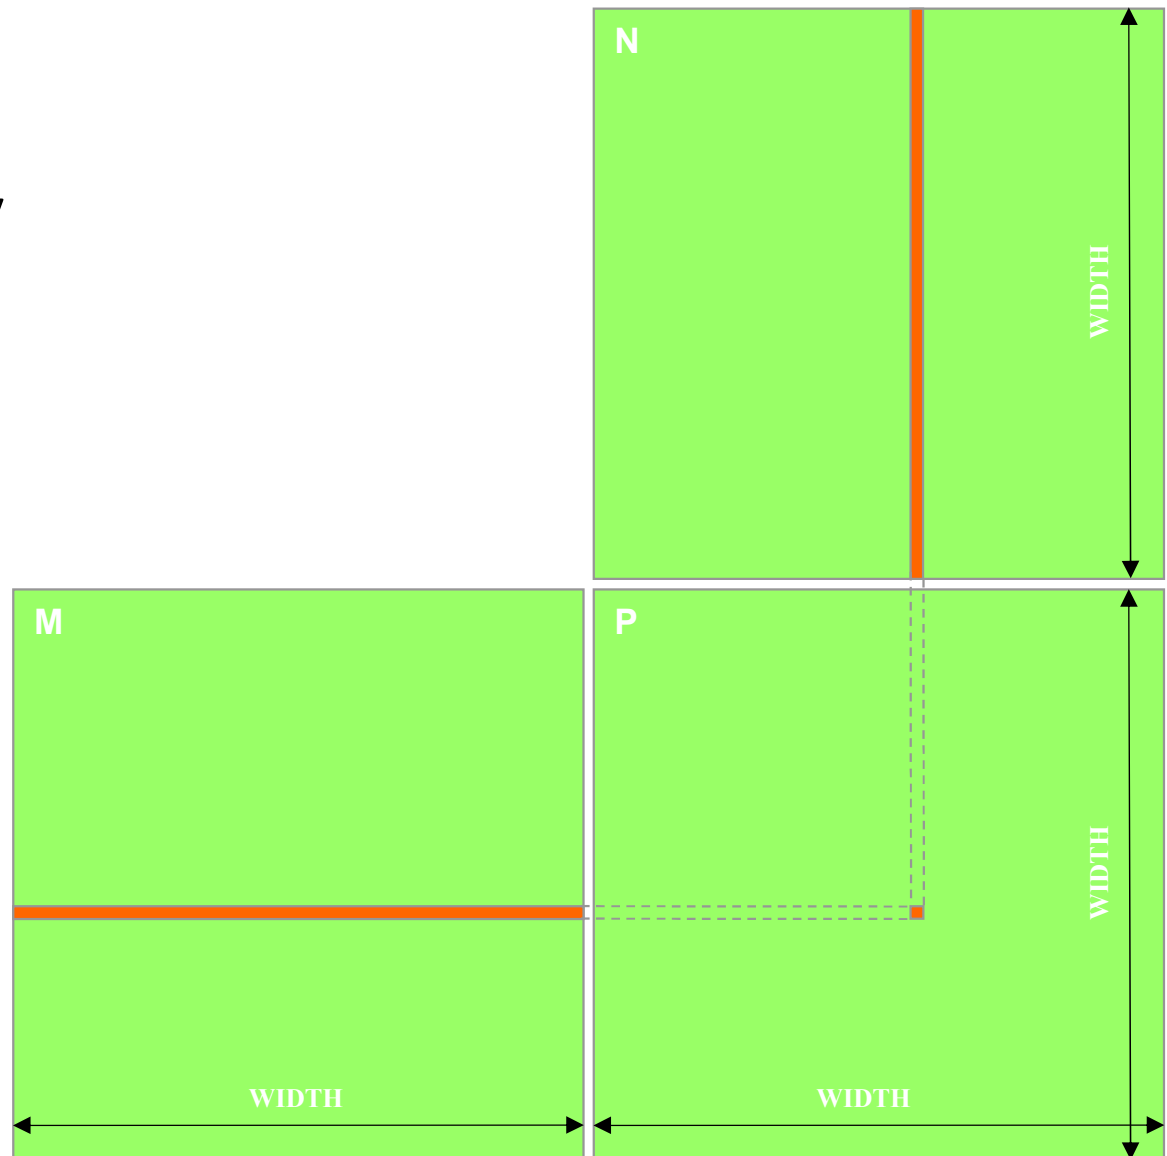
	Executed on the:	Only callable from the:
<code>__device__ float DeviceFunc()</code>	device	device
<code>__global__ void KernelFunc()</code>	device	host
<code>__host__ float HostFunc()</code>	host	host

- `__global__` defines a kernel function. Must return `void`
- `__device__` and `__host__` can be used together
- For functions executed on the device:
 - No static variable declarations inside the function
 - No indirect function calls through pointers

The Hello World of Parallel Programming: Matrix Multiplication

Data Parallelism:

We can safely perform many arithmetic **operations on** the data structures in a **simultaneous** manner.



The Hello World of Parallel Programming: Matrix Multiplication

$M_{0,0}$	$M_{1,0}$	$M_{2,0}$	$M_{3,0}$
$M_{0,1}$	$M_{1,1}$	$M_{2,1}$	$M_{3,1}$
$M_{0,2}$	$M_{1,2}$	$M_{2,2}$	$M_{3,2}$
$M_{0,3}$	$M_{1,3}$	$M_{2,3}$	$M_{3,3}$

M



$M_{0,0}$	$M_{1,0}$	$M_{2,0}$	$M_{3,0}$	$M_{0,1}$	$M_{1,1}$	$M_{2,1}$	$M_{3,1}$	$M_{0,2}$	$M_{1,2}$	$M_{2,2}$	$M_{3,2}$	$M_{0,3}$	$M_{1,3}$	$M_{2,3}$	$M_{3,3}$
-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------

C adopts row-major placement approach
when storing 2D matrix in linear memory address.

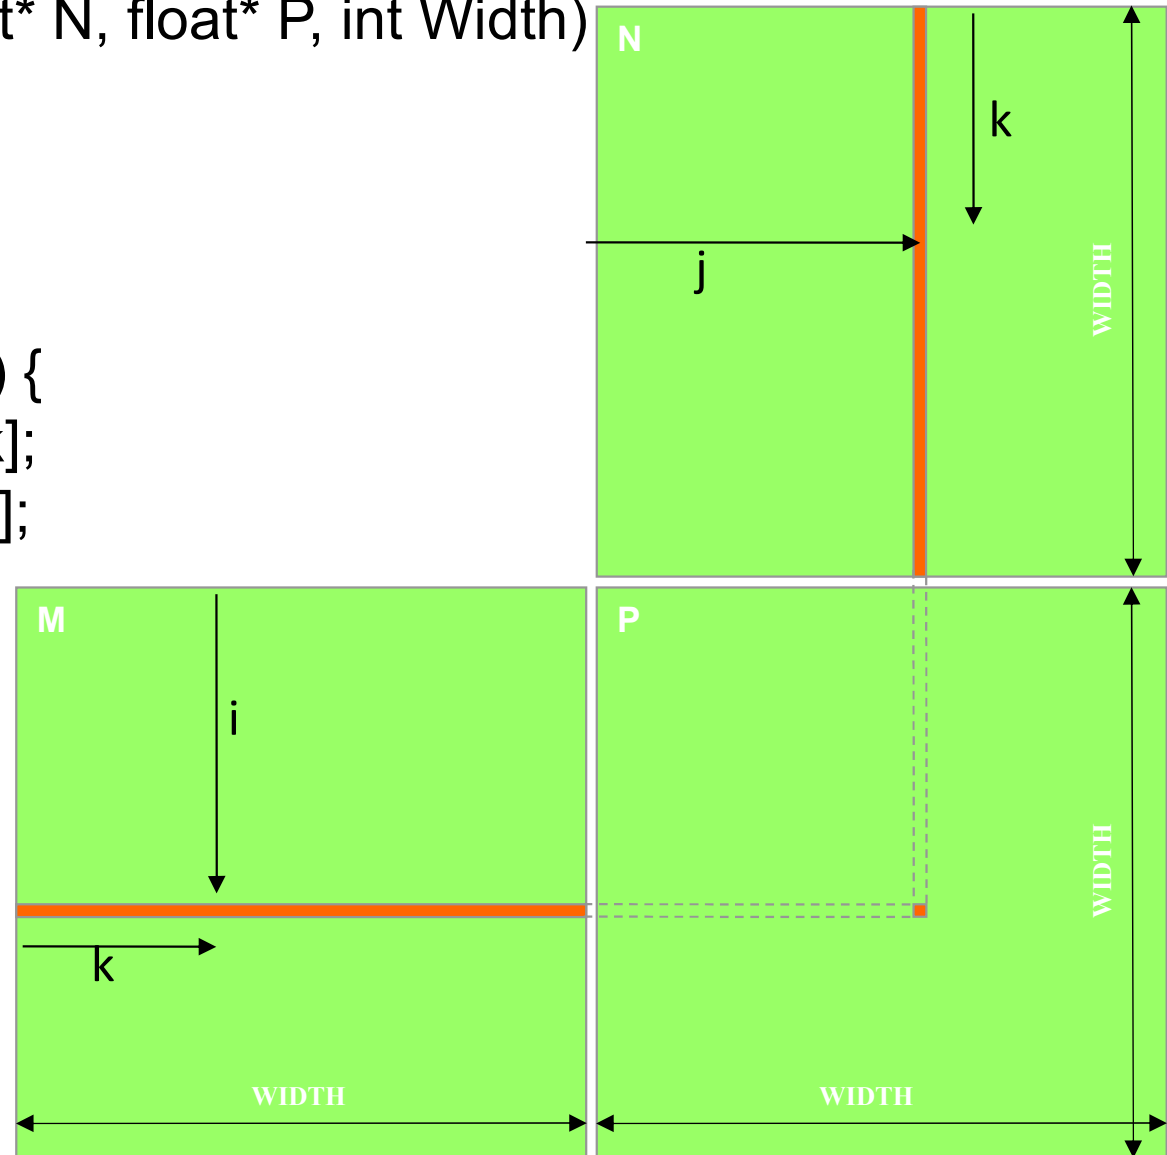
The *Hello World* of Parallel Programming: Matrix Multiplication

```
int main(void) {  
1.  // Allocate and initialize the matrices M, N, P  
    // I/O to read the input matrices M and N  
    ....  
  
2.  // M * N on the device  
    MatrixMultiplication(M, N, P, Width);  
  
3.  // I/O to write the output matrix P  
    // Free matrices M, N, P  
    ...  
    return 0;  
}
```

A Simple main function: executed at the host

The Hello World of Parallel Programming: Matrix Multiplication

```
// Matrix multiplication on the (CPU) host
void MatrixMulOnHost(float* M, float* N, float* P, int Width)
{
    for (int i = 0; i < Width; ++i)
        for (int j = 0; j < Width; ++j) {
            double sum = 0;
            for (int k = 0; k < Width; ++k) {
                double a = M[i * Width + k];
                double b = N[k * Width + j];
                sum += a * b;
            }
            P[i * Width + j] = sum;
        }
}
```



The *Hello World* of Parallel Programming: Matrix Multiplication

```
void MatrixMultiplication(float* M, float* N, float* P, int Width)
{
    int size = Width * Width * sizeof(float);
    float* Md, Nd, Pd;

    1. // Transfer M and N to device memory
       cudaMalloc((void**) &Md, size);
       cudaMemcpy(Md, M, size, cudaMemcpyHostToDevice);
       cudaMalloc((void**) &Nd, size);
       cudaMemcpy(Nd, N, size, cudaMemcpyHostToDevice);

       // Allocate P on the device
       cudaMalloc((void**) &Pd, size);

    2. // Kernel invocation code - to be shown later
       ...

    3. // Transfer P from device to host
       cudaMemcpy(P, Pd, size, cudaMemcpyDeviceToHost);
       // Free device matrices
       cudaFree(Md); cudaFree(Nd); cudaFree (Pd);
}
```

The Hello World of Parallel Programming: Matrix Multiplication

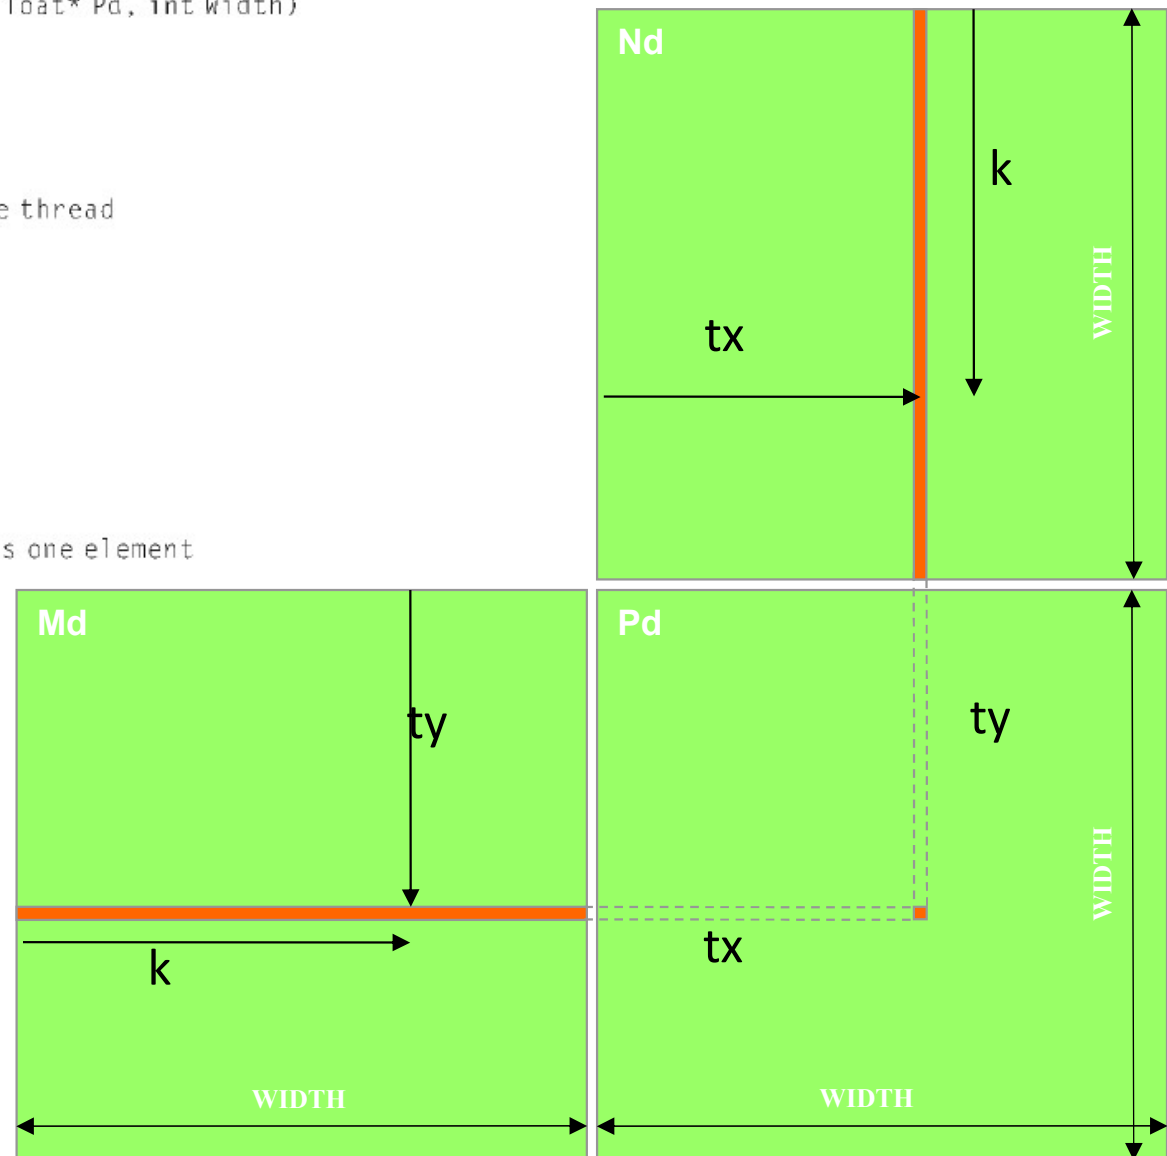
```
// Matrix multiplication kernel - thread specification
__global__ void MatrixMulKernel(float* Md, float* Nd, float* Pd, int Width)
{
    // 2D Thread ID
    int tx = threadIdx.x;
    int ty = threadIdx.y;

    // Pvalue stores the Pd element that is computed by the thread
    float Pvalue = 0;

    for (int k = 0; k < Width; ++k)
    {
        float Mdelement = Md[ty * Width + k];
        float Ndelement = Nd[k * Width + tx];
        Pvalue += Mdelement * Ndelement;
    }

    // Write the matrix to device memory each thread writes one element
    Pd[ty * Width + tx] = Pvalue;
}
```

The Kernel Function



More On Specifying Dimensions

```
// Setup the execution configuration
```

```
    dim3 dimGrid(x, y, z);
```

```
    dim3 dimBlock(x, y, z);
```

```
// Launch the device computation threads!
```

```
MatrixMulKernel<<<dimGrid, dimBlock>>>(Md, Nd, Pd, Width);
```

Important:

- dimGrid and dimBlock are user defined
- **gridDim** and **blockDim** are built-in predefined variable accessible in kernel functions

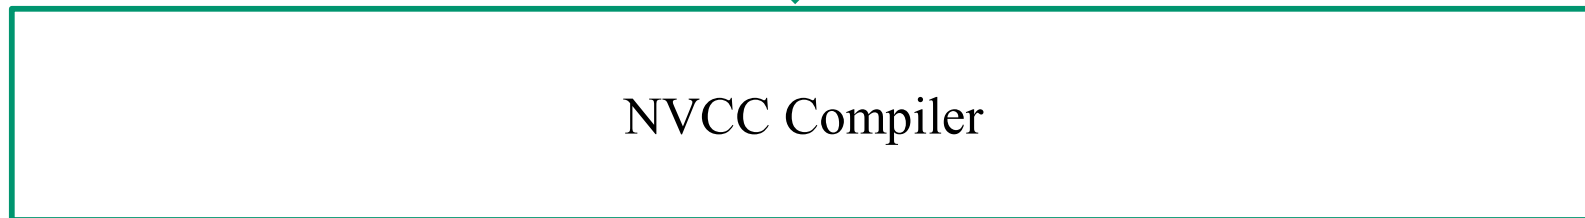
Be Sure To Know:

- Maximum dimensions of a block
- Maximum number of threads per block
- Maximum dimensions of a grid
- Maximum number of blocks per grid

Tools: Compiler: NVCC

`nvcc -o prog prog.cu`

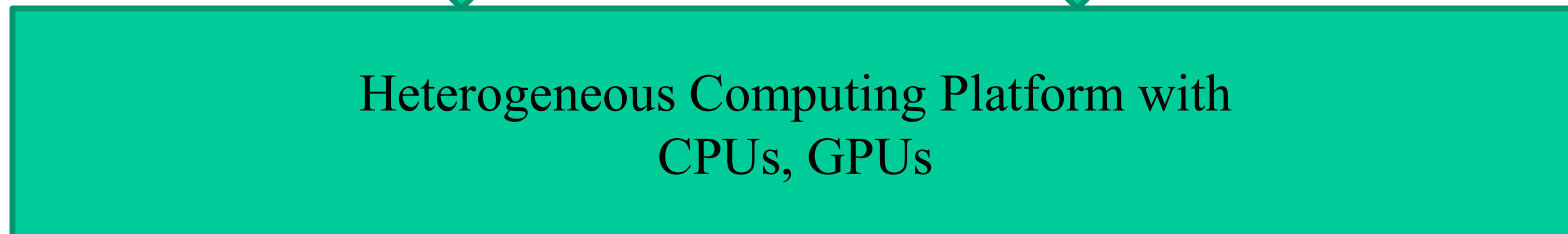
Integrated C programs with CUDA extensions



Host Code



Device Code (PTX)



Conclusions

- Data parallelism is the main source of scalability for parallel programs
- Each CUDA source file can have a mixture of both host and device code.
- What we learned today about CUDA:
 - KernelA<<< nBlk, nTid >>>(args)
 - cudaMalloc()
 - cudaFree()
 - cudaMemcpy()
 - gridDim and blockDim
 - threadIdx and blockIdx
 - dim3