



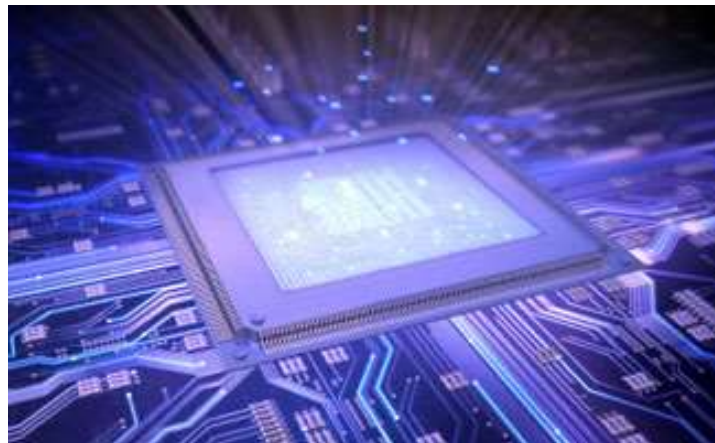
Parallel Computing

OpenMP - I

Mohamed Zahran (aka Z)

mzahran@cs.nyu.edu

<http://www.mzahran.com>



Small and Easy Motivation

```
#include <stdio.h>
#include <stdlib.h>

int main() {

    // Do this part in parallel

    printf( "Hello, World!\n" );

    return 0;
}
```

Small and Easy Motivation

```
#include <stdio.h>
#include <stdlib.h>
#include <omp.h>

int main() {

    omp_set_num_threads(16);

    // Do this part in parallel
    #pragma omp parallel
    {
        printf( "Hello, World!\n" );
    }

    return 0;
}
```

Simple!

Serial Program:

```
void main()
{
    double Res[1000];

    for(int i=0;i<1000;i++) {
        do_huge_comp(Res[i]);
    }
}
```



Parallel Program:

```
void main()
{
    double Res[1000];
    #pragma omp parallel for
    for(int i=0;i<1000;i++) {
        do_huge_comp(Res[i]);
    }
}
```

OpenMP can parallelize many serial programs with **relatively few annotations** that specify parallelism and independence

OpenMP is a small API that **hides cumbersome threading calls with simpler directives**

Interesting Insights About OpenMP




These insights are coming from HPC folks though!

OpenMP

- An **API** for shared-memory parallel programming.
- Designed for systems in which each **thread** can have access to all available memory.
- System is viewed as a collection of cores or CPU's, all of which have access to the same main memory → **shared memory architecture**

A shared memory system



Pragmas

- Special **preprocessor** instructions.
- specified by the C standard for providing additional information to the compiler, beyond what is conveyed in the language itself.
- Compilers that don't support the pragmas ignore them.

#pragma


```

#include <stdio.h>
#include <stdlib.h>
#include <omp.h>

void Hello(void);  /* Thread function */

int main(int argc, char* argv[]) {
    /* Get number of threads from command line */
    int thread_count = strtol(argv[1], NULL, 10);

    # pragma omp parallel num_threads(thread_count)
    Hello();

    return 0;
} /* main */

void Hello(void) {
    int my_rank = omp_get_thread_num();
    int thread_count = omp_get_num_threads();

    printf("Hello from thread %d of %d\n", my_rank, thread_count);

} /* Hello */

```

```
gcc -g -Wall -fopenmp -o omp_hello omp_hello . c
```

```
./omp_hello 4
```

↑
running with 4 threads

Hello from thread 0 of 4
Hello from thread 1 of 4
Hello from thread 2 of 4
Hello from thread 3 of 4

possible
outcomes

↓

Hello from thread 1 of 4
Hello from thread 2 of 4
Hello from thread 0 of 4
Hello from thread 3 of 4

Hello from thread 3 of 4
Hello from thread 1 of 4
Hello from thread 2 of 4
Hello from thread 0 of 4

Another Hello World

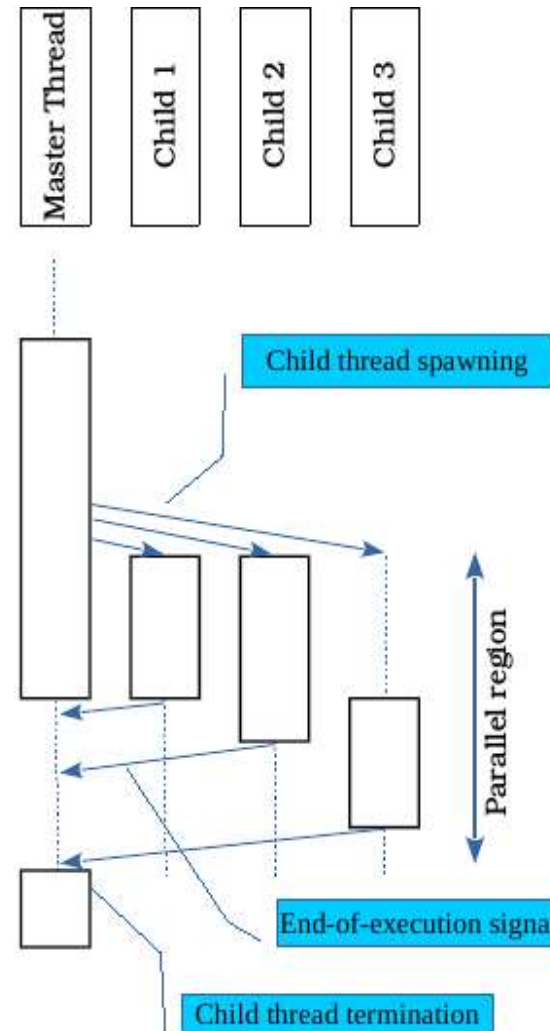
- One of the possible execution sequences:

```
int main (int argc, char **argv)
{
    int numThr = atoi (argv[1]);

    #pragma omp parallel num_threads(numThr)

    cout << "Hello from thread " <<
        Omp_get_thread_num () << endl;


    return 0;
}
```



OpenMP pragmas

- `# pragma omp parallel`
 - Most basic parallel directive.
 - The number of threads that run the following `structured block of code` is determined by the run-time system if the programmer does not specify a number of threads.

A process forking and joining two threads



clause

- Definition: text that modifies a directive.
- The `num_threads` clause can be added to a `parallel` directive.
 - It allows the programmer to specify the number of threads that should execute the following block.

```
# pragma omp parallel num_threads ( thread_count )
```


Clause

Of note...


- There may be system-defined limitations on the number of threads that a program can start.
- The OpenMP standard doesn't guarantee that this will actually start `thread_count` threads.
- **However:** Unless we're trying to start a very large number of threads, we will almost always get the desired number of threads.

Some terminology

- In OpenMP parlance the collection of threads executing the parallel block — the original thread and the new threads — is called a **team**, the original thread is called the **master**, and the additional threads are called **slaves**.



Again: The trapezoidal rule



Serial algorithm

```
/* Input:  a, b, n */  
h = (b-a)/n;  
approx = (f(a) + f(b))/2.0;  
for (i = 1; i <= n-1; i++) {  
    x_i = a + i*h;  
    approx += f(x_i);  
}  
approx = h*approx;
```


A First OpenMP Version

- 1) We identified two type of tasks:
 - a) computation of the areas of individual trapezoids, and
 - b) adding the areas of trapezoids.
- 2) There is no communication among the tasks in the first collection, but each task in the first collection communicates with task 1b.

A First OpenMP Version

- 3) We assumed that there would be many more trapezoids than cores.
- So we aggregated tasks by assigning a contiguous block of trapezoids to each thread.

Assignment of trapezoids to threads



Time	Thread 0	Thread 1
0	<code>global_result = 0 to register</code>	<code>finish my_result</code>
1	<code>my_result = 1 to register</code>	<code>global_result = 0 to register</code>
2	<code>add my_result to global_result</code>	<code>my_result = 2 to register</code>
3	<code>store global_result = 1</code>	<code>add my_result to global_result</code>
4		<code>store global_result = 2</code>

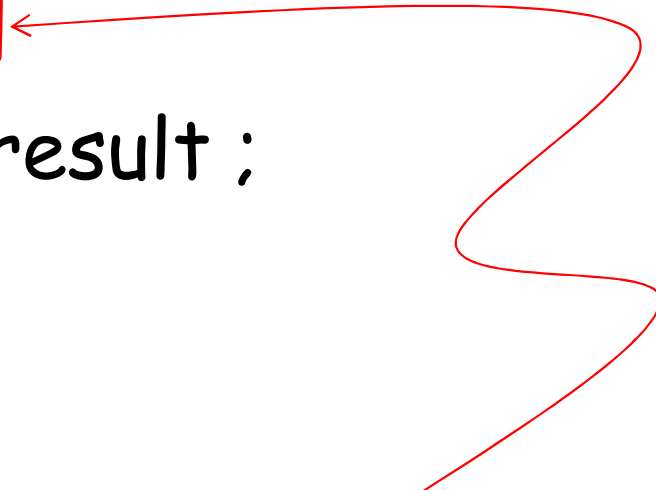
Unpredictable results when two (or more) threads attempt to simultaneously execute:

`global_result += my_result ;`



Mutual exclusion

```
# pragma omp critical  
global_result += my_result ;
```



only one thread can execute
the following structured block at
a time

```

#include <stdio.h>
#include <stdlib.h>
#include <omp.h>

void Trap(double a, double b, int n, double* global_result_p);

int main(int argc, char* argv[]) {
    double    global_result = 0.0;  /* Store result in global_result */
    double    a, b;                 /* Left and right endpoints      */
    int       n;                    /* Total number of trapezoids    */
    int       thread_count;

    thread_count = strtol(argv[1], NULL, 10);
    printf("Enter a, b, and n\n");
    scanf("%lf %lf %d", &a, &b, &n);
    # pragma omp parallel num_threads(thread_count)
    Trap(a, b, n, &global_result);

    printf("With n = %d trapezoids, our estimate\n", n);
    printf("of the integral from %f to %f = %.14e\n",
        a, b, global_result);
    return 0;
} /* main */

```



```

void Trap(double a, double b, int n, double* global_result_p) {
    double h, x, my_result;
    double local_a, local_b;
    int i, local_n;
    int my_rank = omp_get_thread_num();
    int thread_count = omp_get_num_threads();

    h = (b-a)/n;
    local_n = n/thread_count;
    local_a = a + my_rank*local_n*h;
    local_b = local_a + local_n*h;
    my_result = (f(local_a) + f(local_b))/2.0;
    for (i = 1; i <= local_n-1; i++) {
        x = local_a + i*h;
        my_result += f(x);
    }
    my_result = my_result*h;



    # pragma omp critical
        *global_result_p += my_result;
} /* Trap */

```

Another Example

- **Problem:** Count the number of times each ASCII character occurs on a page of text.
- **Input:** ASCII text stored as an array of characters.
- **Output:** A histogram with 128 buckets - one for each ASCII character

Another Example

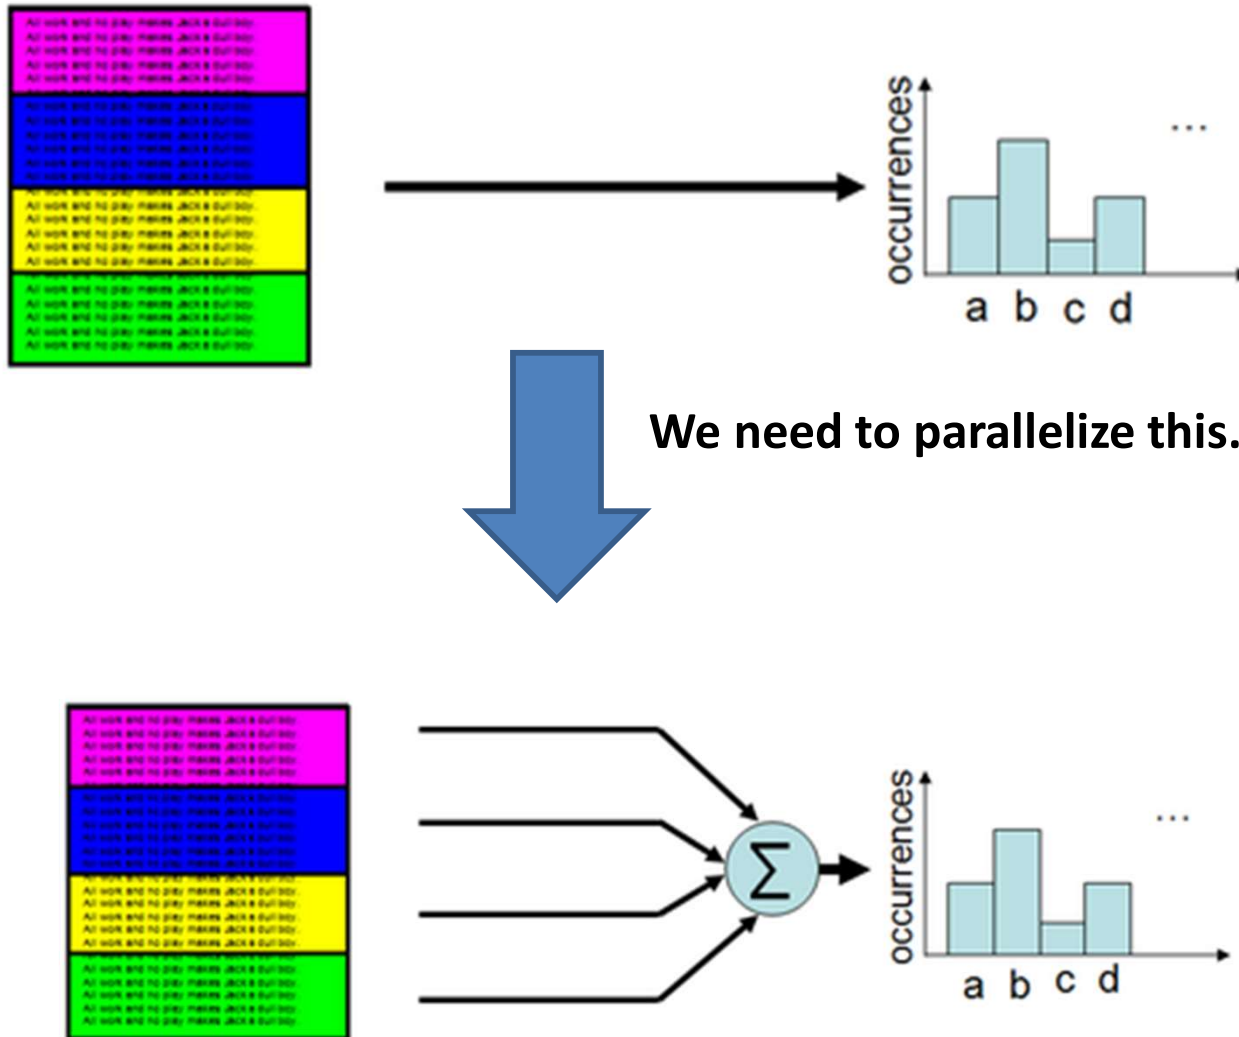


Speed on Quad Core:
10.36 seconds

```
1: void compute_histogram_st(char *page, int page_size, int *histogram){  
2: for(int i = 0; i < page_size; i++){  
3:   char read_character = page[i];  
4:   histogram[read_character]++;  
5: }  
6: }
```

Sequential Version

Another Example



Another Example

```
1: void compute_histogram_st(char *page, int page_size, int *histogram){  
2: #pragma omp parallel for  
3: for(int i = 0; i < page_size; i++){  
4:     char read_character = page[i];  
5:     histogram[read_character]++;  
6: }
```

The above code does not work!! Why?

Another Example

```
1: void compute_histogram_mt2(char *page, int page_size, int *histogram){  
2:  #pragma omp parallel for  
3:  for(int i = 0; i < page_size; i++){  
4:      char read_character = page[i];  
5:      #pragma omp atomic  
6:      histogram[read_character]++;  
7:  }  
8: }
```

Speed on Quad Core:

114.89 seconds

> 10x slower than the single thread version!!

Another Example

```
1: void compute_histogram_mt3(char *page,
                             int page_size,
                             int *histogram, int num_buckets){
2:  #pragma omp parallel
3:  {
4:      int local_histogram[NUM_THREADS][num_buckets];
5:      int tid = omp_get_thread_num();
6:      #pragma omp for nowait
7:      for(int i = 0; i < page_size; i++){
8:          char read_character = page[i];
9:          local_histogram[tid][read_character]++;
10:     }
11:     for(int i = 0; i < num_buckets; i++){
12:         #pragma omp atomic
13:         histogram[i] += local_histogram[tid][i];
14:     }
15: }
16: }
```

Runs in 3.8 secs
Why speedup
is not 4 yet?

Another Example

```
void compute_histogram_mt4(char *page, int page_size,
                           int *histogram, int num_buckets){
1:     int num_threads = omp_get_max_threads();
2:     #pragma omp parallel
3:     {
4:         __declspec (align(64)) int local_histogram[num_threads][num_buckets];
5:         int tid = omp_get_thread_num();
6:         #pragma omp for
7:         for(int i = 0; i < page_size; i++){
8:             char read_character = page[i];
9:             local_histogram[tid][read_character]++;
10:        }
11:
12:        #pragma omp single
13:        for(int t = 0; t < num_threads; t++){
14:            for(int i = 0; i < num_buckets; i++)
15:                histogram[i] += local_histogram[t][i];
16:        }
17: }
```

Speed is
4.42 seconds.
Slower than the
previous version.

Another Example

```
void compute_histogram_mt4(char *page, int page_size,  
                           int *histogram, int num_buckets){  
1:     int num_threads = omp_get_max_threads();  
2:     #pragma omp parallel  
3:     {  
4:         __declspec (align(64)) int local_histogram[num_threads][num_buckets];  
5:         int tid = omp_get_thread_num();  
6:         #pragma omp for  
7:         for(int i = 0; i < page_size; i++){  
8:             char read_character = page[i];  
9:             local_histogram[tid][read_character]++;  
10:        }  
11:  
12:        #pragma omp for  
13:        for(int i = 0; i < num_buckets; i++){  
14:            for(int t = 0; t < num_threads; t++)  
15:                histogram[i] += local_histogram[t][i];  
16:        }  
17: }
```

Speed is
3.60 seconds.


What Can We Learn from the Previous Example?

- Atomic operations
 - They are expensive
 - Yet, they are fundamental building blocks.
- Synchronization:
 - correctness vs performance loss
 - Rich interaction of hardware-software tradeoffs

OpenMP Parallel Programming

1. Start with a *parallelizable* algorithm
 - loop-level parallelism is necessary
2. Implement serially
3. Test and Debug
4. Annotate the code with parallelization (and synchronization) directives
 - Hope for linear speedup
5. Test and Debug

OpenMP uses the fork-join model of parallel execution.




All OpenMP programs begin with a single thread: **master thread** (ID = 0)

FORK: the master thread then creates a team of parallel *threads*.

JOIN: When the team threads complete the statements in the parallel region construct, they synchronize and terminate

Programming Model - Threading

```
int main() {  
    // serial region  
    printf("Hello...");  
    // parallel region  
    #pragma omp parallel  
    {  
        printf("World");  
    }  
    // serial again  
    printf("!");  
}
```



We didn't use `omp_set_num_threads()`, what will be the output?

What we learned so far

- `#include <omp.h>`
- `gcc -fopenmp ...`
- `omp_set_num_threads(x);`
- `omp_get_thread_num();`
- `omp_get_num_threads();`
- `#pragma omp parallel [num_threads(x)]`
- `#pragma omp atomic`
- `#pragma omp single`

Conclusions

- OpenMP is a standard for programming shared-memory systems.
- OpenMP uses both special functions and preprocessor directives called pragmas.
- OpenMP programs start multiple threads rather than multiple processes.
- Many OpenMP directives can be modified by clauses.