



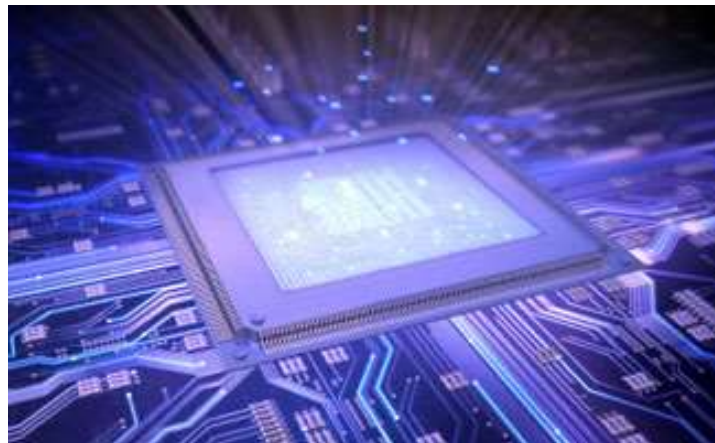
# Parallel Computing

## OpenMP III

Mohamed Zahran (aka Z)

mzahran@cs.nyu.edu

<http://www.mzahran.com>





# MORE ABOUT LOOPS IN OPENMP: SORTING

# Bubble Sort

```
for (list_length = n; list_length >= 2; list_length--)  
    for (i = 0; i < list_length-1; i++)  
        if (a[i] > a[i+1]) {  
            tmp = a[i];  
            a[i] = a[i+1];  
            a[i+1] = tmp;  
        }
```

Loop-carried dependency in inner loop

Loop-carried dependency in outer loop

What can we do?



# Serial Odd-Even Transposition Sort

```
for (phase = 0; phase < n; phase++)  
    if (phase % 2 == 0)  
        for (i = 1; i < n; i += 2)  
            if (a[i-1] > a[i]) Swap(&a[i-1], &a[i]);  
    else  
        for (i = 1; i < n-1; i += 2)  
            if (a[i] > a[i+1]) Swap(&a[i], &a[i+1]);
```

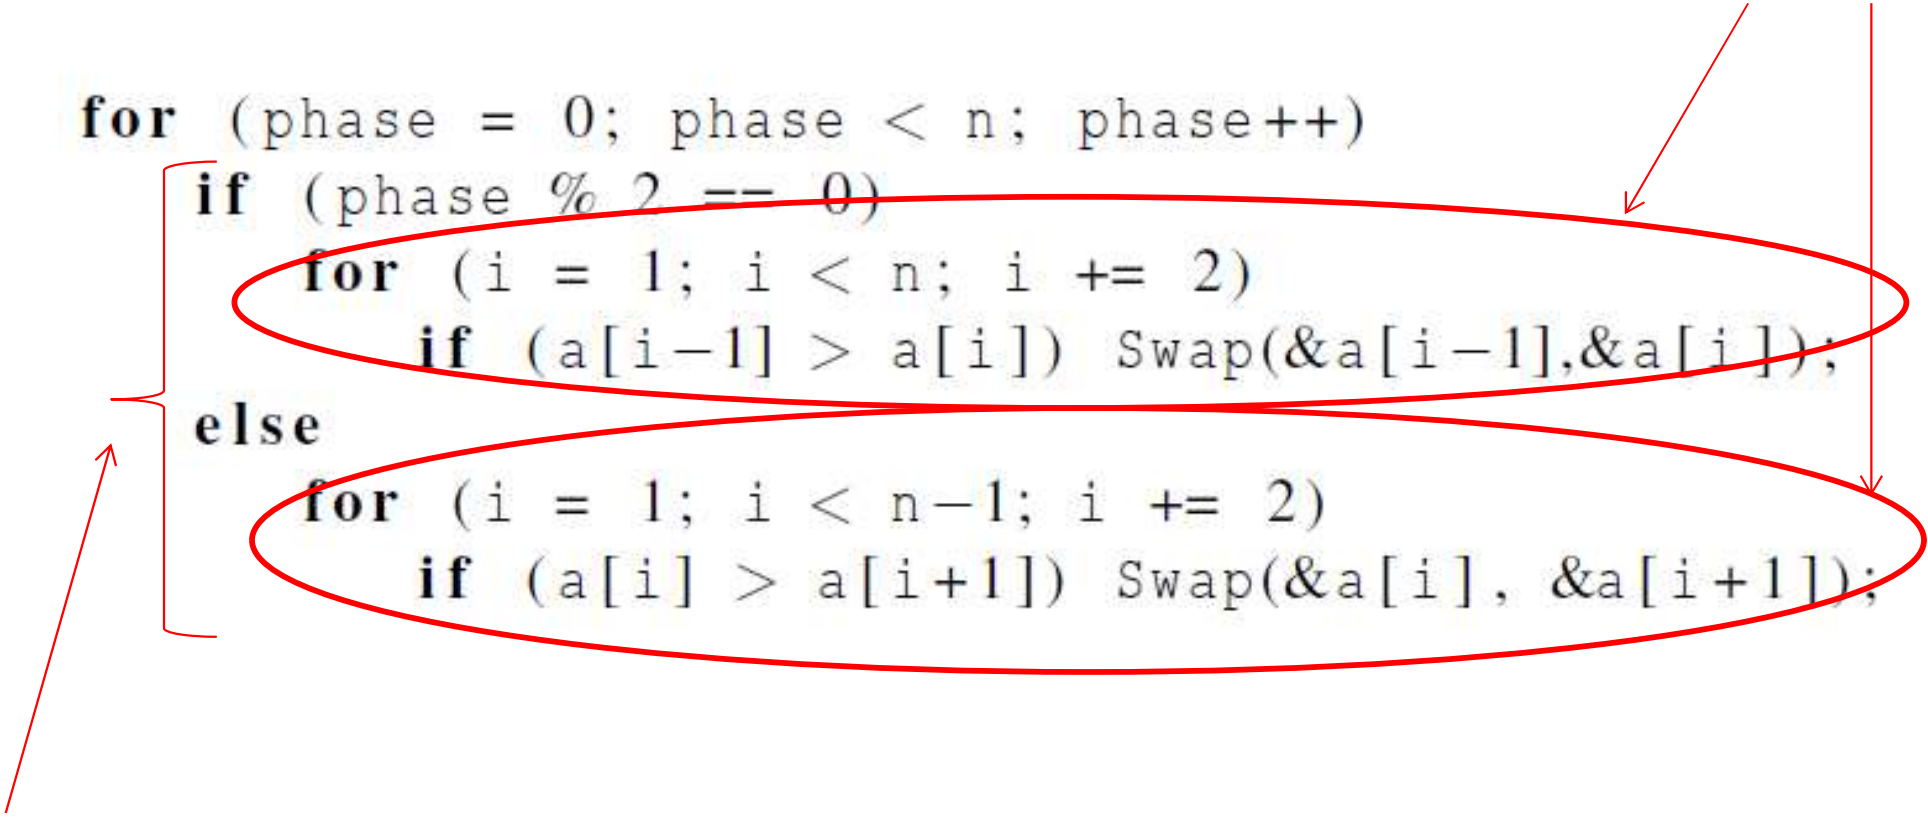
# Serial Odd-Even Transposition Sort

Phase	Subscript in Array			
	0	1	2	3
0	9	$\leftrightarrow$ 7	8	$\leftrightarrow$ 6
	7	9	6	8
1	7	9	$\leftrightarrow$ 6	8
	7	6	9	8
2	7	$\leftrightarrow$ 6	9	$\leftrightarrow$ 8
	6	7	8	9
3	6	7	$\leftrightarrow$ 8	9
	6	7	8	9

# Serial Odd-Even Transposition Sort

No dependence in inner loops

```
for (phase = 0; phase < n; phase++)  
    if (phase % 2 == 0)  
        for (i = 1; i < n; i += 2)  
            if (a[i-1] > a[i]) Swap(&a[i-1], &a[i]);  
    else  
        for (i = 1; i < n-1; i += 2)  
            if (a[i] > a[i+1]) Swap(&a[i], &a[i+1]);
```



Outer-loop carried dependence

# First OpenMP Odd-Even Sort

```
for (phase = 0; phase < n; phase++) {  
    if (phase % 2 == 0)  
#        pragma omp parallel for num_threads(thread_count) \  
            default(none) shared(a, n) private(i, tmp)  
            for (i = 1; i < n; i += 2) {  
                if (a[i-1] > a[i]) {  
                    tmp = a[i-1];  
                    a[i-1] = a[i];  
                    a[i] = tmp;  
                }  
            }  
    else  
#        pragma omp parallel for num_threads(thread_count) \  
            default(none) shared(a, n) private(i, tmp)  
            for (i = 1; i < n-1; i += 2) {  
                if (a[i] > a[i+1]) {  
                    tmp = a[i+1];  
                    a[i+1] = a[i];  
                    a[i] = tmp;  
                }  
            }  
}
```

What if a thread proceeds from  
phase p to phase p+1 before other  
threads?

Performance issue:  
For each outer iteration, OpenMP  
will fork-join threads → Repeated  
overhead per iteration.  
Can we do better?



# Second OpenMP Odd-Even Sort

```
# pragma omp parallel num_threads(thread_count) \
    default(none) shared(a, n) private(i, tmp, phase)
for (phase = 0; phase < n; phase++) {
    if (phase % 2 == 0)
#        pragma omp for
        for (i = 1; i < n; i += 2) {
            if (a[i-1] > a[i]) {
                tmp = a[i-1];
                a[i-1] = a[i];
                a[i] = tmp;
            }
        }
    else
#        pragma omp for
        for (i = 1; i < n-1; i += 2) {
            if (a[i] > a[i+1]) {
                tmp = a[i+1];
                a[i+1] = a[i];
                a[i] = tmp;
            }
        }
    }
}
```

*for directive* does not fork any threads. But uses whatever threads that have been forked before in the enclosing parallel block.




Odd-even sort with two parallel for directives and two for directives.

(Times are in seconds.)

Array of 20,000 elements

thread_count	1	2	3	4
Two parallel <b>for</b> directives	0.770	0.453	0.358	0.305
Two <b>for</b> directives	0.732	0.376	0.294	0.239





# SCHEDULING LOOPS

# Take a look at this:

```
sum = 0.0;  
for (i = 0; i <= n; i++)  
    sum += f(i);
```

- Usually, the default for many OpenMP implementations is to parallelize the above iterations as block of consecutive  $n/\text{thread\_count}$  iterations to each thread.
- What if  $f(i)$  has latency that increases with  $i$ ? What is the best schedule then?

## Example of function *f*.

```
double f(int i) {  
    int j, start = i*(i+1)/2, finish = start + i;  
    double return_val = 0.0;  
  
    for (j = start; j <= finish; j++) {  
        return_val += sin(j);  
    }  
    return return_val;  
} /* f */
```

```

sum = 0.0;
for (i = 0; i <= n; i++)
    sum += f(i);

```

Wouldn't this be better? (why?)

Thread	Iterations
0	$0, n/t, 2n/t, \dots$
1	$1, n/t + 1, 2n/t + 1, \dots$
$\vdots$	$\vdots$
$t - 1$	$t - 1, n/t + t - 1, 2n/t + t - 1, \dots$

Assignment of work  
using cyclic partitioning.

# Results

- $f(i)$  calls the sin function  $i$  times.
- Assume the time to execute  $f(2i)$  requires approximately twice as much time as the time to execute  $f(i)$ .
- $n = 10,000$ 
  - one thread
  - run-time = 3.67 seconds.

# Results

- $n = 10,000$ 
  - two threads
  - default assignment
  - run-time = 2.76 seconds
  - speedup = 1.33
- $n = 10,000$ 
  - two threads
  - cyclic assignment
  - run-time = 1.84 seconds
  - speedup = 1.99



# The Schedule Clause

- Default schedule:

```
sum = 0.0;
# pragma omp parallel for num_threads(thread_count) \
    reduction(+:sum)
for (i = 0; i <= n; i++)
    sum += f(i);
```

- Cyclic schedule:

```
sum = 0.0;
# pragma omp parallel for num_threads(thread_count) \
    reduction(+:sum) schedule(static,1)
for (i = 0; i <= n; i++)
    sum += f(i);
```

# schedule ( type [, chunksize] )

- Type can be:
  - **static**: the iterations can be assigned to the threads before the loop is executed.
  - **dynamic** or **guided**: the iterations are assigned to the threads while the loop is executing.
  - **auto**: the compiler and/or the run-time system determine the schedule.
  - **runtime**: the schedule is determined at run-time.
- The chunksize is a positive integer.

# The Static Schedule Type

Example: twelve iterations, 0, 1, . . . , 11, and three threads

`schedule(static, 1)`

Thread 0 : 0, 3, 6, 9

Thread 1 : 1, 4, 7, 10

Thread 2 : 2, 5, 8, 11

`schedule(static, 2)`

Thread 0 : 0, 1, 6, 7

Thread 1 : 2, 3, 8, 9

Thread 2 : 4, 5, 10, 11

`schedule(static, 4)`

Thread 0 : 0, 1, 2, 3

Thread 1 : 4, 5, 6, 7

Thread 2 : 8, 9, 10, 11

How to implement the default block scheduling using the static schedule?

# The Dynamic Schedule Type

- The iterations are broken up into chunks of **chunksize** consecutive iterations.
- Each thread executes a chunk, and when a thread finishes a chunk, it **requests another one from the runtime system**.
- This continues until all the iterations are completed.
- The chunksize can be omitted. When it is omitted, a **default chunksize of 1** is used.

# The Guided Schedule Type

- Each thread also executes a chunk, and when a thread finishes a chunk, it requests another one.
- As chunks are completed **the size of the new chunks decreases.**
- If no chunksize is specified, the size of the chunks decreases down to 1.
- If chunksize is specified, it decreases down to chunksize, with the exception that the very last chunk can be smaller than chunksize.

## Example:

Assignment of trapezoidal rule iterations 1-9999 using a guided schedule with two threads.

Thread	Chunk	Size of Chunk	Remaining Iterations
0	1 – 5000	5000	4999
1	5001 – 7500	2500	2499
1	7501 – 8750	1250	1249
1	8751 – 9375	625	624
0	9376 – 9687	312	312
1	9688 – 9843	156	156
0	9844 – 9921	78	78
1	9922 – 9960	39	39
1	9961 – 9980	20	19
1	9981 – 9990	10	9
1	9991 – 9995	5	4
0	9996 – 9997	2	2
1	9998 – 9998	1	1
0	9999 – 9999	1	0

# The Runtime Schedule Type

- The system uses the environment variable **OMP\_SCHEDULE** to determine at run-time how to schedule the loop.
- The OMP\_SCHEDULE environment variable can take on any of the values that can be used for a static, dynamic, or guided schedule.
- Example:  
**export OMP\_SCHEDULE = "static,1"**



# Another Way for Controlling the Schedule

- Using the `omp_set_schedule` function. Syntax:

```
void omp_set_schedule (omp_sched_t kind,  
                        int chunk_size);
```

- Where `kind` is one of:

- `omp_sched_static`
- `omp_sched_dynamic`
- `omp_sched_guided`
- `omp_sched_auto`

## Keep in mind:

- There is an overhead in using the schedule directive
- The overhead is higher in dynamic than static schedules
- The overhead of guided is the greatest of all three.
- So: if we get satisfactory performance without schedule then don't use schedule.

# Rules of thumb

- If each iteration requires roughly the same amount of computation → default is best
- If the cost of each iteration increases/decreases linearly as the loop executes → static with small chunksize
- If the cost cannot be determined → you need to try several schedules:  
schedule(runtime) and try different options with OMP\_SCHEDULE

# Question

Can we parallelize the following loop? If yes, do it. If not, why not?

```
a[0] = 0;  
for( i = 1; i < n; i++)  
    a[i] = a[i-1] + i;
```

Hint:

$$a[1] = a[0] + 1$$

$$a[2] = a[1] + 2 = a[0] + 1 + 2$$

$$a[3] = a[2] + 3 = a[0] + 1 + 2 + 3$$

$$a[i] = 0 + 1 + 2 + \dots + i = i(i+1)/2$$

# Conclusions

- OpenMP depends on compiler directives, runtime library, and environment variable.
- The main concept to parallelize a program with OpenMP is how to have independent for-loops.
- Many aspects of OpenMP are still implementation dependent, so you need to be careful!