



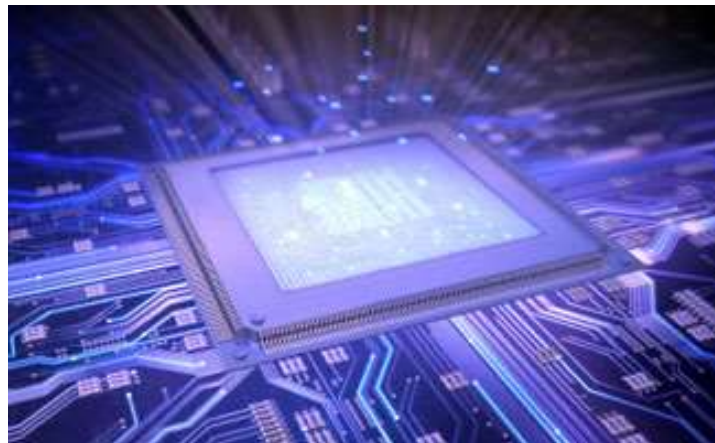
# Parallel Computing

## OpenMP - IV

Mohamed Zahran (aka Z)

mzahran@cs.nyu.edu

<http://www.mzahran.com>






# PRODUCERS AND CONSUMERS

# Queues

- A natural data structure to use in many multithreaded applications.
- The two main operations: enqueue and dequeue
- For example, suppose we have several “producer” threads and several “consumer” threads.
  - Producer threads might “produce” requests for data.
  - Consumer threads might “consume” the request by finding or generating the requested data.

# Example of Usage: Message-Passing

- Each thread could have a **shared message queue**, and when one thread wants to "send a message" to another thread, it could enqueue the message in the destination thread's queue.
- A thread could receive a message by dequeuing the message at the head of its message queue.



# Example of Usage: Message-Passing

Each thread executes the following:

```
for (sent_msgs = 0; sent_msgs < send_max; sent_msgs++) {  
    Send_msg();  
    Try_receive();  
}
```

```
while (!Done())  
    Try_receive();
```

# Send\_msg()

```
msg = random();  
dest = random() % thread_count;  
# pragma omp critical  
  Enqueue(queue, dest, my_rank, msg);
```

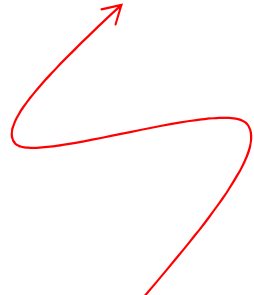
# Try\_receive()

```
if (queue_size == 0) return;  
else if (queue_size == 1)  
#    pragma omp critical  
    Dequeue(queue, &src, &msg);  
else  
    Dequeue(queue, &src, &msg);  
Print_message(src, msg);
```

When queue size is 1, dequeue affects the tail pointer.

# Termination Detection

```
queue_size = enqueued - dequeued;  
if (queue_size == 0 && done_sending == thread_count)  
    return TRUE;  
else  
    return FALSE;
```



each thread increments this after  
completing its for loop



# Startup (1)

- When the program begins execution, a single thread, the master thread, will get command line arguments and allocate an **array of message queues**: one for each thread.
- This array needs to be **shared among the threads**.
- Each thread allocates its queue in the array.

# Startup (2)

- One or more threads may finish allocating their queues before some other threads.
- We need an explicit barrier so that when a thread encounters the barrier, it blocks until all the threads in the team have reached the barrier.

```
# pragma omp barrier
```

# The Atomic Directive

- Higher performance than critical
- It can only protect critical sections that consist of **a single C assignment statement**.

```
# pragma omp atomic
```

- The statement must have one of the following forms:

```
x <op>= <expression>;  
x++;  
++x;  
x--;  
--x;
```

Must not reference X

+, \*, -, /, &, ^, |, <<, or >>

# Critical Sections

```
# pragma omp critical(name)
```

- OpenMP provides the option of adding a name to a critical directive:
- When we do this, **two blocks protected with critical directives with different names can be executed simultaneously.**
- However, the **names are set during compilation**, and we may want a different critical section for each thread's queue.

# Locks

- A lock consists of a data structure and functions that allow the programmer to explicitly **enforce mutual exclusion** in a critical section.



# Locks: main actions

*/\* Executed by one thread \*/*

Initialize the lock data structure;

. . .

*/\* Executed by multiple threads \*/*

Attempt to lock or set the lock data structure;

Critical section;

Unlock or unset the lock data structure;

. . .

*/\* Executed by one thread \*/*

Destroy the lock data structure;

# Locks: main actions

void omp\_**init**\_lock(omp\_lock\_t \* lock\_p);

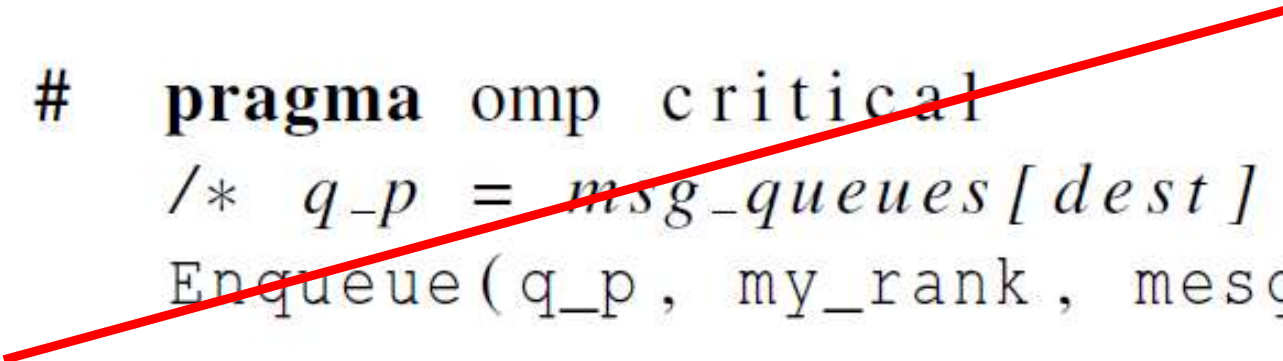
void omp\_**set**\_lock(omp\_lock\_t \* lock\_p);

void omp\_**unset**\_lock(omp\_lock\_t \* lock\_p);

void omp\_**destroy**\_lock(omp\_lock\_t \* lock\_p);

# Using Locks in the Message-Passing Program

```
# pragma omp critical  
/* q_p = msg_queues[dest] */  
Enqueue(q_p, my_rank, msg);
```



```
/* q_p = msg_queues[dest] */  
omp_set_lock(&q_p->lock);  
Enqueue(q_p, my_rank, msg);  
omp_unset_lock(&q_p->lock);
```



# Using Locks in the Message-Passing Program


```
# pragma omp critical  
/* q_p = msg_queues[my_rank] */  
Dequeue(q_p, &src, &msg);
```

```
/* q_p = msg_queues[my_rank] */  
omp_set_lock(&q_p->lock);  
Dequeue(q_p, &src, &msg);  
omp_unset_lock(&q_p->lock);
```


# Some Caveats

1. You shouldn't mix the different types of mutual exclusion for a single critical section.
  - i.e. do not mix atomic and critical for the same variable update
2. There is no guarantee of fairness in mutual exclusion constructs.
  - A thread can be blocked forever!
3. It can be dangerous to "nest" mutual exclusion constructs.

# Dividing Work Among Threads




# Dividing Work Among Threads



```
#pragma omp parallel for  
for_loop
```

# Dividing Work Among Threads




```
#pragm omp parallel
#pragma omp sections
{
    #pragma omp section
        structured_block

    #pragma omp section
        structured_block
}
```

```

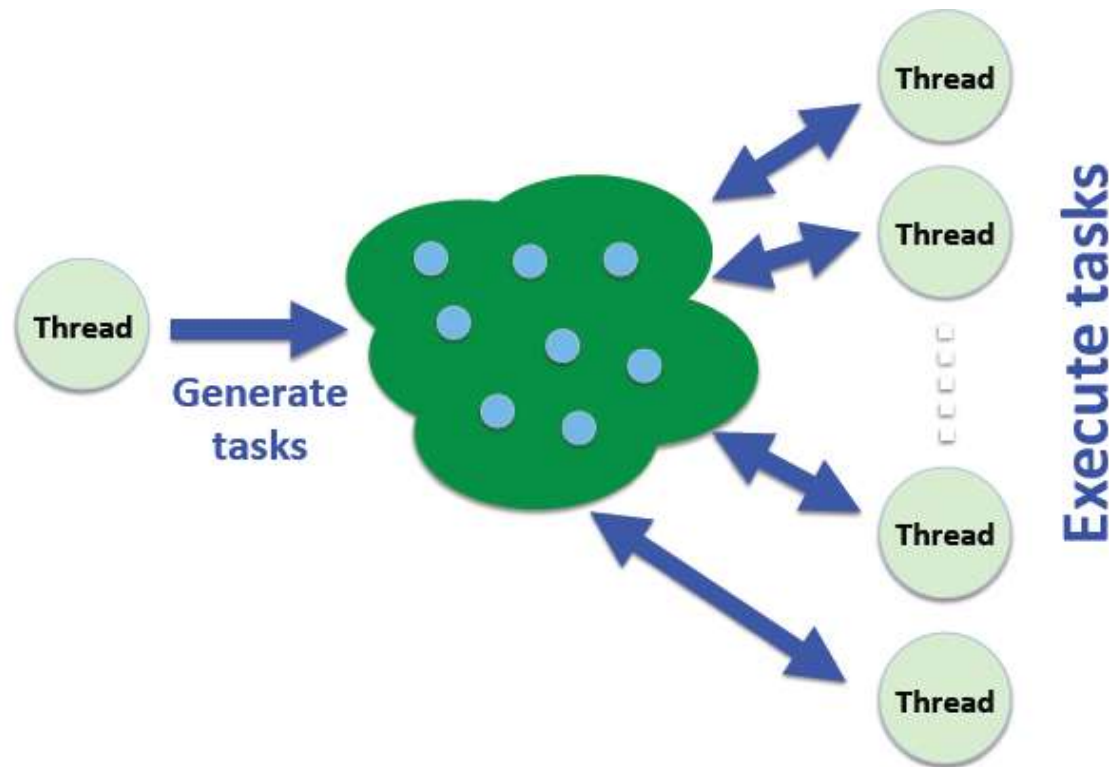
#pragma omp parallel
#pragma omp sections
{
    #pragma omp section
    { { a=...;
      b=...; }
    #pragma omp section
    { c=...;
      d=...; }
    #pragma omp section
    { e=...;
      f=...; }
    #pragma omp section
    { g=...;
      h=...; }
} /*omp end sections*/
} /*omp end parallel*/

```



# Tasks

- Feature added to version 3.0 of OpenMP
- A task is: an independent unit of work
- A thread is assigned to perform a task.



Source:  
Ruud van der Pass  
SC'13

# Tasks Example

```
#pragma omp parallel {
```

```
    #pragma omp single {
```

```
        node *p = head_of_list;
```

```
        while (p) {
```

```
            #pragma omp task private(p)
```

```
                process(p);
```

```
                p = p->next;
```

```
            } // end while
```

```
        } //end pragma single
```

```
    } // end pragma parallel
```

Threads start executing  
tasks at that point.



Implicit barrier





# Task Synchronization

#pragma omp barrier

#pragma omp taskwait

- explicitly waits on the completion of  
child tasks

Example:  
Write a program that prints either  
"A Race Car" or "A Car Race"

```
#include <stdlib.h>
#include <stdio.h>

int main(int argc, char *argv[]) {

    printf("A ");
    printf("race ");
    printf("car ");

    printf("\n");
    return(0);
}
```

```
$ cc -fast hello.c
$ ./a.out
A race car
$
```

*What will this program print ?*

Source:  
Ruud van der Pass  
SC'13

Example:  
Write a program that prints either  
"A Race Car" or "A Car Race"

```
#include <stdlib.h>
#include <stdio.h>

int main(int argc, char *argv[]) {

    #pragma omp parallel
    {
        printf("A ");
        printf("race ");
        printf("car ");

    } // End of parallel region

    printf("\n");
    return(0);
}
```

*What will this program print  
using 2 threads ?*

Source:  
Ruud van der Pass  
SC'13

Example:  
Write a program that prints either  
"A Race Car" or "A Car Race"

```
#include <stdlib.h>
#include <stdio.h>

int main(int argc
```

*What will this program print  
using 2 threads ?*

```
    #pragma omp parallel
    {
        #pragma omp single
        {
            printf("A ");
            printf("race ");
            printf("car ");
        }
    } // End of parallel region

    printf("\n");
    return(0);
}
```

Source:  
Ruud van der Pass  
SC'13

Example:  
Write a program that prints either  
"A Race Car" or "A Car Race"

```
int main(int argc, char *argv[]) {  
    #pragma omp parallel  
    {  
        #pragma omp single  
        {  
            printf("A ");  
            #pragma omp task  
            {printf("race ");}  
            #pragma omp task  
            {printf("car ");}  
        }  
    } // End of parallel region  
  
    printf("\n");  
    return(0);  
}
```

*What will this program print  
using 2 threads ?*

Source:  
Ruud van der Pass  
SC'13

Example:  
Write a program that prints either  
"A Race Car" or "A Car Race"

```
int main(int argc, char *argv[]) {  
    #pragma omp parallel  
    {  
        #pragma omp single  
        {  
            printf("A ");  
            #pragma omp task  
            {printf("race ");}  
            #pragma omp task  
            {printf("car ");}  
            printf("is fun to watch ");  
        }  
    } // End of parallel region  
  
    printf("\n");  
    return(0);  
}
```

***What will this program print  
using 2 threads ?***

A is fun to watch race car  
\$ ./a.out

A is fun to watch race car  
\$ ./a.out

A is fun to watch car race

Source:  
Ruud van der Pass  
SC'13



Example:  
Write a program that prints either  
"A Race Car" or "A Car Race"

```
int main(int argc, char  
#pragma omp parallel  
{  
    #pragma omp single  
    {  
        printf("A ");  
        #pragma omp task  
        {printf("car ");}  
        #pragma omp task  
        {printf("race ");}  
        #pragma omp taskwait  
        printf("is fun to watch ");  
    }  
} // End of parallel region  
  
printf("\n"); return(0);  
}
```


*What will this program  
print using 2 threads ?*

```
A car race is fun to watch  
$ ./a.out  
A car race is fun to watch  
$ ./a.out  
A race car is fun to watch  
^
```

Source:  
Ruud van der Pass  
SC'13

# Dividing Work Among Threads

Specifies that the enclosed code is to be executed by **only one thread** in the team.



```
#pragma omp single [clause ...]
```

```
structured_block
```



# Conclusions

- We have seen three mechanisms to enforce mutual exclusion: atomic, critical, and locks
  - atomic is fastest but with limitations
  - critical can name sections but at compile time
  - locks are slowest but sometimes are the only option